

Complexity, and the Array

DSA27 — Lecture 02 · Data Structures and Algorithms

Dr. Haitham A. El-Ghareeb

Week 2 · Fall 2026

Contents

Where We Are	2
Last week, in one line	2
Today	2
What Do We Count?	3
Two ways to answer “how fast?”	3
The model: basic steps	3
Counting a loop	3
Counting a nested loop	4
Throw away what does not matter	4
Big-O, Big-Ω, Big-Θ	5
Big-O: an upper bound	5
A proof, once	6
Big-Ω and Big-Θ	6
The classes you will meet	7
Best, Worst, Average	8
One algorithm, many costs	8
Measured: same n, different cost	8
Time and space	9
Reading the Cost Off the Code	9
Five rules	9
Rule 5: the hidden loop	10
The Array	11
What an array is	11
In C, the array is the machine	11
Arrays in Python: four candidates	12
Python stores references	12
The Course Array	13
dsa/array.py — given to you	13
What it will not do	13
The operations, and what they cost	14
Insertion, slot by slot	15
Measured, not asserted	15

Two dimensions, one block	16
The Rule, From Today	16
No list, dict or set inside a structure	17
The same idea, all term	18
This Week	18
Exercises: dsa/array_ops.py	18
Homework 2 — before Lecture 03	19
Summary	20
Seven things to keep	20
Next	20
Sources and further reading	20

How to read this document. This is the handout for Lecture 02. It holds everything on the slides, plus what I said out loud and the reasoning behind each claim. Two things are new this week and both matter for the rest of the term: the **language of cost** (Big-O and its relatives), and the **Array**, the one storage primitive every structure in this course will be built on.

Slides: DSA27-L02-slides.pdf · Code: dsa/array.py, dsa/array_ops.py · Tests: tests/test_array.py, tests/test_array_ops.py

Where We Are

Last week, in one line

An **ADT** is a contract. A **data structure** keeps the contract. Choosing between structures **is the engineering** — and you cannot choose without knowing what each one **costs**.

Today: **how to state a cost**, and **the first structure whose costs we state**.

Lecture 01 ended with a promise: complexity comes before the structures, because every choice of structure is a choice about cost. This lecture keeps that promise twice over. The first half gives you the vocabulary — Big-O, Big-Ω, Big-Θ, best, worst and average case, time and space. The second half applies it to the simplest structure there is, the **array**, and introduces the rule that governs every structure you write from now on.

Today

1. What we count, and why we throw most of it away
2. Big-O, Big-Ω, Big-Θ — the definitions, and one proof
3. Best, worst and average case; time and space
4. Rules for reading the cost off code
5. **The array**: what it is in memory, and what it costs
6. Arrays in Python — and **the course Array**
7. **The rule** from today on: no list, dict or set inside a structure

What Do We Count?

Two ways to answer “how fast?”

Measure	Count
time it with a clock	count the basic steps
true for this machine, this input, today	true for every machine
gives you the constants	gives you the shape
<code>viz.complexity.measure</code>	pen and paper

You need both. This lecture is mostly about counting — and ends by measuring.

A stopwatch answers “how long did it take?”, which depends on the processor, the Python version, what else was running, and the input you happened to pick. A count answers “how does the work **grow** as the input grows?”, which depends only on the algorithm.

The count is what lets you predict. If an algorithm does n^2 steps and takes one second at $n = 10,000$, it will take about a hundred seconds at $n = 100,000$ — on any machine. The stopwatch cannot tell you that; the count can.

The model: basic steps

A **basic step** takes a constant amount of time, whatever the input size:

- arithmetic and comparison on small numbers: $a + b$, $x < y$
- assigning a variable: $x = y$
- reading or writing **one** array slot: $a[i]$, $a[i] = v$
- calling a function, returning from one (not what the function then does)

The cost of an algorithm is **the number of basic steps**, as a function of the input size **n**.

This is the *RAM model* (random-access machine): memory is a long row of slots, and reaching any slot costs the same. It is a simplification — real machines have caches, and Lecture 01 warned that Python integers are not fixed-size — but it is the right simplification, because it keeps what matters for large n and throws away what does not.

Note the third bullet. **Reading one array slot is a basic step**. That is not an assumption we get for free; it is a property of how arrays are laid out in memory, and the second half of this lecture explains why it is true.

What is n ? Whatever measures the size of the input: the number of elements in an array, the number of characters in a string, the number of nodes in a graph. Sometimes there are two: an $n \times m$ matrix, or a graph with V nodes and E edges.

Counting a loop

```
def total(arr):
    result = 0                # 1 step
```

```

for i in range(len(arr)):  # the body runs n times
    result += arr[i]       # 2 steps: read, add-and-assign
return result             # 1 step

```

$$T(n) = 2n + 2$$

Count it honestly and you get something like $2n + 2$. Count it more pedantically — the loop itself compares and increments a counter — and you get $4n + 3$, or $5n + 2$. Nobody agrees on the exact constants, and **it does not matter**, for a reason the next slide makes precise: every one of those formulas is “a constant times n , plus a constant”, and they all grow the same way.

Counting a nested loop

```

def count_pairs(arr):
    n, pairs = len(arr), 0
    for i in range(n):
        for j in range(i + 1, n):  # runs n-1, n-2, ..., 1, 0 times
            if arr[i] == arr[j]:
                pairs += 1
    return pairs

```

$$(n - 1) + (n - 2) + \dots + 1 + 0 = \frac{n(n - 1)}{2} = \frac{1}{2}n^2 - \frac{1}{2}n$$

This is exercise 7 from Lab 02, `count_pairs`, and now you can see where the formula comes from. The inner loop runs $n - 1$ times for $i = 0$, $n - 2$ for $i = 1$, and so on down to 0. Adding $1 + 2 + \dots + (n - 1)$ is the sum Gauss is said to have done as a schoolboy: pair the first with the last, and you have $(n - 1)/2$ pairs each summing to n .

The leading term is $\frac{1}{2}n^2$. Double n and the work goes up about four times.

Throw away what does not matter

n	$\frac{1}{2}n^2$	$\frac{1}{2}n$	share of the total from $\frac{1}{2}n^2$
10	50	5	91%
1,000	500,000	500	99.9%
1,000,000	5×10^{11}	500,000	99.9999%

For large n , **the fastest-growing term is everything**. And the constant in front of it ($\frac{1}{2}$) depends on the machine and on how you count. So we drop both:

$$\frac{1}{2}n^2 - \frac{1}{2}n \longrightarrow n^2$$

Two things are being thrown away, for two different reasons.

Lower-order terms go because they stop mattering. By $n = 1,000$ the $\frac{1}{2}n$ term is a rounding error.

Constant factors go because they are not a property of the algorithm. The same algorithm has a different constant in C and in Python, on a laptop and on a server, counted by me and counted by you. What stays the same everywhere is the *shape*: n^2 , n , $\log n$.

This is the deal Big-O offers. You give up the constants, and in return you get a statement that is true on every machine and that tells you what happens when n grows. It is a very good deal for large n and a misleading one for small n — which is why we still measure.

Big-O, Big-Ω, Big-Θ


Big-O: an upper bound

$$f(n) = O(g(n))$$

means: there are constants $c > 0$ and n_0 such that

$$f(n) \leq c \cdot g(n) \quad \text{for every } n \geq n_0.$$

“From some point on, f never grows faster than a constant times g .”



Read the definition slowly, because every part of it is doing a job.

- **“There are constants c and n_0 ”** — you get to choose them. Proving $f(n) = O(g(n))$ means finding *one* pair that works.
- **“ $c \cdot g(n)$ ”** — the constant factor is allowed. That is how the definition throws constants away: $1000n$ is $O(n)$, because $c = 1000$ works.

- “**for every $n \geq n_0$** ” — only large n counts. Below n_0 , f may be as badly behaved as it likes. That is how lower-order terms disappear.
- “ $\leq$ ” — it is an **upper bound**. It says f grows *no faster* than g . It does not say f grows *as fast as* g .

That last point surprises people: $n = O(n^2)$ is **true**. It is just not very informative, like saying “the lecture room holds fewer than a million people”. When we say “linear search is $O(n)$ ” we mean, informally, the *tightest* upper bound we can prove — but the notation itself only promises “no worse than”.

In Arabic you will see رتبة التعقيد (order of complexity) or simply O الكبيرة.

A proof, once

Claim. $3n + 5 = O(n)$.

Proof. Choose $c = 4$ and $n_0 = 5$. For every $n \geq 5$:

$$3n + 5 \leq 3n + n = 4n.$$

So $3n + 5 \leq 4 \cdot n$ for all $n \geq 5$. ■

The only step that needs thought is “ $5 \leq n$ ”, which is true exactly because we chose $n_0 = 5$. Other choices work too: $c = 8$ and $n_0 = 1$ ($3n + 5 \leq 3n + 5n = 8n$ for $n \geq 1$). The definition asks for *some* c and n_0 , not the best ones.

The same idea shows that **any polynomial is Big-O of its highest power**: $a_k n^k + \dots + a_1 n + a_0 = O(n^k)$. Replace every lower power by n^k (for $n \geq 1$ each is smaller), add up the coefficients, and that sum is your c .

And to show something is **not** $O(g)$, show that no c can work: n^2 is not $O(n)$, because $n^2 \leq c \cdot n$ would need $n \leq c$ for *every* large n , and no constant is bigger than every n .

You will be asked to do one proof like this in the exam. You will be asked, far more often, to *read off* the Big-O of a piece of code — which the rules later in this lecture make mechanical.

Big-Ω and Big-Θ

Notation	Means	Like
$f = O(g)$	$f \leq c \cdot g$ eventually	$\leq$ — an upper bound
$f = \Omega(g)$	$f \geq c \cdot g$ eventually	$\geq$ — a lower bound
$f = \Theta(g)$	both: $c_1 g \leq f \leq c_2 g$	$=$ — a tight bound

$\frac{1}{2}n^2 - \frac{1}{2}n$ is $O(n^2)$, $\Omega(n^2)$ and therefore $\Theta(n^2)$. It is also $O(n^3)$ — true, and useless.

Ω (Omega) is the mirror image of O : there are $c > 0$ and n_0 such that $f(n) \geq c \cdot g(n)$ for all $n \geq n_0$. “ f grows *at least* as fast as g .”

Θ (Theta) is both at once. $f = \Theta(g)$ means f is sandwiched between two constant multiples of g : it grows *exactly* like g , up to constants. That is usually what people mean when they say “Big-O” in conversation.

In this course, and in most interviews, “the complexity is $O(n)$ ” is understood as “ $\Theta(n)$, and that is tight”. In a proof or an exam answer, use the symbol you can actually prove.

One more use of Ω that matters later: **a lower bound on a problem**, not an algorithm. Any sorting algorithm that works only by comparing elements needs $\Omega(n \log n)$ comparisons in the worst case. No cleverness gets under it. In Week 9, counting sort gets under it by *not comparing* — which is exactly the loophole the bound leaves open.

The classes you will meet

Big-O	Name	Example in this course	$n = 1,000,000$
$O(1)$	constant	<code>a[i]</code> ; push onto a stack	1
$O(\log n)$	logarithmic	binary search (Week 8)	20
$O(n)$	linear	linear search; insert into an array	10^6
$O(n \log n)$	linearithmic	merge sort, heap sort (Week 10)	2×10^7
$O(n^2)$	quadratic	bubble, selection, insertion sort	10^{12}
$O(2^n)$	exponential	all subsets (Week 3)	—




Figure 1: The same curves as Lecture 01. Note the logarithmic vertical axis.

Learn this table. The right-hand column is the one to feel: at a billion simple steps a second, 10^6 steps is a millisecond, 2×10^7 is a fiftieth of a second, and 10^{12} is **more than a quarter of an**

hour. For 2^n at $n = 1,000,000$ there is no useful number — the universe is younger.

Where does $\log n$ come from? From **halving**. Lab 01's Checkpoint 4 halved 100 down to 1 in six steps; halving n down to 1 takes about $\log_2 n$ steps. Any algorithm that throws away half of what is left at each step is $O(\log n)$.

Best, Worst, Average

One algorithm, many costs

```
def find(arr, target):
    for i in range(len(arr)):
        if arr[i] == target:
            return i
    return -1
```

Case	When	Comparisons
Best	target is <code>arr[0]</code>	$1 - O(1)$
Worst	target is last, or absent	$n - O(n)$
Average	present, every position equally likely	$\frac{n+1}{2} - O(n)$

The cost of an algorithm is not one function of n ; it depends on *which* input of size n you give it. So we name three:

Worst case — the maximum over all inputs of size n . This is the default: when someone says “linear search is $O(n)$ ” they mean the worst case. It is a **guarantee**: no input can do worse.

Best case — the minimum. Rarely useful on its own. “My sort is $O(n)$ in the best case” usually means “on input that was already sorted”, which is not a reason to choose it.

Average case — the expected cost over some *distribution* of inputs. It needs an assumption, and you must state it. For linear search, if the target is present and equally likely to be in any of the n positions, the expected number of comparisons is $(1 + 2 + \dots + n)/n = (n + 1)/2$. Still $O(n)$: on average you search half the array, and half is a constant.

Do not confuse these with O , Ω and Θ . “Worst case” says *which input*; “Big- O ” says *what kind of bound*. “The worst case of linear search is $\Theta(n)$ ” is a perfectly sensible sentence that uses both.

Measured: same n , different cost

This is the cost of inserting one element into an array of 20,000, depending on *where*. At index 0 every element must shift; at the end, none do. Same operation, same n , a factor of thousands between best and worst. The straight line is the cost being exactly proportional to the number of elements after the insertion point, $n - i$. Averaged over every position that is $n/2 - O(n)$.

You will produce the points on that line yourself, with your own `insert_at` from `dsa/array_ops.py`.




Figure 2: Insertion into an array of 20,000 elements, at nine different positions. Real timings on the course Array.

Time and space

Space complexity: how the *extra* memory an algorithm needs grows with n .

Function	Extra space	Why
<code>find(arr, x)</code>	$O(1)$	two variables, whatever n is
<code>reverse_in_place(arr)</code>	$O(1)$	swaps within the array
<code>resized(arr, 2n)</code>	$O(n)$	a whole new array
<code>merge_sorted(a, b)</code>	$O(n + m)$	the output array
<code>recursive total(values)</code>	$O(n)$	one stack frame per element

We usually count **auxiliary** space — memory beyond the input itself. Reversing an array in place needs two indices and one temporary: $O(1)$. Copying it needs n new slots: $O(n)$.

The last row is the one people forget. **Recursion uses space.** Every call that has not yet returned holds a frame on the call stack (Lab 03, Part 10). A recursive function n calls deep uses $O(n)$ stack space even if it creates no data at all — and Python stops it at about 1,000 frames.

Time and space often trade against each other. Merge sort (Week 10) buys its $O(n \log n)$ time with $O(n)$ extra space; heap sort gets the same time in $O(1)$ extra space and pays in other ways. Knowing which you are paying for is know-why.

Reading the Cost Off the Code

Five rules

1. **Sequence — add.** $O(f)$ then $O(g)$ is $O(f + g)$ = the bigger one.
2. **Loop — multiply.** A loop of n iterations with an $O(g)$ body is $O(n \cdot g)$.

3. **Nested loops — multiply again.** Two full nested loops over n : $O(n^2)$.
4. **Halving — logarithm.** A loop that halves what is left is $O(\log n)$.
5. **Function calls cost what the function costs** — including built-ins.

Rule 1: an $O(n)$ loop followed by an $O(n^2)$ loop is $O(n + n^2) = O(n^2)$. The bigger term swallows the smaller.

Rule 2 needs care about *what the body costs*. A body of three basic steps is $O(1)$, so n iterations cost $O(n)$. But a body that itself does $O(n)$ work — see rule 5 — makes the loop $O(n^2)$.

Rule 3 is rule 2 applied twice. The triangular loop (j from $i + 1$) is still $O(n^2)$: it does half the work of the square loop, and half is a constant.

Rule 4 is where $\log n$ comes from:

```
steps = 0
while n > 1:
    n //= 2
    steps += 1          # runs about log2(n) times
```

Rule 5 is the one that catches Python programmers, and it is why the second half of this lecture exists.

Rule 5: the hidden loop

```
def has_duplicate(values):
    for i in range(len(values)):
        if values[i] in values[i + 1:]:    # <- how much does this line cost?
            return True
    return False
```

One line, **two hidden $O(n)$ loops**: the slice **copies** up to n elements, and **in** **searches** them one by one. Inside an n -iteration loop: $O(n^2)$.

A call is only $O(1)$ if you **know** it is $O(1)$.

This is the most important slide for Python programmers, and the reason this course is built the way it is.

`values[i + 1:]` looks like one step. It builds a new list and copies every element after i into it — $O(n)$. `x in some_list` looks like one step. It walks the list comparing — $O(n)$. `list.insert(0, x)` looks like one step. It shifts every element — $O(n)$. `list.pop(0)`, `list.remove(x)`, `list.index(x)`, `sorted(xs)`: all look like one step, none is.


Python's built-ins are superbly engineered, and their costs are documented. But they are *hidden* inside one short line, and a student who has never had to write them does not see them. You cannot count the cost of code you cannot see.

So, from today, you are going to see it.

The Array

What an array is

A **fixed number** of **equal-sized slots**, stored **side by side** in memory.



$$\text{address of } a[i] = \text{base} + i \times \text{slot size}$$

One multiplication and one addition — **whatever i is, whatever n is**. That is why $a[i]$ is $O(1)$.

Everything about the array follows from that picture.

$O(1)$ indexing comes from two properties together: the slots are **contiguous** (no gaps, one block), and they are **equal-sized**. Then the address of slot i is arithmetic, not a search. Take away either property and indexing stops being $O(1)$ — a linked list (Week 5) keeps its elements anywhere in memory, so reaching element i means following i links.

Fixed size is the price. The block was reserved at a particular size, and the memory right after it may already belong to something else. The array cannot grow in place. To get a bigger one you must reserve a new block and copy — $O(n)$. That is the whole of Week 4.

Why indices start at 0. Because i is an *offset* from the base: element 0 is zero slots from the start. Languages that start at 1 must subtract 1 somewhere.

Array in Arabic is متجاورة. You will also hear *contiguous memory* as متجاورة ذاكرة.

In C, the array is the machine

```
int a[6] = {17, 4, 42, 8, 15, 23}; // 6 × 4 bytes, reserved once
a[2] = 99;                        // base + 2 × 4: one instruction
a[6] = 0;                         // no check – writes over whatever is next
```

- The size is part of the type, fixed at compile time
- Every slot holds an `int`, so every slot is 4 bytes
- **No bounds check:** an out-of-range index is not an error, it is a bug

C is the language that most directly exposes what an array is — which is why it is worth two minutes, even though we write Python. (Kernighan and Ritchie, the book Lecture 01 recommended, spends a chapter on it.)

The last line is famous. C does not check the index. `a[6]` computes `base + 6 × 4` and writes there, over whatever variable happens to live next. This is the root of a whole family of security holes, *buffer overflows*. Every array in this course **does** check, and raises `IndexError` — the extra comparison is $O(1)$ and worth it.

Arrays in Python: four candidates

	Grows?	Holds	Hides the cost?
<code>list</code>	yes — a dynamic array	references to any objects	yes : <code>append</code> , <code>insert</code> , <code>pop(0)</code> , <code>in</code> , slicing
<code>array.array</code>	yes	numbers only, one type, packed	yes: same methods as <code>list</code>
<code>numpy.ndarray</code>	no	numbers, one type, packed	yes: whole-array operations in C
<code>ctypes array</code>	no	C values, or references	no — <code>index</code> , <code>assign</code> , nothing else

Python does not have a plain, C-style array as a built-in type. What it has:

list is itself a **dynamic array** — the structure you will build in Week 4. Under the hood it is exactly a block of references plus a size and a capacity. Every convenient method on it is an algorithm with a cost, already written for you.


array.array (the standard-library `array` module) stores numbers packed tightly — `array('i', [1, 2, 3])` uses 4 bytes per element instead of a reference to an `int` object. It is compact, and useful for large numeric data. But it can `append`, `insert` and `pop` just like a `list`: it is a dynamic array too.

NumPy arrays are fixed-size and packed, and the backbone of data science and of most of what the AI program does with data. But their power is whole-array operations — `a + b` adds a million numbers in one line, in C. Wonderful for computation; useless for learning what one insertion costs.

ctypes arrays are real C arrays, reached from Python. They have a fixed size, and you can index them and assign to them — nothing else. That is exactly the primitive we want.

Python stores references

A Python array stores references, so its slots can be equal even when the objects are not



A Python array slot does not hold 42; it holds **a reference to** the object 42. Slots are equal-sized — one reference each — even when the objects are not. That is how one array holds an int, a string and a list, and still indexes in $O(1)$.

This is the ladder from Lecture 01 again: a name is bound to an object, and a collection is a collection of references. The slot size is the size of one reference (8 bytes on a 64-bit machine), so the address arithmetic works.

The cost is a second hop: `a[i]` finds the reference in $O(1)$, then follows it to the object, which lives wherever Python put it. Packed arrays (`array.array`, NumPy) avoid that hop by storing the numbers themselves, which is why they are faster for numerical work — and why they can hold only one type.

The Course Array

`dsa/array.py` — given to you

```
from dsa.array import Array

a = Array(5)           # 5 slots, all None. Fixed for ever.
a[0] = 42              # O(1)
a[0]                  # O(1) -> 42
len(a)                # 5: the capacity, not "how many are used"
list(a)               # iterate: [42, None, None, None, None]

Array.from_values([3, 1, 2]) # a full array, for tests and the REPL
Array(3, fill=0)          # every slot starts as 0
```

A real C array of references (ctypes), wrapped so that it **checks its bounds**.

`dsa/array.py` is infrastructure, like `viz/`: already written, read it, do not change it. It is thirty lines, and every one is worth reading.

Three details:

- **`len(a)` is the capacity.** An `Array(5)` always has length 5, whether you have put anything in it or not. Tracking how many slots are *in use* — the **size** — is your job, in a separate variable. That split, capacity against size, is the heart of the dynamic array.
- **Every slot starts as `None`** (or the `fill` you give). A raw ctypes slot starts out `NULL` — reading it raises a confusing `ValueError` — so the class fills them.
- **`from_values`** copies an iterable into a new, exactly-full `Array`. It is for building test data and experimenting at the REPL. Do not use it inside a structure to dodge the rules — converting to a tuple and back is exactly the kind of hidden $O(n)$ this lecture is about.

What it will not do

```
>>> a = Array(3)
>>> a[3]
```

```

IndexError: Array index 3 out of range for length 3 ...
>>> a[-1]
IndexError: ... # no negative indices, as in C
>>> a[0:2]
TypeError: Array does not support slicing; copy with a loop
>>> a.append(1)
AttributeError: 'Array' object has no attribute 'append'

```

No growing. No shrinking. No inserting. No slicing. No searching methods. **Every one of those is an algorithm — and you are going to write them.**

Each refusal is deliberate.

No negative indices, because an array index is an offset from the base, and there is nothing before the base. When you build `DynamicArray` in Week 4 you will *implement* negative indexing (`arr[-1]` is the last element), and you will see that it is a translation, `i + size`, on top of an array that has none.

No slicing, because a slice is a copy — $O(k)$ — dressed as one step. If you want a copy, write the loop and see the cost.

No append, because an array cannot grow. That is the definition.

The operations, and what they cost


Operation	How	Cost
read / write slot <code>i</code>	address arithmetic	$O(1)$
find a value	check slots one by one	$O(n)$
insert at <code>i</code>	shift <code>i .. size-1</code> right , then write	$O(n)$
remove at <code>i</code>	shift <code>i+1 .. size-1</code> left	$O(n)$
insert / remove at the end	nothing to shift	$O(1)$
grow	new array, copy everything	$O(n)$

This table is the array's **contract** — its ADT, with costs. Every row is a function in `dsa/array_ops.py` or a consequence of one, and every structure you build this term inherits these costs from the array underneath it.

Look at the two rows about the end. Inserting and removing at the *end* is $O(1)$, because nothing has to move. That single fact decides:

- that a **stack** should keep its top at the *end* of an array (Week 6);
- that a **queue** cannot simply remove from the front of an array — which is why Week 7 needs a ring buffer;
- that a **heap** adds at the end and repairs upward (Week 12).

Insertion, slot by slot



insert X at index 2: shift $5 - 2 = 3$ elements right, starting from the END

```
def insert_at(arr, size, index, value):
    for i in range(size, index, -1): # from the END, backwards
        arr[i] = arr[i - 1]
    arr[index] = value
    return size + 1
```

This is the core of `insert_at` from this week's exercise, without the error checks (the full contract is in the docstring: `OverflowError` when full, `IndexError` for a bad index).

Why backwards? Shift forwards — copy `arr[2]` into `arr[3]` first — and you have overwritten `d` before moving it; every slot ends up holding `c`. Moving right, you must start at the right. Moving left (removal), you start at the left. Trace both on paper once; it is the most common bug in this week's exercises.

The loop runs $\text{size} - \text{index}$ times. Index 0: size moves, the worst case. Index size : zero moves, which is just writing at the end.

Measured, not asserted

Three operations on the course `Array`, timed with `viz.complexity.measure` as n doubles from 1,024 to 32,768. Both axes are logarithmic, so a straight line climbing one grid step per doubling is $O(n)$, and a flat line is $O(1)$.

- **Indexing** stays at about a quarter of a microsecond whatever n is. (One access is too fast to time on its own, so the figure times 10,000 and divides — an honest measurement has to know its own limits.)
- **Linear search for a missing value** doubles when n doubles: $O(n)$.
- **Inserting at index 0** also doubles: $O(n)$ — and it sits *above* the search, because each step does a read *and* a write where search does a read and a comparison. Same Big-O, different constant: exactly the thing Big-O throws away, and exactly the thing a measurement shows you.

At $n = 32,768$, one index is roughly 25,000 times cheaper than one front insertion. That is the





Figure 3: Real timings on the course Array. Each step right doubles n.

difference between $O(1)$ and $O(n)$, in microseconds.

Two dimensions, one block



A matrix is stored as **one** array, row after row: element (r, c) lives at index $r \times \text{cols} + c$.

Memory is one-dimensional, so a 2-D array has to be flattened. **Row-major** order — row 0, then row 1, ... — is what C and NumPy use by default; Fortran and MATLAB use column-major. The formula is address arithmetic again, so (r, c) is still $O(1)$.

The alternative is an array of arrays: an Array of rows slots, each holding an Array of cols. It works, costs an extra hop per access, and lets rows have different lengths. `transpose_flat` in this week's exercises uses the flat layout, so that you work the index arithmetic by hand.

The Rule, From Today

No list, dict or set inside a structure

A data structure in `dsa/` keeps its data **only** in:

1. the course **Array** (`dsa/array.py`), or
2. **node objects** — `Node`, `TreeNode`, `Entry`, ..., or
3. **another structure you have already built** in `dsa/`.

Never in a Python `list`, `dict`, `set` or `collections.deque`.

Lists remain the **interface**: functions may take a list and return a list, and the tests use them freely.

This is the rule that governs the rest of the course. Here is exactly how the repository applies it, so that you can see the structures stacking up week by week:

Week	Structure	Stores its data in
4	<code>DynamicArray</code>	an <code>Array</code> , replaced by a bigger one when full
5	<code>LinkedList</code>	<code>Node</code> objects
6	<code>Stack</code>	your <code>DynamicArray</code> — top at the end
7	<code>SlowQueue</code> / <code>CircularQueue</code>	your <code>DynamicArray</code> / an <code>Array</code> as a ring
11	<code>BinarySearchTree</code>	<code>TreeNode</code> objects
12	<code>MinHeap</code> , <code>PriorityQueue</code>	your <code>DynamicArray</code> / your <code>MinHeap</code>
13	<code>ChainingHashMap</code> / <code>OpenAddressingHashMap</code>	an <code>Array</code> of <code>Entry</code> chains / two <code>Arrays</code>
14	<code>Graph</code> / <code>MatrixGraph</code>	your hash map and <code>DynamicArrays</code>

Read down the right-hand column. By Week 14 your graph is built on your hash map, which is built on your array; your breadth-first search runs on your queue, which runs on your array. Nothing in the stack is borrowed. When something is slow, you will know which line of *your* code made it slow.

Why not lists? Not because lists are bad — they are excellent, and you will use them for the rest of your career. Because a list hides, behind one convenient method, exactly the cost this course exists to teach. A stack built on `list.append` teaches you nothing about why a stack is $O(1)$. A stack built on an array you had to grow yourself does.

What is still allowed. Python lists as *function inputs and return values* — `bfs` returns a list of nodes, sorting functions take one — because that is the convenient interface for tests and for you. Plain local variables. Tuples to return two values. What is not allowed is a `list`, `dict` or `set` as the *storage* of a structure, or as the working storage of an algorithm that has a structure of yours available (use your `CircularQueue` for BFS, not a list).

The TAs will check this when they mark. A structure that passes every test by hiding a `list` inside scores as not done.

The same idea, all term

Array	(given, $O(1)$ index, fixed size)
└─ DynamicArray	(Week 4: grow by copying – amortised $O(1)$ append)
└─ Stack	(Week 6)
└─ SlowQueue	(Week 7 – <code>pop(0)</code> is $O(n)$, measure it)
└─ MinHeap	(Week 12)
└─ PriorityQueue	
└─ CircularQueue	(Week 7 – $O(1)$ both ends, on a ring)
└─ hash maps	(Week 13)
└─ Graph	(Week 14, with DynamicArray)

Every arrow is “is built on”. Linked lists and trees sit beside this tree, built from nodes rather than arrays — and Week 5 exists to compare the two families directly.

This Week

Exercises: `dsa/array_ops.py`

Function	Target
<code>find(arr, target, size=None)</code>	$O(n)$ worst, $O(1)$ best
<code>is_sorted(arr, size=None)</code>	$O(n)$, stop early
<code>insert_at(arr, size, index, value)</code>	$O(n)$ — shift right, from the end
<code>remove_at(arr, size, index)</code>	$O(n)$ — shift left
<code>reverse_in_place(arr)</code>	$O(n)$ time, $O(1)$ space
<code>rotate_left(arr, k)</code>	$O(n)$ time, $O(1)$ space — three reversals
<code>resized(arr, new_length)</code>	$O(n)$ — the dynamic array’s one move
<code>merge_sorted(a, b)</code>	$O(n + m)$ — half of merge sort
<code>transpose_flat(arr, rows, cols)</code>	$O(\text{rows} \times \text{cols})$

```
pytest tests/test_array.py -v      # the Array itself – passes already
pytest tests/test_array_ops.py -v  # your exercises
```

Every function takes an `Array` and works with indices only: no slicing, no converting to a list, no `list` scratch space. Where a function builds a new array, it builds an `Array`.

Several of them take a `size`: the array is only partly filled, and the slots from `size` on are free. That is how `insert_at` and `remove_at` can change “how many elements there are” without the array changing length.

`rotate_left` is the one that needs an idea. Rotating `[1 2 3 4 5]` left by 2: reverse the first 2 $\rightarrow$ `[2 1 3 4 5]`; reverse the rest $\rightarrow$ `[2 1 5 4 3]`; reverse the whole $\rightarrow$ `[3 4 5 1 2]`. Three reversals, $O(n)$ time, no second array. Convince yourself on paper *why* it works before you write it.

Homework 2 — before Lecture 03

1. **Implement** `dsa/array_ops.py` until `pytest tests/test_array_ops.py` passes.
2. **Measure** your `insert_at` at index 0 and at the end, for $n = 1,000 \dots 32,000$, with `viz.complexity.measure`, and plot both with `plot_growth(..., reference=["1", "n"])`. One paragraph: do the curves match the table?
3. **Prove** $5n^2 + 3n + 7 = O(n^2)$: give c and n_0 and show the inequality.
4. **Read off** the Big-O of the three snippets in the handout, and say which rule you used for each.
5. **Lab 02** exercises, if not finished in the lab.

For item 2, the Lecture 01 example shows the shape of the call. A starting point:

```
from dsa.array import Array
from dsa.array_ops import insert_at
from viz.complexity import measure, plot_growth

sizes = [1000, 2000, 4000, 8000, 16000, 32000]

def make(n):
    arr = Array(n + 1)
    for i in range(n):
        arr[i] = i
    return arr, n

front = measure(lambda a: insert_at(a[0], a[1], 0, -1), sizes, make)
end = measure(lambda a: insert_at(a[0], a[1], a[1], -1), sizes, make)
plot_growth({"insert at 0": front, "insert at end": end}, reference=["1", "n"])
```

For item 4, the three snippets:

```
# (a)
for i in range(n):
    for j in range(10):
        total += i * j

# (b)
i = n
while i > 0:
    for j in range(n):
        count += 1
    i //= 2

# (c)
for i in range(n):
    if find(arr, i) != -1:           # arr has n elements
        hits += 1
```

Hand in items 2–4 as one page (a PDF, or a notebook with the plot) through the channel the TAs announce.

Summary

Seven things to keep

1. Count **basic steps** as a function of n ; keep the **shape**, drop the constants and the lower-order terms.
2. $f = O(g)$: **eventually** $f \leq c \cdot g$. Ω is $\geq$; Θ is both.
3. **Worst case** is the guarantee. Say which case you mean.
4. **Sequence adds, loops multiply, halving is $\log n$** — and a call costs what the function costs.
5. An array is **contiguous, equal-sized slots**: $O(1)$ index, **fixed size**, $O(n)$ insert and search.
6. The course **Array**: index, assign, length. Nothing else.
7. From today: structures store data in **Array, nodes, or your own structures** — never list, dict or set.

Next

Week 3 — Recursion. Base case, recursive case, and the call stack — the first place where *how* you solve a problem decides its complexity.

Lab 03 this week ends with a bridge into it.

Sources and further reading

- **Python Software Foundation.** *TimeComplexity* — the documented costs of list, collections.deque, set and dict operations in CPython. The page to check before you believe any one-line Python operation is $O(1)$.
- **Python Software Foundation.** *ctypes* — *A foreign function library for Python*, section “Arrays”. The mechanism under `dsa/array.py`.
- **B. W. Kernighan and D. M. Ritchie.** *The C Programming Language*, 2nd ed., Prentice Hall, 1988, chapter 5, “Pointers and Arrays” — the address arithmetic on the array slides, from the people who designed it.
- **T. H. Cormen, C. E. Leiserson, R. L. Rivest and C. Stein.** *Introduction to Algorithms*, 3rd ed., MIT Press, 2009, chapter 3, “Growth of Functions” — the formal definitions of O , Ω and Θ used here.
- **M. T. Goodrich, R. Tamassia and M. H. Goldwasser.** *Data Structures and Algorithms in Python*, Wiley, 2013, chapters 3 and 5 — algorithm analysis, and array-based sequences built on ctypes, which is the approach this course takes.

Every figure in this lecture is generated by `tools/figures_102.py`; the two timing figures are real measurements and will differ slightly on your machine. Run the script and compare.