

Dynamic Arrays

DSA27 — Lecture 04 · Data Structures and Algorithms

Dr. Haitham A. El-Ghareeb

Week 4 · Fall 2026

Contents

Where We Are	2
Two weeks ago, one problem	2
Today	2
Capacity and Size	2
An Array plus a count	3
Reading and writing are still $O(1)$	3
Growing	4
Allocate, copy, switch	4
How Much to Grow?	5
Option 1: one slot at a time	5
Option 2: a constant step — say, 100 at a time	5
Option 3: double	5
The three side by side	6
Amortised Analysis	6
One expensive append, many cheap ones	7
Amortised cost: the definition	7
Amortised is not “average case”	8
Any constant factor works	8
Shrinking	9
Give memory back — carefully	9
Python’s list	9
Measured: CPython does exactly this	9
All the Operations	10
DynamicArray: the contract with costs	10
What it is built for	11
This Week	11
Exercises: dsa/dynamic_array.py	11
Homework 4 — before Lecture 05	12
Summary	12

Seven things to keep	12
Next	12
Sources and further reading	12

How to read this document. *This is the handout for Lecture 04. It holds everything on the slides, plus what I said out loud. This week you build the first real data structure of the course — the one Python's `list` already is — and meet the idea that makes it work: **amortised** cost.*

Slides: `DSA27-L04-slides.pdf` · Code: `dsa/dynamic_array.py` · Tests: `tests/test_dynamic_array.py`

Where We Are

Two weeks ago, one problem

The course `Array` gives $O(1)$ indexing — and a **fixed size**.

```
a = Array(4)
a[4] = "e"           # IndexError: there is no slot 4, and there never will be
```

Python's `list` grows for ever, and `append` is fast. **How?**

Lecture 02 introduced the array and its one great limitation: the size is decided when it is created. The memory right after the block may belong to something else, so it cannot grow in place.

Yet you have used Python lists since your first week, appending millions of items, and `append` never seemed slow. Today we build the structure that explains it — the **dynamic array** — on nothing but the course `Array`, and we prove why `append` is fast *on average* even though it is occasionally very slow.

Lecture 03's recursion is not needed today; complexity from Lecture 02 is needed throughout.


Today

1. Capacity and size: an `Array` plus a count
2. Growing: allocate, copy, switch
3. **How much** to grow — by 1, by 100, or $\times 2$?
4. **Amortised** analysis: why `append` is $O(1)$ on average
5. Shrinking, and how not to thrash
6. Python's own `list`, measured
7. The rest of the operations, and their costs

Capacity and Size

An Array plus a count

A dynamic array is an Array plus a count of how much of it is used



- **capacity** — how many slots the Array underneath has (`len(self._block)`)
- **size** — how many of them hold data (`self._size`, and what `len()` returns)
- Always $0 \leq \text{size} \leq \text{capacity}$

This is the split you already met in Lecture 02's `insert_at(arr, size, ...)`: the array has a length, and separately you track how much of it is in use. A dynamic array simply keeps both inside one object.

```
class DynamicArray:
    def __init__(self, values=(), growth=2):
        self._size = 0           # slots in use
        self._capacity = 1       # slots reserved
        self._block = make_block(1) # an Array(1)
        ...
```

That is the skeleton you are given in `dsa/dynamic_array.py`. `make_block` is just `Array(capacity)`.

While there is room, appending is one write into the next free slot and one increment: $O(1)$. The whole question of this lecture is what happens when there is no room.

Dynamic array in Arabic: المصفوفة الديناميكية — the array whose size changes.

Reading and writing are still $O(1)$

```
def __getitem__(self, index):
    if index < 0:
        index += self._size           # -1 means the last USED slot
    if not 0 <= index < self._size:
        raise IndexError(index)
    return self._block[index]
```

Negative indices are **translated**: the course Array has none, so the dynamic array provides them — in $O(1)$.


Two details that the tests check:

- The bounds are **0 to size – 1**, not 0 to capacity – 1. Slots beyond size exist in the Array but hold nothing the user put there; reading one must be an `IndexError`, exactly as with a Python list.
- `arr[-1]` is the last *used* element, so the translation adds size, not capacity.

That is essentially the whole of `__getitem__` — it is here because the translation is the idea worth seeing. `__setitem__` follows the same shape and is yours to write, along with everything that follows.

Growing

Allocate, copy, switch



A full array cannot grow in place. So:

1. allocate a **new, bigger** Array
2. **copy** every element across — $O(\text{size})$
3. **switch** `self._block` to the new Array (the old one is garbage-collected)

That is `_resize(capacity)` in the exercise, and it is the one genuinely expensive thing a dynamic array does:

```
def _resize(self, capacity):
    block = make_block(capacity)           # 1. allocate
    for i in range(self._size):             # 2. copy — O(size)
        block[i] = self._block[i]
    self._block = block                     # 3. switch
    self._capacity = capacity
    self.resize_count += 1                  # the tests and the notebook read this
```

(This one is shown in full because it is mechanical; deciding *when* to call it and *with what capacity* is the exercise.)

Step 3 is where Python's garbage collection, from Lecture 01, quietly does work: nothing refers to the old Array any more, so its memory is reclaimed. In C you would call `free` on it yourself — and forgetting would leak memory.

append now reads: *if full, resize; then write at size and increment*. A single append that triggers a resize costs $O(n)$. So is append $O(n)$? That depends entirely on **how much bigger** the new Array is.

How Much to Grow?

Option 1: one slot at a time

Capacity 1, 2, 3, 4, ... — **every** append is a resize.

Copies for n appends:

$$0 + 1 + 2 + \dots + (n - 1) = \frac{n(n - 1)}{2} = \Theta(n^2)$$

Every append is $O(n)$. 100,000 appends $\approx$ 5 **billion** copies.

The most economical in memory — never a wasted slot — and hopeless in time. The k -th append copies the $k - 1$ elements already there. It is the triangular sum from Lecture 02 again.

Option 2: a constant step — say, 100 at a time

Resize only every 100 appends — surely much better?

Copies for n appends: $100 + 200 + 300 + \dots \approx \frac{n^2}{2 \times 100}$

Still $\Theta(n^2)$ — just 100 times smaller. 100,000 appends $\approx$ 50 million copies.

This is the trap that looks like a fix. Growing by a constant c divides the number of resizes by c , but each resize still copies everything, and the amounts copied still grow linearly: $c, 2c, 3c, \dots$. A sum of n/c terms growing linearly is quadratic. The constant c only changes the constant in front of the n^2 — the part Big-O throws away.

Counted exactly for $n = 100,000$ appends starting from capacity 1: growing by 1 copies 4,999,950,000 elements; growing by 10 copies 499,960,000. The second is ten times smaller, and still hopeless.

Option 3: double

Capacity 1, 2, 4, 8, 16, ... — resize only when size reaches a power of two.

Copies for n appends:

$$1 + 2 + 4 + \dots + 2^k < 2n \quad (2^k < n)$$


$\Theta(n)$ **in total**. 100,000 appends: 131,071 copies.

Resizes happen at sizes 1, 2, 4, 8, ..., and each copies that many elements. The copies form a **geometric** series, and a geometric series is dominated by its last term: $1 + 2 + 4 + \dots + 2^k =$

$2^{k+1} - 1 < 2 \cdot 2^k$. The last resize copied fewer than n elements, so all the copying together is fewer than $2n$.

Counted exactly for $n = 100,000$: 131,071 copies — against 5 billion for growing by one. That is the whole of this lecture in two numbers.

The three side by side




Counted, not timed: the total number of writes and copies for n appends, for each growth rule, on a logarithmic axis. Growing by 1 and growing by 100 are the same shape — quadratic — shifted apart by a constant. Doubling is a different shape: it stays under the dotted line $3n$, for every n .

This is the plot your `resize_count` lets you draw in the notebook `notebooks/04-dynamic-array.ipynb`.

Amortised Analysis

One expensive append, many cheap ones



The red bars are resizes: expensive, but **exponentially rare**. The green line — the average so far — never rises above **3**.

Each bar is the cost of one append with doubling, counting one unit per write and one per copied element. Most appends cost 1. Append number 2, 3, 5, 9, 17, 33, 65, ... — one more than a power of two — pays for copying everything first.

The expensive appends are far apart, and they get further apart exactly as fast as they get more expensive. That balance is the whole trick.

Amortised cost: the definition

The **amortised cost** of an operation is the **total cost of any sequence** of n operations, **divided by n** .

For append with doubling: total $< 3n$, so the amortised cost is $< 3 = O(1)$.

- One individual append can still cost $O(n)$ — the **worst case** is unchanged.
- But **no sequence** of n appends can cost more than $3n$ — a **guarantee**.

Amortise comes from finance: paying off a large cost in small instalments over time. (In Arabic, الإطفاء or التكلفة المستهلكة.) The large cost here is a resize; the instalments are the cheap appends that come before it.

The proof — the aggregate method. Take any n appends starting from an empty array of capacity 1.

- Every append writes one element: n writes.
- Resizes happen when size reaches 1, 2, 4, ..., 2^k , with $2^k < n$, and copy that many elements: $1 + 2 + \dots + 2^k = 2^{k+1} - 1 < 2n$ copies.
- Total: fewer than $n + 2n = 3n$. Divided by n : fewer than 3 per append. ■

The same proof, as a bank account. Charge every append 3 coins. One pays for its own write. The other two are saved *on the element just written*. When the array doubles from capacity c to $2c$, the $c/2$ elements written since the last resize each hold 2 saved coins — c coins in all, exactly

enough to copy the c elements. The account never goes negative, so 3 coins per append always suffice.

Amortised is not “average case”

	Average case (Lecture 02)	Amortised
Averages over	possible inputs	the operations in one sequence
Needs	an assumption about the inputs	nothing
Guarantee	“usually”	always , for every sequence

append is **worst-case** $O(n)$, **amortised** $O(1)$. Both statements are true and both are useful.

Average-case analysis says: if the inputs are random in a certain way, the expected cost is this. It can be wrong for the inputs you actually get.

Amortised analysis makes no assumption at all. It says: whatever you do, any n appends in a row cost at most $3n$. There is no unlucky input that breaks it.

When would the worst case still matter? When *one* slow operation is a problem in itself — a game that must draw a frame every 16 milliseconds, a controller that must respond in real time. Then an occasional $O(n)$ pause is a bug even though the average is excellent, and real-time systems reserve capacity in advance. For almost everything else, amortised $O(1)$ is as good as $O(1)$.

Any constant factor works

Growth rule	Total copies for n appends	Amortised append	Worst unused space
+ 1	$\approx n^2/2$	$O(n)$	none
+ c (constant)	$\approx n^2/2c$	$O(n)$	c slots
$\times 2$	$< 2n$	$O(1)$	about half
$\times 1.5$	$< 3n$ (roughly)	$O(1)$	about a third

What makes it $O(1)$ is growing by a **factor** — a percentage of the current size — not the particular factor.

For any factor $r > 1$ the copies form a geometric series with ratio r , whose sum is at most about $n \cdot \frac{r}{r-1}$ — a constant times n . A smaller factor copies a little more often but wastes less memory; a larger one wastes more memory but copies less. That is a time–space trade-off with no single right answer, which is why real implementations differ: many use 2, some 1.5, and CPython uses about 1.125 (next slide).

(In the exercise the factor is `self.growth`, 2 by default. Array needs a whole number of slots, so a factor like 1.5 needs rounding up — `int(capacity * 1.5) + 1` — if you experiment with it.)

Shrinking

Give memory back — carefully

Pop until the array is nearly empty, and most of the Array is wasted. Shrink when size falls to **one quarter** of capacity — and then **halve**.

Why a quarter, not a half?

Start: capacity 8, size 8 (full)	halve at ½ full	halve at ¼ full
append → size 9	double to 16: copy 8	double to 16: copy 8
pop → size 8	halve to 8: copy 8	nothing ($8 > 16/4$)
append → size 9	double to 16: copy 8	nothing
pop → size 8	halve to 8: copy 8	nothing
repeat...	every operation $O(n)$	$O(1)$


Growing at “full” and shrinking at “half full” leaves no gap between the two thresholds. An append that just crossed the growth boundary is immediately followed by a pop that crosses the shrink boundary, and the array copies itself back and forth on every single operation. That is **thrashing**, and it turns an amortised $O(1)$ structure into an $O(n)$ one on a perfectly ordinary sequence.

Shrinking at a quarter leaves slack on both sides. After a halving, the array is half full, so it needs about $n/4$ more pops, or $n/4$ more appends, before it resizes again — and those operations have paid for the copy. Amortised $O(1)$ for any mix of appends and pops.

The `DynamicArray` exercise does **not** require shrinking; the tests never check capacity after pop. It is an extension for Homework 4 — and the question bank has a practice problem on it.

Python’s list

Measured: CPython does exactly this



Read straight from the interpreter with `sys.getsizeof`: capacity 4, 8, 16, 24, 32, 40, 52, 64, 76, 92, ... — grow by about **one eighth**, plus a little.

A Python list's memory is a fixed header plus one 8-byte reference per slot of capacity, so `sys.getsizeof` reveals the capacity:

```
import sys
empty = sys.getsizeof([])
values = []
for i in range(20):
    values.append(i)
    print(len(values), (sys.getsizeof(values) - empty) // 8)
```

Run it on your machine. On Python 3.13 the capacity steps through 4, 8, 16, 24, 32, 40, 52, 64, ... — slightly more than one eighth extra each time, rounded, with bigger jumps while the list is small. CPython chooses a small factor to waste little memory, because Python programs hold enormous numbers of lists. It is still a *factor*, so `append` is still amortised $O(1)$ — the same argument as this lecture, with a different constant.

That is Lecture 01's promise kept: the doubling that "happens in C, a few layers below your code", now visible, and now yours.

All the Operations

DynamicArray: the contract with costs

Operation	Cost	Why
<code>a[i], a[i] = x, len(a)</code>	$O(1)$	address arithmetic; a stored count
<code>append(x)</code>	amortised $O(1)$, worst $O(n)$	occasional resize, paid for by the rest
<code>pop()</code> from the end	$O(1)$	nothing moves
<code>insert_at(i, x)</code>	$O(n)$	shift <code>i .. size-1</code> right (and maybe resize)
<code>pop(i)</code> from the middle	$O(n)$	shift <code>i+1 .. size-1</code> left
<code>x in a</code> , search	$O(n)$	one by one
iterate	$O(n)$	

Exactly the costs of Python's `list` — because it **is** one.

Compare with Lab 03's table of list methods: every row matches. You now know *why* each one costs what it does, because you have written the code that pays it.

`insert_at` and `pop(i)` are Lecture 02's `insert_at` and `remove_at` from `dsa/array_ops.py`, moved inside the class — with one addition: `insert_at` must resize first if the array is full.

Remember the storage rule: `_block` is an `Array` and nothing else. No `list` anywhere in the class, not even as scratch space during a resize.

What it is built for

From here on, the dynamic array is **storage** for other structures:

Week	Structure	Uses the DynamicArray as
6	Stack	the items — top at the end , so push and pop are $O(1)$
7	SlowQueue	the items — pop(0) from the front, $O(n)$: measure it
12	MinHeap	the complete binary tree, stored level by level
14	Graph	each node's list of neighbours

A bug in your DynamicArray becomes a bug in all four. **Test it well now.**

This is the storage rule from Lecture 02 paying off. Because the stack, the slow queue, the heap and the graph are all built on the structure you write this week, their costs come straight from this lecture's table: a stack is $O(1)$ because append and pop at the end are; the slow queue is slow because pop(0) shifts; a heap can add a leaf in amortised $O(1)$ before sifting it up.

This Week

Exercises: dsa/dynamic_array.py

Method	Target	Watch out for
append(value)	amortised $O(1)$	resize before writing; grow by <code>self.growth</code>
<code>_resize(capacity)</code>	$O(n)$	copy only size elements; count <code>resize_count</code>
<code>__getitem__</code> , <code>__setitem__</code>	$O(1)$	bounds are 0 .. size-1; translate negatives
<code>insert_at(index, value)</code>	$O(n)$	resize if full; shift from the end
<code>pop(index=-1)</code>	$O(1)$ at the end, else $O(n)$	IndexError when empty; clear the freed slot

```
pytest tests/test_dynamic_array.py -v
```

`test_doubling_keeps_resizes_logarithmic`: 1,000 appends must resize fewer than 20 times.

The constructor, `__len__`, `__iter__`, `__repr__` and `make_block` are given. The resize test is the one that separates doubling from a constant step: with doubling, 1,000 appends from capacity 1 resize exactly 10 times ($2^{10} = 1024$); with "+ 1" they resize 999 times.

After this week, `pytest tests/test_stack_queue.py` will start to make sense: the stack and the slow queue already use your DynamicArray, so they depend on it being right.

Homework 4 — before Lecture 05

1. **Implement** `dsa/dynamic_array.py` until `pytest tests/test_dynamic_array.py` passes.
2. **Measure.** Make a copy of your class that grows by $+ 10$ instead of $\times$ growth. Time 1,000 ... 32,000 appends for both with `viz.complexity.measure`, and plot with `reference=["n", "n^2"]`. Plot `resize_count` against `n` for both as well.
3. **Prove.** With growth factor 3 (capacity 1, 3, 9, 27, ...), show that `n` appends cost $O(n)$ in total.
4. **Extend.** Add shrinking to your class: halve the capacity when `pop` leaves the array one quarter full. Write a test that alternates `append` and `pop` at the boundary, and check with `resize_count` that it does not thrash.

For item 2, the starting point is the `measure` example in the Lecture 02 handout; `make` should return an empty array and the function should `append` `n` items to it. Keep both classes in a scratch file or the notebook, not in `dsa/dynamic_array.py`, so the tests keep checking the real one.

For item 3, follow the aggregate proof: where do the resizes happen, how much does each one copy, and what does the geometric series sum to?

Summary

Seven things to keep

1. A dynamic array is an **Array plus a size**; `capacity` $\geq$ `size`.
2. Full? **Allocate** a bigger Array, **copy** everything, **switch** — $O(n)$.
3. Grow by a **constant** $\rightarrow \Theta(n^2)$ for `n` appends. Grow by a **factor** $\rightarrow \Theta(n)$.
4. **Amortised** $O(1)$: any `n` appends cost $< 3n$. The worst single append is still $O(n)$.
5. Amortised is a **guarantee over a sequence**, not an average over inputs.
6. Shrink at **one quarter** full, not one half — or it **thrashes**.
7. Python's `list` is exactly this, growing by about **1/8**.

Next

Week 5 — Linked lists. The opposite trade: no copying, ever, and $O(1)$ insertion at the front — paid for with $O(n)$ indexing.

Sources and further reading

- **T. H. Cormen, C. E. Leiserson, R. L. Rivest and C. Stein.** *Introduction to Algorithms*, 3rd ed., MIT Press, 2009, chapter 17, "Amortized Analysis" — the aggregate, accounting and potential methods, with dynamic tables (section 17.4) as the worked example, including the one-quarter shrinking rule.
- **M. T. Goodrich, R. Tamassia and M. H. Goldwasser.** *Data Structures and Algorithms in Python*, Wiley, 2013, section 5.3, "Dynamic Arrays and Amortization" — a dynamic array built on a `ctypes` array, the approach this course takes.

- **CPython source.** `Objects/listobject.c`, function `list_resize` — the over-allocation rule you measured with `sys.getsizeof`.
- **R. E. Tarjan.** "Amortized Computational Complexity", *SIAM Journal on Algebraic and Discrete Methods* 6(2), 1985 — the paper that named and formalised amortised analysis.

Every figure in this lecture is generated by `tools/figures_104.py`. The cost charts are counted; the CPython chart is read from your interpreter, so run the script to see your own Python's growth pattern.