

Linked Lists

DSA27 — Lecture 05 · Data Structures and Algorithms

Dr. Haitham A. El-Ghareeb

Week 5 · Fall 2026

Contents

Where We Are	2
Last week, one weakness	2
Today	2
Nodes and References	2
A node: a value and a link	2
Build one by hand	3
Two layouts, one difference	4
Walking the List	4
The traversal loop	4
Indexing is $O(n)$ — that is the lesson	5
Inserting	5
At the front: $O(1)$	5
In the middle: walk, then splice	6
At the end: $O(n)$ — unless you keep a tail	6
Removing	7
From the front: $O(1)$	7
By value: skip over it	7
Reversing in Place	8
Three references	8
Array Against Linked List	9
The costs, side by side	9
Measured, not asserted	9
Memory: a node is not free	10
Variants	10
Doubly linked, and sentinels	10
This Week	11
Exercises: dsa/linked_list.py	11
Homework 5 — before Lecture 06	11
Summary	12

Seven things to keep	12
Next	12
Sources and further reading	12

How to read this document. *This is the handout for Lecture 05. It holds everything on the slides, plus what I said out loud. The linked list is the exact mirror of the dynamic array: every operation that was cheap there is expensive here, and the other way round. It is also the first structure built from **nodes and references** rather than from an array — the building blocks of trees and graphs later.*

Slides: DSA27-L05-slides.pdf · Code: dsa/linked_list.py · Tests: tests/test_linked_list.py

Where We Are

Last week, one weakness

The dynamic array: $O(1)$ to index, amortised $O(1)$ to append —

and $O(n)$ to insert at the **front**, because **everything shifts**.

Today: a structure where nothing ever shifts, and nothing is ever copied.

Every array-based structure pays for contiguity. Keeping elements side by side is what makes `a[i]` a single address calculation — and it is also why inserting at position 0 moves every other element, and why growing means copying the whole array.

A linked list gives up contiguity entirely. Each element lives in its own small object, anywhere in memory, and knows only where the next one is. Insertion becomes a matter of changing two references. Indexing becomes a walk.

Linked list in Arabic: القائمة المتصلة.

Today

1. Nodes and references
2. Walking a list — and why indexing is $O(n)$
3. Inserting: at the front, in the middle, at the end
4. Removing — and who frees the node
5. Reversing in place, with three references
6. Array against linked list: costs, memory, measured
7. Variants: tail pointer, doubly linked, sentinel


Nodes and References

A node: a value and a link

```
class Node:
    __slots__ = ("value", "next")
```

```
def __init__(self, value, next=None):
    self.value = value
    self.next = next          # another Node, or None at the end
```

A linked list: nodes, each holding a value and a reference to the next



The list object keeps only **head** (the first node) and **_size**.

Node is given to you in `dsa/linked_list.py`, exactly as above. `__slots__` tells Python that a node has only these two attributes, which saves memory — it matters when a list has a million nodes.

The picture is the whole structure. There is no block, no capacity, no index arithmetic: just a chain of references ending in `None`. The `LinkedList` object holds the entry point, `head`, and a count, `_size`.

This is also the course storage rule from Lecture 02 at work: the linked list stores its data in **node objects** — one of the three things a structure is allowed to use. No `list` anywhere.

Build one by hand


```
>>> from dsa.linked_list import Node
>>> c = Node(1)
>>> b = Node(7, c)
>>> a = Node(3, b)
>>> a.value, a.next.value, a.next.next.value
(3, 7, 1)
>>> a.next.next.next is None
True
```

A reference **is** the arrow in the picture. `a.next` is `b`.

Do this at the REPL before you write any of the exercise. Everything that follows is only a matter of which reference you change, and in what order.

Lecture 01 said a Python variable is a name bound to an object, and Lab 01 made you draw the arrows for `b = a`. A linked list is that idea turned into a data structure: `node.next` is just another name bound to another object.

Two layouts, one difference



The array's slots are contiguous, so the address of item i is arithmetic. The nodes can be anywhere — they are created one at a time, wherever Python has room — so the only way to find item i is to start at head and follow i references.

The addresses in the figure are illustrative, but the point is real: this is why indexing a linked list is $O(n)$, and it is also why a linked list is slower in practice than its Big-O suggests. Modern processors fetch memory in blocks and are fast at walking contiguous data (the *cache*); jumping from node to node defeats that.

Walking the List

The traversal loop

```
node = self.head
while node is not None:
    # ... use node.value ...
    node = node.next
```

Every linked-list algorithm is this loop, plus a little bookkeeping.

- `node is not None` — stop **after** the last node
- `node = node.next` — the only way to move

Learn this loop until you write it without thinking. `__iter__` in the given code is exactly this loop with `yield node.value` in the middle. `find(value)` is this loop with an index counter and an early return. `__getitem__(index)` is this loop stopped after `index` steps.

Two classic mistakes:

- `while node.next is not None:` stops **on** the last node instead of after it — right when you need the last node (appending), wrong when you need to visit every node.
- Forgetting `node = node.next` — an infinite loop. Press Ctrl+C.

Watch it move. Set a breakpoint inside this loop in `dsa/linked_list.py`, press F5, choose `pytest: current test file`, and step with F10. In the Variables panel, node jumps from object to object. `viz.draw.draw_linked_list(list(ll), highlight=i)` draws the same thing.

Indexing is $O(n)$ — that is the lesson

	DynamicArray[i]	LinkedList[i]
How	base + i × slot size	follow next i times from head
Cost	$O(1)$	$O(i)$ — $O(n)$ worst case

Same interface, `a[i]`. **Completely different cost.**

```
for i in range(len(ll)):
    print(ll[i])
```

$O(n)$ each time: $O(n^2)$ for the whole loop!

Iterate with `for value in ll`: — one walk, $O(n)$ in total.


This is the most important slide for your own code. Python lets `LinkedList` support `ll[i]` through `__getitem__`, and it looks exactly like array indexing. But each `ll[i]` starts again from head. A loop that indexes 0, 1, 2, ... does $0 + 1 + 2 + \dots + (n - 1)$ steps — the triangular sum again, $O(n^2)$.

An interface does not tell you the cost. The ADT does — and for a linked list, positional access is $O(n)$ by contract.

`len(ll)` on the other hand is $O(1)$, because the list keeps `_size` up to date. Without it, `len` would be a walk too. That is a design decision with a price: every method that adds or removes a node must remember to update `_size`.

Inserting

At the front: $O(1)$



```


new = Node(value)
new.next = self.head    # 1. Link the new node to the old first node
self.head = new         # 2. then make it the first node
    
```

Two references change, whatever the length of the list. No shifting: compare Lecture 02's array insertion, which moved every element.

Order matters. Do step 2 first and `self.head` already points at the new node, so step 1 sets `new.next = new` — a node pointing at itself, and the rest of the list is lost (unreachable, then garbage-collected). With references, the rule is: **link the new node in before you let go of the old one.**

Both steps also fit in one line — `self.head = Node(value, self.head)` — because Python evaluates the right-hand side, including the old `self.head`, before assigning. Then `self._size += 1`.

In the middle: walk, then splice



To put a new node at index i , stop at the node **before** it (index $i - 1$), then:

```

new.next = prev.next
prev.next = new
    
```

A singly linked list can only be changed from the node **before** the change, because only that node holds the reference that must be redirected. So insertion at index i walks $i - 1$ steps to find `prev`: $O(n)$ in the worst case. The splice itself is two assignments, $O(1)$.

This is the honest comparison with an array: the array *finds* position i in $O(1)$ and pays $O(n)$ to *make room*; the list pays $O(n)$ to *find* position i and $O(1)$ to make room. If you already hold a reference to the right node — which happens often, for example while walking — the list's insertion is $O(1)$ and the array's is still $O(n)$.

`insert_at(0, value)` has no `prev`; it is `push_front`. Handle it first, then walk for index ≥ 1 . The exercise must accept `index == len(self)` (insert at the end) and raise `IndexError` otherwise out of range.

At the end: $O(n)$ — unless you keep a tail

`append` must find the last node: walk from `head` until `node.next is None` — $O(n)$.

Keep a second reference, `_tail`, to the last node, and `append` becomes $O(1)$:

```
self._tail.next = new
self._tail = new
```

The price: **every** method that can change the last node must keep `_tail` correct.

The exercise asks for `append` in $O(n)$, with the tail pointer as a challenge in the docstring. If you add it, audit every method: `push_front` on an empty list must set the tail too; `pop_front` of the only node must clear it; `remove` of the last node must move it back — and in a **singly** linked list, “back” means walking to find the new last node, $O(n)$. `reverse` swaps head and tail.

That is the general lesson of this lecture: every extra reference buys speed somewhere and costs care everywhere.

Removing

From the front: $O(1)$


```
if self.head is None:
    raise IndexError("pop from empty list")
value = self.head.value
self.head = self.head.next
self._size -= 1
return value
```

The old first node is no longer referenced by anything — so it is **gone**.

Who frees the node? In C, you would call `free(old)`, and forgetting is a memory leak while freeing too early is a crash. In Python, the garbage collector reclaims any object nothing refers to — Lecture 01’s memory models. You simply stop pointing at it.

That is convenient and it hides the structure, which is why the debugger is worth using this week: watch `self.head` change and the old node disappear from view.

By value: skip over it



Find the node **before** the one to remove, then `prev.next = prev.next.next`.

The **head** is the special case: it has no `prev`.

`remove(value)` removes the **first** node holding `value` and returns `True`, or returns `False` if there is none. The shape:


1. If the list is empty, return `False`.
2. If the head holds `value`, move head on, as in `pop_front`.
3. Otherwise walk with `prev`, checking `prev.next.value`, and when it matches, `prev.next = prev.next.next`.

Do not forget `_size`. And compare the cost with the array: finding the value is $O(n)$ in both; removing it is $O(1)$ here and $O(n)$ there, because the array must close the gap.

The special case for the head is annoying enough that there is a standard trick to remove it — the sentinel node, at the end of this lecture.

Reversing in Place

Three references



Keep **prev** (the reversed part) and **node** (the rest). One step moves one node across the line:

1. remember `node.next` — you are about to overwrite it
2. point `node.next` back at `prev`
3. advance: `prev` becomes `node`, `node` becomes the remembered next

At the end, `prev` is the new head. $O(n)$ time, $O(1)$ extra space.

This is the classic linked-list problem, and it appears in interviews for a reason: it tests whether you can hold an **invariant** in your head. Before every step, `prev` is the head of a correctly reversed list of the nodes already visited, and `node` is the head of the part still to do. The step moves exactly one node from the second to the first, and the invariant still holds.

Start: `prev = None` (nothing reversed yet), `node = self.head`. Stop when `node` is `None`. Then `self.head = prev`.

Step 1 is the one people forget. Once you have set `node.next = prev`, the only reference to the rest of the list is gone — unless you saved it first.

The test `test_reverse_relinks_rather_than_rebuilding` checks that you reversed the **same** node objects. Building a new list of new nodes is $O(n)$ extra space and fails the test — the point is to re-link, not to copy.

Draw three or four nodes on paper, and do the three moves by hand, before you write a line.

Array Against Linked List

The costs, side by side

Operation	DynamicArray	LinkedList
read / write position i	$O(1)$	$O(n)$
insert or remove at the front	$O(n)$	$O(1)$
append at the end	amortised $O(1)$	$O(n)$ — $O(1)$ with a tail
remove from the end	$O(1)$	$O(n)$ singly linked
insert / remove given the position's neighbour	$O(n)$	$O(1)$
search for a value	$O(n)$	$O(n)$
<code>len</code>	$O(1)$	$O(1)$ — because of <code>_size</code>
memory per element	one reference	a whole node

Read the table as a set of trade-offs, not a ranking. Neither structure is better. Each is $O(1)$ exactly where the other is $O(n)$ — which is why this course builds both, and why Week 6 (stacks) and Week 7 (queues) will ask you to choose between them.

“Remove from the end” is $O(n)$ for a singly linked list even with a tail pointer: after removing the last node, you need the new last node, and only the node before it knows where it is. A **doubly** linked list fixes that.

Measured, not asserted




Figure 1: Real timings, reference implementations on the course Array and on nodes.

Both panels are log–log. At $n = 32,768$:

- inserting at the front of an Array took about 23,000 μs (every element shifted); `push_front` took about 0.15 μs — independent of n ;
- reading the middle element of an Array took about 0.2 μs ; walking to the middle of a linked list took about 370 μs .

The flat lines are averaged over 1,000 operations each, because a single $O(1)$ operation is too fast for the clock — the same caution as in Lecture 02.

Memory: a node is not free

```
>>> import sys
>>> sys.getsizeof(Node(1))      # with __slots__, on 64-bit CPython 3.13
48
```

An array slot is **8 bytes** (one reference). A node is **48** — six times as much, per element, **plus** the value object both of them point to.

Every node is a full Python object with a header, a type pointer and two references. A dynamic array stores only the references, side by side, with some unused capacity. For a million small values, that difference is tens of megabytes.

Memory layout also matters for speed: an array's references are contiguous and read efficiently in blocks; nodes are scattered, and every `node.next` may be a jump to a different place in memory. That is why, in practice, arrays often beat linked lists even for operations where the Big-O is equal.

When is a linked list the right choice? When you insert and remove at the front a lot (a stack, Week 6); when you splice in and out at positions you already hold (an LRU cache, a text editor's buffer); when you must never pay for a big copy at an unlucky moment; and as the **chains** inside a hash table (Week 13), where each bucket is a short linked list of Entry nodes.

Variants

Doubly linked, and sentinels

Doubly linked: each node also holds `prev`.

- remove a node you hold in $O(1)$ — no need to find its predecessor
- walk both ways; remove from the end in $O(1)$
- costs one more reference per node, and twice the bookkeeping

Sentinel (dummy) node: a permanent, valueless node before the first real one.

- every real node now has a predecessor — **no special case for the head**

Circular: the last node points back to the first — useful for round-robin.

Python's `collections.deque` — the queue from Lab 03 — is a doubly linked list of fixed-size blocks: $O(1)$ at both ends like a linked list, with less memory overhead per element like an array. Real systems often combine the two ideas.

With a sentinel, remove has no special case:

```
prev = self.sentinel # sentinel.next is the first real node
while prev.next is not None:
    if prev.next.value == value:
        prev.next = prev.next.next
        self._size -= 1
        return True
    prev = prev.next
return False
```

The exercise does **not** use a sentinel — `self.head` is the first real node, as the tests check — so you will write the head case yourself. The trick is worth knowing for the exam and for later structures.

This Week

Exercises: `dsa/linked_list.py`

Method	Target	The trap
<code>push_front(value)</code>	$O(1)$	link first, then move head
<code>append(value)</code>	$O(n)$ (tail: $O(1)$ challenge)	the empty list has no last node
<code>insert_at(index, value)</code>	$O(n)$	index 0 is <code>push_front</code> ; stop at index - 1
<code>pop_front()</code>	$O(1)$	<code>IndexError</code> when empty
<code>remove(value)</code>	$O(n)$	the head has no prev; return True/False
<code>find(value)</code>	$O(n)$	-1 when absent
<code>reverse()</code>	$O(n)$, $O(1)$ space	save next before you overwrite it
<code>__getitem__(index)</code>	$O(n)$	<code>IndexError</code> out of range

Every method that adds or removes a node: **update _size**.

```
pytest tests/test_linked_list.py -v
```

Node, the constructor, `__len__`, `__iter__` and `__repr__` are given. Draw before you code: for each method, sketch the nodes and references before and after, and write the reference assignments in the order that never loses the rest of the list.

Test the edge cases by hand: the empty list, a one-node list, the first node, the last node. Almost every linked-list bug lives in one of those four.

Homework 5 — before Lecture 06

1. **Implement** `dsa/linked_list.py` until `pytest tests/test_linked_list.py` passes.

2. **Tail pointer.** Add `_tail` and make `append` $O(1)$. List every method you had to change, and why.
3. **Measure.** Time `push_front` against `DynamicArray.insert_at(0, x)`, and `ll[n // 2]` against `arr[n // 2]`, for $n = 1,000 \dots 32,000$. Do your curves match the figure in this handout?
4. **On paper.** Trace `reverse()` on `[1, 2, 3, 4]`: draw `prev`, `node` and all the links after each step.

For item 3, remember to average fast operations over many repetitions, as the figure's legend says — a single $O(1)$ call is below the clock's resolution.

Summary

Seven things to keep

1. A linked list is **nodes** — a value and a next reference — plus a head.
2. Every algorithm is the **traversal loop**: `while node is not None: ... node = node.next.`
3. **Indexing is $O(n)$** ; iterate instead of indexing in a loop.
4. Front insert and remove: $O(1)$. Link the new node in **before** moving head.
5. Changes happen **from the node before**; the head is the special case (or use a sentinel).
6. Reverse with **three references**; save next first.
7. Array and linked list are **mirror images**: each is $O(1)$ where the other is $O(n)$ — choose by the operations you need.

Next

Week 6 — Stacks. One ADT, and the question this lecture prepares: array or linked list underneath? Then the applications that make stacks essential.

Sources and further reading

- **M. T. Goodrich, R. Tamassia and M. H. Goldwasser.** *Data Structures and Algorithms in Python*, Wiley, 2013, chapter 7, "Linked Lists" — singly, circular and doubly linked lists, with sentinels, and the comparison with array-based sequences.
- **T. H. Cormen, C. E. Leiserson, R. L. Rivest and C. Stein.** *Introduction to Algorithms*, 3rd ed., MIT Press, 2009, section 10.2, "Linked lists" — including sentinels.
- **B. W. Kernighan and D. M. Ritchie.** *The C Programming Language*, 2nd ed., Prentice Hall, 1988, chapter 6, "Structures" — self-referential structures: the same node, with `malloc` and `free` visible.
- **CPython source.** `Modules/_collectionsmodule.c` — deque as a doubly linked list of fixed-size blocks.

Every figure in this lecture is generated by `tools/figures_l05.py`. The timing figure is real measurement and will differ slightly on your machine.