

Stacks

DSA27 — Lecture 06 · Data Structures and Algorithms

Dr. Haitham A. El-Ghareeb

Week 6 · Fall 2026

Contents

Where We Are	2
Two structures, one question	2
Today	2
The Stack ADT	2
Last in, first out	2
Two Honest Implementations	3
The top goes where the $O(1)$ end is	3
...and one dishonest one	4
Application 1: Balanced Brackets	4
Every closer must match the most recent opener	5
Application 2: Postfix	5
Expressions with no parentheses	5
Evaluating postfix with one stack	6
Application 3: Infix to Postfix	6
The shunting-yard algorithm	6
Trace: $3 + 4 * 2$	7
Trace: $(3 + 4) * 2$	7
Application 4: The Call Stack	8
You have been using a stack since Week 3	8
This Week	9
Exercises: <code>dsa/stack.py</code>	9
Homework 6 — before Lecture 07	9
Summary	9
Seven things to keep	9
Next	10
Sources and further reading	10

How to read this document. This is the handout for Lecture 06. It holds everything on the slides, plus what I said out loud. The stack is the simplest ADT in the course, and one of the

most useful: the structure is almost trivial, and the interesting part is what you do with it. It is also the first time you choose between the two structures of Weeks 4 and 5.

Slides: `DSA27-L06-slides.pdf` · Code: `dsa/stack.py` · Tests: `tests/test_stack_queue.py` (the stack half)

Where We Are

Two structures, one question

Week 4: the **dynamic array** — cheap at the **end**. Week 5: the **linked list** — cheap at the **front**.

Today: an ADT that only ever touches **one end** — so either structure can keep its contract in $O(1)$. Which, and how?

Lecture 01 separated the **ADT** — the contract — from the **data structure** that keeps it, and promised that choosing between structures is the engineering. This is the first week you do that choice for real, with two structures you have built yourself.

Stack in Arabic: المكسدس.


Today

1. The Stack ADT: push, pop, peek — LIFO
2. Two honest implementations, and one dishonest one
3. Balanced brackets
4. Postfix: expressions without parentheses
5. Infix to postfix: the shunting-yard algorithm
6. The stack you have used all term: the call stack

The Stack ADT

Last in, first out

A stack: push, pop and peek all act on the top — and only the top



Operation	Meaning	Cost
<code>push(x)</code>	put x on top	$O(1)$
<code>pop()</code>	remove and return the top — IndexError if empty	$O(1)$

Operation	Meaning	Cost
<code>peek()</code>	return the top without removing it — <code>IndexError</code> if empty	$O(1)$
<code>is_empty()</code> , <code>len(s)</code>		$O(1)$

That table **is** the ADT: operations, their meaning, their cost — and nothing about storage. Everything a stack can do happens at one end, the **top**. There is no indexing, no search, no access to the middle, by design: a restriction that turns out to be exactly what a surprising number of problems need.


Everyday stacks: a pile of plates; the *Undo* history of an editor (the last edit is the first undone); the *Back* button of a browser; and, most important for us, the call stack of Lecture 03.

Empty stacks. Popping or peeking an empty stack is an error, and the contract says which one: `IndexError`, like `list.pop()` on an empty list. The tests check it. Returning `None` instead would be worse than an error: a `None` can travel a long way through a program before anyone notices.

Two Honest Implementations

The top goes where the $O(1)$ end is

One ADT, two honest implementations — the top goes where the $O(1)$ end is



- On a **dynamic array**: `top` = the **end**. `append` and `pop()` — nothing moves.
- On a **linked list**: `top` = the **head**. `push_front` and `pop_front`.

Both are honest stacks: both keep every operation $O(1)$. They differ in the details that Weeks 4 and 5 taught you to see:


	Array, top at end	Linked list, top at head
push	amortised $O(1)$ — the occasional resize copies	$O(1)$ worst case
pop, peek	$O(1)$	$O(1)$
memory per item	one reference (+ spare capacity)	a whole node (48 bytes)

	Array, top at end	Linked list, top at head
memory locality	contiguous	scattered

`dsa/stack.py` uses the first: `self._items` is your own `DynamicArray` from Week 4. That is the storage rule from Lecture 02 at work — a structure built on a structure you built — and it means a bug in last week's `DynamicArray` shows up in this week's `Stack`.

...and one dishonest one

Put the top at **index 0** of an array, and **every** push and pop shifts the whole stack: $O(n)$ each.



Real timings of n pushes followed by n pops, on reference implementations. The two honest designs grow linearly — one unit of slope on the log–log plot, $O(n)$ for n operations, $O(1)$ each. The top-at-index-0 design has twice the slope: $O(n^2)$ in total. At $n = 8,192$ that is roughly 50 seconds against a few milliseconds.


The two honest lines are not on top of each other: here the linked version is faster, because every access to the course `Array` goes through a Python-level bounds check (`dsa/array.py`), while a node's `next` is a plain attribute. That is a constant factor — the thing Big-O ignores and measurement shows. Python's own `list`, whose checks run in C, would beat both.

Same ADT, same interface, same tests pass — and one design is quadratic. The tests check *what* your stack does. Only you can check *how*.

Application 1: Balanced Brackets

Every closer must match the most recent opener

`is_balanced("({})")`: push each opener; each closer must match the top



- **Opener** ([{ : push it.
- **Closer**)] } : the stack must be **non-empty**, and the popped opener must be its **partner**.
- **End of text**: the stack must be **empty**.
- Anything else: ignore.

Why a stack? Because the rule is “the most recently opened bracket must close first” — last in, first out. `([ ])` fails exactly because `)` arrives while the most recent opener is `[`.

The three ways to be unbalanced correspond to the three checks:

Input	What happens	Which check
<code>([])</code>	<code>)</code> pops <code>[</code> — not its partner	the partner check
<code>)</code>	a closer arrives with the stack empty	the non-empty check
<code>(</code>	the text ends with <code>(</code> still on the stack	the empty-at-the-end check

Forgetting any one of them passes most tests and fails one. That is precisely what `tests/test_stack_queue.py::test_is_balanced` is built to catch — read its cases before you write the function.

The cost: every character is looked at once, and each push and pop is $O(1)$: $O(n)$ time, and $O(n)$ space in the worst case `(((((...)))))`.

This is how every compiler and editor checks brackets, and how an HTML or XML parser checks that tags nest.

Application 2: Postfix

Expressions with no parentheses

Infix (usual)	Postfix (reverse Polish)
<code>3 + 4</code>	<code>3 4 +</code>

Infix (usual)	Postfix (reverse Polish)
3 + 4 * 2	3 4 2 * +
(3 + 4) * 2	3 4 + 2 *
8 - 3 - 2	8 3 - 2 -

In **postfix**, each operator comes **after** its operands. No parentheses, no precedence rules — the order says everything.

Postfix notation was introduced by the Australian philosopher and computer scientist Charles Hamblin in the 1950s, building on the prefix (“Polish”) notation of the logician Jan Łukasiewicz — hence *reverse Polish notation*, RPN. Hewlett-Packard calculators used it for decades, and the Java and Python virtual machines are essentially postfix machines: they run a stack of operands.

Compare the second and third rows. In infix, the parentheses change the meaning. In postfix, the *order* of the tokens carries that meaning instead, so no parentheses are needed at all.

Evaluating postfix with one stack

- **Number:** push it.
- **Operator:** pop **two** values, apply, push the result.
- **End:** exactly one value is left — the answer.

Token	Action	Stack (top on the right)
3	push	3
4	push	3 4
2	push	3 4 2
*	pop 2 and 4, push 8	3 8
+	pop 8 and 3, push 11	11

The operand order trap. For 8 3 -, the first value popped is 3 and the second is 8, and the answer is 8 - 3 — **second popped minus first popped**. Get it the wrong way round and addition and multiplication still work (they commute), while subtraction and division silently give wrong answers. A test with - or / is the only thing that catches it.

Malformed input. If an operator finds fewer than two values, or more than one value is left at the end, the input was not valid postfix. Raise `ValueError`.

This is `evaluate_postfix` in `dsa/translation.py` — an exercise for Week 15, where the whole calculator comes together. You can write it now: it needs only this week’s stack. $O(n)$ time.

Application 3: Infix to Postfix

The shunting-yard algorithm

Read the infix tokens left to right, with an **output list** and an **operator stack**:

1. **Number** → output.
2. **Operator** → first pop to the output every operator on the stack with **higher or equal precedence**; then push this one.
3. **(** → push.
4. **)** → pop to the output until the **(**; discard the **(**.
5. **End** → pop everything left to the output.

Precedence: * / above + -.

Edsger Dijkstra described this algorithm in 1961, and named it after a railway shunting yard: numbers go straight through the junction, while operators wait on a siding — the stack — until it is their turn.

The idea: an operator must wait until its right operand has been output. It can leave the siding only when a *later* operator of lower or equal precedence arrives (which means its operand is complete), or a **)** closes its group, or the input ends.

Equal precedence pops too — that is what makes **-** and **/** **left associative**: $8 - 3 - 2$ means $(8 - 3) - 2$. When the second **-** arrives, the first **-** (equal precedence) is popped to the output first, giving $8\ 3\ -\ 2\ -$. (An operator like ******, which is right associative, pops only on strictly higher precedence. The exercise does not include it.)

Trace: $3 + 4 * 2$

Token	Output	Stack
3	3	
+	3	+
4	3 4	+
*	3 4	+ * ← + is lower, so it stays
2	3 4 2	+ *
end	3 4 2 * +	

Trace: $(3 + 4) * 2$


Token	Output	Stack
(		(
3, +, 4	3 4	(+
)	3 4 +	← pop to (
*, 2	3 4 + 2	*
end	3 4 + 2 *	

These are exactly the two cases in `tests/test_stack_queue.py::test_infix_to_postfix`. Trace $8 - 3 - 2$ and $2 * 3 + 4$ yourself before writing the function: one checks the “equal precedence pops” rule, the other checks that a lower-precedence operator empties the higher ones first.

The whole calculator. Infix → (shunting-yard) → postfix → (evaluation) → a value: two passes, two stacks, $O(n)$, no tree and no recursion. Week 15 builds the other route — a recursive parser and an expression tree — and compares them. Declared in the 2013/2014 bylaws as “the principles of language translation”, and this is where it starts.

Application 4: The Call Stack

You have been using a stack since Week 3



Each call **pushes** a frame; each return **pops** one. Recursion **is** a stack.

Lecture 03’s figure, seen again: the call stack is a stack of frames, managed by Python for you. That is why any recursive algorithm can be rewritten with a loop and an explicit stack — the stack you now know how to build. Week 14 does exactly that for depth-first search: `dfs_recursive`, and `dfs_iterative` with your own `Stack`, visiting the nodes in the same order.

An example you can write now: summing a nested list without recursion.

```
from dsa.stack import Stack

def nested_sum(nested):
    total, pending = 0, Stack([nested])
    while not pending.is_empty():
        item = pending.pop()
        if isinstance(item, list):
            for inner in item:
                pending.push(inner)          # "call" on every inner item
        else:
            total += item
    return total
```

Every push is a recursive call that has not happened yet; every pop is one being made. And unlike the call stack, this one has no 1,000-frame limit.

Undo and redo use two stacks: doing an action pushes it on the undo stack; *Undo* pops it and pushes it on the redo stack; *Redo* moves it back; any new action clears the redo stack.

This Week

Exercises: dsa/stack.py

Function	Target	The trap
<code>Stack.push(value)</code>	$O(1)$ amortised	top at the end of <code>self._items</code>
<code>Stack.pop()</code>	$O(1)$	<code>IndexError</code> when empty
<code>Stack.peak()</code>	$O(1)$	<code>IndexError</code> when empty; do not remove
<code>is_balanced(text)</code>	$O(n)$	three checks: partner, non-empty, empty at the end
<code>infix_to_postfix(tokens)</code>	$O(n)$	pop on higher or equal precedence; (and)

```
pytest tests/test_stack_queue.py -v -k "stack or balanced or postfix"
```

Your Stack stands on your DynamicArray: `tests/test_dynamic_array.py` must pass first.

`__init__`, `is_empty`, `__len__` and `__repr__` are given. `is_balanced` and `infix_to_postfix` should use your Stack — not a Python list as a stack, which is what the storage rule forbids for working storage when a structure of yours is available. Their **result** — the list of output tokens from `infix_to_postfix` — is a Python list, because that is the interface the tests use.

Both algorithms are $O(n)$: each token is pushed at most once and popped at most once.

Homework 6 — before Lecture 07

1. **Implement** `dsa/stack.py` until the stack tests pass.
2. **Trace** `infix_to_postfix` on `8 - 3 - 2, 2 * 3 + 4` and `( 1 + 2 ) * ( 3 - 4 ) / 5`, as tables. Then evaluate each result by hand with the postfix algorithm.
3. **Write** `evaluate_postfix` in `dsa/translation.py` (it only needs your Stack), and check it against your traces.
4. **Design**. Sketch an undo/redo system for a text editor with two stacks. What happens to the redo stack when the user types after undoing?

For item 3, `pytest tests/test_translation.py -v -k postfix` runs just the postfix tests; the rest of that file belongs to Week 15.

Summary

Seven things to keep

1. A stack is **LIFO**: push, pop and peek act on the **top** only, all $O(1)$.
2. Empty pop and peek raise **IndexError**.
3. On an array the top is the **end**; on a linked list the **head**. Both honest.

4. Top at index 0 of an array: $O(n)$ per operation — same tests pass, wrong design.
5. **Balanced brackets**: push openers; each closer must match the top; empty at the end.
6. **Postfix**: push numbers; an operator pops two — **second popped is the left operand**.
7. **Shunting-yard**: operators wait on a stack and leave on **higher or equal** precedence, or at).

Next

Week 7 — Queues. First in, first out: the ADT that needs **both** ends — and the week your `DynamicArray` finally meets an operation it is bad at.

Sources and further reading

- **E. W. Dijkstra.** “Algol 60 translation: An Algol 60 translator for the X1 and making a translator for Algol 60”, Mathematisch Centrum report MR 35, 1961 — where the shunting-yard algorithm was first described.
- **C. L. Hamblin.** “Translation to and from Polish notation”, *The Computer Journal* 5(3), 1962 — reverse Polish notation and its evaluation with a stack.
- **M. T. Goodrich, R. Tamassia and M. H. Goldwasser.** *Data Structures and Algorithms in Python*, Wiley, 2013, chapter 6, “Stacks, Queues, and Deques” — array-based stacks, and bracket and HTML tag matching.
- **T. H. Cormen, C. E. Leiserson, R. L. Rivest and C. Stein.** *Introduction to Algorithms*, 3rd ed., MIT Press, 2009, section 10.1, “Stacks and queues”.

Every figure in this lecture is generated by `tools/figures_106.py`; the call stack figure is Lecture 03’s. The timing figure is real measurement and will differ slightly on your machine.