

Queues

DSA27 — Lecture 07 · Data Structures and Algorithms

Dr. Haitham A. El-Ghareeb

Week 7 · Fall 2026

Contents

Where We Are	2
One end was easy. Now we need both.	2
Today	2
The Queue ADT	2
First in, first out	3
The Obvious Design	3
SlowQueue: a DynamicArray, as it is	4
Idea 1: don't shift — move the front	4
The Ring Buffer	5
Idea 2: wrap around	5
The two operations	5
Trace: capacity 4	6
Full or empty?	7
When the ring is full	7
Growing: the trap	8
Two More Honest Queues	8
The linked queue	8
A queue from two stacks	9
Measured	9
The gap is the lesson	10
Four queues, compared	10
Where Queues Are Used	11
FIFO is fairness — and breadth	11
The deque: both ends, both ways	11
This Week	12
Exercises: dsa/queue.py	12
Homework 7 — before Lecture 08	12
Summary	13
Seven things to keep	13
Next	13

Sources and further reading	13
---------------------------------------	----

How to read this document. *This is the handout for Lecture 07. It holds everything on the slides, plus what I said out loud. A queue looks like a stack with one word changed, and that one word, first instead of last, is enough to break the obvious design. This week you measure that break yourself, then fix it with the ring buffer, the most useful trick in the course that fits in one line of arithmetic.*

Slides: `DSA27-L07-slides.pdf` · Code: `dsa/queue.py` · Tests: `tests/test_stack_queue.py` (the queue half)

Where We Are

One end was easy. Now we need both.

Week 6: the **stack** touches **one** end. Put the top at the cheap end, done.

Today: the **queue** adds at one end and removes at the **other**.

- A dynamic array is cheap at the **end** only.
- A singly linked list is cheap at the **head** (and, with a tail, for adding at the end).

So which end goes where, and what happens to the expensive end?

Lecture 06 closed with a promise: this is the week your `DynamicArray` finally meets an operation it is bad at. A stack could hide the array's weak end simply by never using it. A queue cannot: whatever you do, one of its two operations lands on the front of the array.

Queue in Arabic: الطابور.


Today

1. The Queue ADT: enqueue, dequeue — FIFO
2. The obvious design, `SlowQueue`, and its $O(n)$ dequeue
3. The ring buffer: `CircularQueue`, every operation $O(1)$
4. Full or empty? — and growing a ring without scrambling it
5. The linked queue, and a queue made of two stacks
6. Measured: the gap is the lesson
7. Where queues are used — and the deque

The Queue ADT

First in, first out

A queue: in at the back, out at the front — first in, first out (FIFO)



Operation	Meaning	Cost
<code>enqueue(x)</code>	add x at the back	$O(1)$
<code>dequeue()</code>	remove and return the front (<code>IndexError</code> if empty)	$O(1)$
<code>is_empty()</code> , <code>len(q)</code>		$O(1)$

That table is the ADT: operations, meaning, cost, and nothing about storage — exactly as with the stack. Many textbooks add a `peek` (also called `front` or `first`) that returns the front without removing it; the exercise does not need it, but it is one line once the rest works.

Everyday queues: a line at a bank or a bakery, where the first to arrive is the first served; a printer's job list; the messages waiting for a WhatsApp server to deliver them; the keys you press while a program is busy, which arrive in the order you pressed them. The rule is **fairness in order of arrival** — and that is exactly what a stack does not give you (W6-M20).

Empty queues. Dequeuing an empty queue raises `IndexError`, like popping an empty stack. `tests/test_stack_queue.py::test_dequeue_on_empty_raises` checks it for both classes.


Stack against queue, in one line each:

	Adds at	Removes at	Order out
Stack	the top	the same end	reversed (LIFO)
Queue	the back	the other end	preserved (FIFO)

The Obvious Design

SlowQueue: a DynamicArray, as it is

SlowQueue: enqueue is append ($O(1)$); dequeue is pop(0) — $O(n)$



- enqueue = append at the end: amortised $O(1)$.
- dequeue = pop(0): every remaining element shifts left — $O(n)$.


This is SlowQueue in `dsa/queue.py`, and it is a correct queue: it passes every test. The class is deliberately named for its flaw, because you are going to measure it.

Could we swap the ends? Put the back at index 0 and the front at the end: then dequeue is pop() at the end, $O(1)$, but enqueue becomes insert_at(0, x), which shifts everything — $O(n)$ again. With a plain array, **whichever** end is the front, one of the two operations costs $O(n)$. That is the difference from the stack, which could simply put all its work at the cheap end.

Python's own list has the same problem: `list.pop(0)` is $O(n)$. That is why Lab 03 told you to use `collections.deque` for a queue, and promised you would measure the difference this week.

Idea 1: don't shift — move the front

Idea 1 — don't shift, move a head index: $O(1)$, but the queue crawls to the right



Keep a head index. Dequeue reads `block[head]` and adds 1 to head: $O(1)$.


But the used part crawls right, and the slots it leaves behind are never used again.

Dequeue is now $O(1)$ — nothing moves. The price is that dequeued slots at the front are dead. In the figure, eight enqueues and two dequeues leave only six items in an eight-slot block, yet there is no room for a ninth at the end. A long-running queue with this design grows its block for ever, even if it never holds more than three items at a time — a memory leak by design.

You could compact now and then — shift everything back to index 0 when the dead space gets large — and that is a legitimate design with an amortised argument like Lecture 04's. But there is a neater answer that never moves anything.

The Ring Buffer

Idea 2: wrap around



After the last slot comes slot 0 again:

- **front** at head; **back** at $(\text{head} + \text{size}) \% \text{capacity}$
- every step forward is $(i + 1) \% \text{capacity}$

The block is still a plain, fixed Array. What changes is how we **read** it: index arithmetic *modulo the capacity* makes the end of the block join up with the beginning, so the slots freed at the front are reused as soon as the back reaches them. This is called a **ring buffer** or **circular buffer**, and it is everywhere in systems programming: keyboard buffers, network cards, audio streaming, logging the last N events. Wherever data arrives at one rate and leaves at another, there is probably a ring buffer between them.

`dsa/queue.py` stores three things besides the block: `_capacity`, `_head` (the index of the front item) and `_size` (how many items are in the queue). It **does not** store the tail; the tail is computed:

$$\text{tail} = (\text{head} + \text{size}) \bmod \text{capacity}$$

In the figure, `head = 5` and `size = 5` in a block of 8, so the tail — the slot the next enqueue writes to — is $(5 + 5) \bmod 8 = 2$. The queue, front to back, is C D E F G: slots 5, 6, 7, then round to 0 and 1. `__iter__` in `dsa/queue.py` is given to you and walks exactly this way: $(\text{head} + \text{offset}) \% \text{capacity}$ for each offset from 0 to `size - 1`.

The two operations

enqueue(x)

1. full? — see the next slides

2. write `x` at `(head + size) % capacity`
3. `size += 1`

dequeue()

1. empty? — `IndexError`
2. read the value at `head`; clear that slot
3. `head = (head + 1) % capacity; size -= 1`

Nothing moves. Every step is a little arithmetic: $O(1)$.

Those six steps are the whole exercise, apart from the challenge; turning them into Python is your job. Two details are worth saying out loud.

Why clear the slot? Setting the dequeued slot back to `None` is not needed for the queue to be correct: that slot is outside the live part of the ring and will be overwritten later. But until then the block would still hold a reference to the dequeued object, and Python's garbage collector could not free it. Clearing it is the same care as `DynamicArray.pop` setting the vacated slot to `None`.

The modulo is the whole trick. Without it, `head + 1` would walk off the end of the `Array` and raise `IndexError`. `test_circular_queue_wraps_around` fills a ring of capacity 3, dequeues once and enqueues again: the new value must land in the slot at index 0, which only the `%` makes possible.


Trace: capacity 4

Operation	Block	head	size	Returns
start	_ _ _ _	0	0	
enqueue A, B, C	A B C _	0	3	
dequeue	_ B C _	1	2	A
dequeue	_ _ C _	2	1	B
enqueue D	_ _ C D	2	2	
enqueue E	E _ C D	2	3	← wrapped to slot 0
enqueue F	E F C D	2	4	← full
dequeue	E F _ D	3	3	C

Read the block as a ring from `head`: after "enqueue F" it is C D E F — slots 2, 3, 0, 1. E and F were written into the slots that A and B vacated; no element ever moved. This trace was produced by running a reference implementation, and every row can be checked with `tail = (head + size) % 4`: before "enqueue E", `tail = (2 + 2) % 4 = 0`.

Full or empty?

Full or empty? The one real trap in a ring buffer



When `head == tail`, the ring is **either** empty **or** full. Head and tail alone cannot tell which.

Look at the start of the trace (head 0, size 0, tail 0) and after “enqueue F” (head 2, size 4, tail $(2 + 4) \% 4 = 2$). In both, the next write would go to the slot head points at. A ring that stores only head and tail cannot tell these two situations apart. There are three standard fixes:

1. **Keep a count** — `_size`. Empty is `size == 0`, full is `size == capacity`. This is what `dsa/queue.py` does, and why `is_empty` and `is_full` are given to you as one-liners.
2. **Leave one slot unused**. Full is defined as `(tail + 1) % capacity == head`, so `head == tail` can only mean empty. Costs one slot; common in hardware and in C, where a count shared between two threads is awkward.
3. **A flag** set on the enqueue that makes `head == tail`, cleared on the next dequeue.

Every one of them works. The bug is to have none, which is a classic exam question (W7-B3).

When the ring is full


Two honest choices:

- **Refuse**: raise an error. A *bounded* queue — fixed memory, like a real buffer.
- **Grow** (the challenge): allocate a bigger Array and copy — as `DynamicArray` does, doubling so that enqueue stays amortised $O(1)$.

The docstring of `CircularQueue.enqueue` calls growing a **challenge**. The tests never overfill a ring, so raising `IndexError("queue is full")` is enough to pass them. But you should grow: Week 11’s `level_order` and Week 14’s `bfs` use your `CircularQueue`, and neither knows in advance how many items it will hold.

Both choices are used in practice. An audio driver’s buffer is bounded on purpose: if the consumer falls behind, old data is dropped or the producer waits, but memory never grows. A general-purpose queue, like Python’s `deque`, grows.

Growing: the trap



Copying slot by slot keeps the wrapped order — so the queue reads **wrong**. Copy in **queue order**, and restart head at 0.


This is the trap the docstring warns about (“a naive copy scrambles the order once the data has wrapped around”). In the figure, the full ring of capacity 4 holds C D E F starting at slot 2. Copying `new[i] = old[i]` puts E and F at slots 0 and 1 of the new block and C D at 2 and 3; with head still 2 and a capacity of 8, the queue now reads C, D, then two empty slots — E and F are lost behind the front. If the ring has not wrapped yet (head = 0), the naive copy happens to work, which is exactly why this bug survives casual testing.

The fix **unrolls** the ring while copying: the i -th item of the queue, which lives at $(\text{head} + i) \% \text{old_capacity}$, goes to slot i of the new block, for i from 0 to $\text{size} - 1$. Then head becomes 0, and capacity the new capacity (the last is easy to forget, and then every $\%$ still wraps at the old size). The front goes to slot 0, and the rest follow in order. Doubling the capacity each time gives the same amortised $O(1)$ enqueue as Lecture 04: over n enqueues, the copies total less than $2n$.

Two More Honest Queues

The linked queue

A linked queue: front at the head, back at the tail — never the other way round



Front at the **head**, back at the **tail**: both operations $O(1)$ worst case.

A singly linked list with a **tail** reference can add at either end in $O(1)$, but can remove only at the head in $O(1)$: removing the last node needs the node **before** it, which takes a walk (Lecture

05, W5-K3). So the queue's front must be the head (where we remove) and its back the tail (where we add). The other way round, every dequeue is $O(n)$.

- **enqueue:** make a node. If the queue is empty, it becomes both head and tail; otherwise link it after tail, and it becomes the new tail.
- **dequeue:** IndexError if empty; otherwise take the head's value and move head to the next node.

The one trap is the end of dequeue: when the queue becomes empty, tail must be reset to None too, or it keeps pointing at a node that is no longer in the queue and the next enqueue links onto it (W7-B4).

A queue from two stacks

- **inbox:** enqueue pushes here.
- **outbox:** dequeue pops here. When outbox is **empty**, first pour **all** of inbox into it.

Pouring reverses the order — and LIFO reversed is FIFO.


One dequeue may pour n items, so a single operation can cost $O(n)$. But each item is pushed onto inbox once, moved to outbox once, and popped from outbox once: three $O(1)$ steps over its whole life. So n operations cost $O(n)$ in total, and each is **amortised** $O(1)$ — the aggregate argument of Lecture 04 again, and of next_greater in W6-C3.

The rule “pour only when outbox is empty” is essential. Pouring on every dequeue is still correct, but each dequeue then moves everything twice, $O(n)$ each (W7-K3). It is W7-C1 in the question bank.

Why would anyone build a queue this way? In languages where the standard structures are **immutable**, as in Haskell or Erlang, a stack is a linked list that can be shared safely, and two of them make the standard functional queue. It is also a favourite interview question, and a good test of whether you understand amortised cost.

Measured

The gap is the lesson



Real timings of `dsa/queue.py` — the instructor’s reference solution, built on the course Array exactly as yours will be — against a simple linked queue. The ring is created with room for all n items, so growing plays no part. Left: n enqueues followed by n dequeues. The two honest queues rise with slope 1 on the log–log plot, $O(n)$ for n operations; the `SlowQueue` rises with slope 2, $O(n^2)$. At $n = 4,000$ that was about 6.5 seconds against 9 milliseconds for the ring: some 700 times slower, and the gap doubles every time n doubles.

Right: the cost of **one** dequeue (paired with one enqueue, to keep the size at n) on a queue already holding n items. The ring and the linked queue are flat lines: $O(1)$, however long the queue. The `SlowQueue` line is a straight line of slope 1: $O(n)$ per dequeue, about 15 milliseconds at $n = 16,000$ against roughly 1.4 microseconds for the ring. (The $O(1)$ lines are averages over 2,000 operations, because one operation is too quick to time on its own.)

Constant factors. The linked queue beats the ring by four to five times here, because every ring operation does a % and goes through the Python-level bounds check of `dsa/array.py`, while a node’s `next` is a plain attribute. The same effect as in Lecture 06. Big-O says both are $O(1)$; measurement says how big the 1 is.

This is the demonstration the course plan promised in Week 7: the same tests pass for both classes, and one of them is quadratic. Plot your own with the 07-queues notebook once `dsa/queue.py` works.

Four queues, compared

	enqueue	dequeue	memory
<code>SlowQueue</code> — array, <code>pop(0)</code>	amortised $O(1)$	$O(n)$	compact
Array, head index, no wrap	amortised $O(1)$	$O(1)$	dead slots grow for ever
<code>CircularQueue</code> — ring buffer	$O(1)$ (amortised if it grows)	$O(1)$	compact, reuses slots
Linked, head + tail	$O(1)$	$O(1)$	a node per item

	enqueue	dequeue	memory
Two stacks	amortised $O(1)$	amortised $O(1)$	two arrays

The ring buffer is the usual choice when the maximum size is known or growth is rare: it is contiguous, it allocates nothing per item, and it never moves data except while growing. The linked queue has the best worst case, one allocation per item. Python's `collections.deque` combines the two: a doubly linked list of fixed blocks of 64 slots, so it allocates rarely, like an array, and never shifts, like a list of nodes.

Where Queues Are Used

FIFO is fairness — and breadth

- **Scheduling:** print jobs, requests to a web server, processes sharing a CPU (round-robin).
- **Buffers:** keyboard input, network packets, audio — ring buffers between a producer and a consumer; “the last k values” for a moving average.
- **Breadth-first search** (Week 14): visit everything 1 step away, then 2 steps, ... — the queue keeps the frontier in order.
- **Level-order traversal** of a tree (Week 11): the same idea on a tree.

Round-robin. An operating system gives each ready process a short time slice. A process whose slice runs out goes to the back of the ready queue; the next one is dequeued from the front. Everyone gets a turn, in order. The same dequeue-then-enqueue step solves the Josephus problem — people in a circle, every k -th one leaves — which is W7-C2: put everyone in a queue, then move $k - 1$ people from the front to the back and remove the k -th, until one is left.

Breadth-first order. A queue processes things in the order they were *discovered*. Start from one node, enqueue its neighbours, then theirs: everything one step away is dequeued before anything two steps away. That is why BFS finds shortest paths (by number of edges), and why `level_order` in `dsa/tree.py` visits a tree row by row. Replace the queue with a stack and the same loop becomes depth-first — Week 14 puts them side by side.

A small example you can run now: the binary numbers 1 to n , in order, generated breadth-first (W7-C5). Start with “1”; each time you dequeue a string s , enqueue $s + “0”$ and $s + “1”$.

The last k values. A `CircularQueue(k)` that dequeues the oldest value whenever a new one arrives always holds the latest k : the window of a moving average, the last k log lines, the last k readings of a sensor. Keep a running total beside it and each new average costs $O(1)$ instead of $O(k)$ — W7-C3.

The deque: both ends, both ways

A **deque** (double-ended queue, say “deck”): add and remove at **both** ends, all $O(1)$.

- only one end used → a **stack**
- add at one end, remove at the other → a **queue**

In Python: `collections.deque` — `append`, `appendleft`, `pop`, `popleft`.

The name *deque* was coined by E. J. Scheppe, according to Knuth (*The Art of Computer Programming*, vol. 1, §2.2.1). A ring buffer makes a deque with no new ideas: adding at the front means moving head **back** one slot, `head = (head - 1) % capacity`. Python's `%` always returns a non-negative result for a positive capacity, so `(0 - 1) % 8` is 7: the step back from slot 0 wraps round to the last slot, exactly as needed. (In C and Java, `-1 % 8` is `-1`, and you must write `(head - 1 + capacity) % capacity` instead.) W7-S3 draws a deque on a ring, step by step.

Where next. Week 12 builds a **priority queue**: a queue that serves the most **urgent** item rather than the oldest one. The ADT has the same two operations, but a completely different structure underneath, the heap. Same name, different contract.

This Week

Exercises: `dsa/queue.py`

Method	Target	The trap
<code>SlowQueue.enqueue</code>	amortised $O(1)$	<code>append</code> on <code>self._items</code>
<code>SlowQueue.dequeue</code>	$O(n)$ — on purpose	<code>IndexError</code> when empty; <code>pop(0)</code>
<code>CircularQueue.enqueue</code>	$O(1)$	<code>tail = (head + size) % capacity</code> ; full?
<code>CircularQueue.dequeue</code>	$O(1)$	<code>IndexError</code> when empty; head wraps
<i>challenge</i> : grow when full	amortised $O(1)$	copy in queue order; <code>head = 0</code>

```
pytest tests/test_stack_queue.py -v -k queue
```

`SlowQueue` stands on your `DynamicArray`, so `tests/test_dynamic_array.py` must pass first. `CircularQueue` stands directly on the course `Array`: its `__init__`, `is_empty`, `is_full`, `__len__`, `__iter__` and `__repr__` are given to you — read `__iter__` first, since it already contains the ring arithmetic.

Four tests cover the queues: FIFO order and the empty `IndexError` for both classes, a wrap-around at capacity 3, and ten interleaved operations that take a capacity-4 ring round twice.

Homework 7 — before Lecture 08

1. **Implement** `dsa/queue.py` until the queue tests pass — then the challenge: grow instead of raising, and test it on a ring that has wrapped.
2. **Measure** `SlowQueue` against `CircularQueue` in `notebooks/07-queues.ipynb` and plot both. Estimate how long the `SlowQueue` would take for $n = 10^6$.
3. **Trace** a `CircularQueue(capacity=3)` through: `enqueue 1, 2, 3`; `dequeue`; `enqueue 4`; `dequeue`; `dequeue`; `enqueue 5, 6`. Show the block, head and size after each step.
4. **Design**. A queue from two stacks: write it, and explain why it is amortised $O(1)$ (W7-C1).

For item 2, a sensible range is $n = 500$ to 8,000; beyond that the `SlowQueue` takes minutes. Fit the slope of your log-log line and extrapolate: slope 2 means 10 times more items cost 100 times more time.

For item 3, check your final answer with the given `__iter__`: `list(q)` should read the queue front to back.

Summary

Seven things to keep

1. A queue is **FIFO**: enqueue at the **back**, dequeue at the **front**, both $O(1)$.
2. Empty dequeue raises **IndexError**.
3. On a plain array, one end is always the front: `pop(0)` or `insert(0)` — $O(n)$.
4. The **ring buffer**: head moves, nothing shifts; `tail = (head + size) % capacity`.
5. `head == tail` means **full or empty**: keep size (or waste one slot).
6. Growing a ring: copy in **queue order**, then `head = 0`.
7. Linked queue: front at the **head**, back at the **tail**. Two stacks: **amortised** $O(1)$.

Next

Week 8 — Searching (*midterm window*). Seven ways to find a value, and what a **sorted** array buys you: from $O(n)$ to $O(\log n)$.

Sources and further reading

- **D. E. Knuth.** *The Art of Computer Programming*, vol. 1, *Fundamental Algorithms*, 3rd ed., Addison-Wesley, 1997, §2.2.1 “Stacks, queues, and dequeues” and §2.2.2 “Sequential allocation” — queues in a circular block, and the origin of the word *deque*.
- **M. T. Goodrich, R. Tamassia and M. H. Goldwasser.** *Data Structures and Algorithms in Python*, Wiley, 2013, chapter 6, “Stacks, Queues, and Deques” — an array-based circular queue that grows, written in Python.
- **T. H. Cormen, C. E. Leiserson, R. L. Rivest and C. Stein.** *Introduction to Algorithms*, 3rd ed., MIT Press, 2009, section 10.1, “Stacks and queues”, and exercise 10.1-6, a queue from two stacks.
- **CPython source.** `Modules/_collectionsmodule.c` — deque as a doubly linked list of 64-slot blocks.

Every figure in this lecture is generated by `tools/figures_l07.py`. The timing figure is real measurement and will differ slightly on your machine.