

Searching

DSA27 — Lecture 08 · Data Structures and Algorithms

Dr. Haitham A. El-Ghareeb

Week 8 · Fall 2026

Contents

Where We Are	2
From structures to algorithms	2
Today	2
Linear Search	2
Look at everything	2
Binary Search	3
Halve what is left	3
The algorithm, and the promise it keeps	3
How many comparisons?	4
Easy to state, hard to get right	5
Recursive binary search	5
Bounds: Searching for Positions	6
lower_bound and upper_bound	6
Three More Searches	6
Jump search: $O(\sqrt{n})$	7
Exponential search: $O(\log i)$	7
Interpolation search: guess, don't halve	8
Measured	8
Eight functions, one picture	9
When is sorting first worth it?	9
The eight functions	10
This Week	10
Exercises: <code>dsa/searching.py</code>	10
Homework 8 — before Lecture 09	11
The midterm window	11
Summary	11
Seven things to keep	11
Next	12
Sources and further reading	12

How to read this document. This is the handout for Lecture 08. It holds everything on the slides, plus what I said out loud. Eight functions, one idea: an array that is kept **sorted** can be searched without looking at most of it. This week is also the midterm window, so the handout ends with a map of what weeks 1–7 expect of you.

Slides: `DSA27-L08-slides.pdf` · Code: `dsa/searching.py` · Tests: `tests/test_searching.py`

Where We Are

From structures to algorithms

Weeks 4–7 built **structures**: dynamic array, linked list, stack, queue.

Weeks 8–10 are about **algorithms** on the simplest of them, the array:

- **Week 8**: find a value — **searching**
- **Weeks 9–10**: put the values in order — **sorting**

Today's question: **what does it buy you to keep an array sorted?**

Every structure so far had a find or an in, and every one of them was $O(n)$: walk until you see the value. This week shows that the walk is unavoidable only when you know **nothing** about where the value might be. An array whose values are in increasing order tells you a great deal, and each search below uses that knowledge in a different way.

Search in Arabic: البحث الثنائي *Binary search*: البحث الثنائي

Today

1. Linear search: the baseline
2. Binary search: halving — and why it is so easy to get wrong
3. lower_bound and upper_bound: positions, duplicates, insertion
4. Jump, exponential and interpolation search
5. Measured: eight functions, one picture
6. When is sorting first worth it?

Linear Search

Look at everything

Linear search: no order needed, and no shortcut — $O(n)$

look at each element in turn



`linear_search(values, 8) → 4`
5 comparisons here; n when the target is last or missing

- Works on **any** sequence, sorted or not.

- Found at index i : $i + 1$ comparisons. Missing: **n**. Worst case $O(n)$.

`linear_search` is the one function in `dsa/searching.py` that must work on **unsorted** data (`test_linear_search_works_on_unsorted_data`). Return the index of the **first** match, or -1 when there is none.

Best, worst, average. Best case: the target is first, 1 comparison. Worst case: it is last or missing, n comparisons. Average, if the target is present and equally likely to be at any position: $(1 + 2 + \dots + n)/n = (n + 1)/2$ comparisons — still $\Theta(n)$. Halving the constant does not change the class (Lecture 02).

On sorted data a linear search can stop early, as soon as it passes a value larger than the target: a miss then costs, on average, half the array instead of all of it. Still $O(n)$. To do fundamentally better, you must stop looking at the elements **in order**.

Linear search is not a bad algorithm. For a few dozen elements it is often the fastest search there is; it needs no sorting, no extra memory, and it is what Python's `in` and `list.index` do.

Binary Search

Halve what is left

Binary search for 41: each comparison discards half of what is left

step 1 16 left	<table> <tr> <td>3</td><td>6</td><td>8</td><td>12</td><td>15</td><td>19</td><td>21</td><td>24</td><td>27</td><td>31</td><td>34</td><td>38</td><td>41</td><td>45</td><td>48</td><td>52</td> </tr> <tr> <td>0</td><td>1</td><td>2</td><td>3</td><td>4</td><td>5</td><td>6</td><td>7</td><td>8</td><td>9</td><td>10</td><td>11</td><td>12</td><td>13</td><td>14</td><td>15</td> </tr> </table>	3	6	8	12	15	19	21	24	27	31	34	38	41	45	48	52	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	lo=0 hi=15 mid=7 24 < 41 → lo = 8
3	6	8	12	15	19	21	24	27	31	34	38	41	45	48	52																			
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15																			
step 2 8 left	<table> <tr> <td>3</td><td>6</td><td>8</td><td>12</td><td>15</td><td>19</td><td>21</td><td>24</td><td>27</td><td>31</td><td>34</td><td>38</td><td>41</td><td>45</td><td>48</td><td>52</td> </tr> <tr> <td>0</td><td>1</td><td>2</td><td>3</td><td>4</td><td>5</td><td>6</td><td>7</td><td>8</td><td>9</td><td>10</td><td>11</td><td>12</td><td>13</td><td>14</td><td>15</td> </tr> </table>	3	6	8	12	15	19	21	24	27	31	34	38	41	45	48	52	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	lo=8 hi=15 mid=11 38 < 41 → lo = 12
3	6	8	12	15	19	21	24	27	31	34	38	41	45	48	52																			
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15																			
step 3 4 left	<table> <tr> <td>3</td><td>6</td><td>8</td><td>12</td><td>15</td><td>19</td><td>21</td><td>24</td><td>27</td><td>31</td><td>34</td><td>38</td><td>41</td><td>45</td><td>48</td><td>52</td> </tr> <tr> <td>0</td><td>1</td><td>2</td><td>3</td><td>4</td><td>5</td><td>6</td><td>7</td><td>8</td><td>9</td><td>10</td><td>11</td><td>12</td><td>13</td><td>14</td><td>15</td> </tr> </table>	3	6	8	12	15	19	21	24	27	31	34	38	41	45	48	52	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	lo=12 hi=15 mid=13 45 > 41 → hi = 12
3	6	8	12	15	19	21	24	27	31	34	38	41	45	48	52																			
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15																			
step 4 1 left	<table> <tr> <td>3</td><td>6</td><td>8</td><td>12</td><td>15</td><td>19</td><td>21</td><td>24</td><td>27</td><td>31</td><td>34</td><td>38</td><td>41</td><td>45</td><td>48</td><td>52</td> </tr> <tr> <td>0</td><td>1</td><td>2</td><td>3</td><td>4</td><td>5</td><td>6</td><td>7</td><td>8</td><td>9</td><td>10</td><td>11</td><td>12</td><td>13</td><td>14</td><td>15</td> </tr> </table>	3	6	8	12	15	19	21	24	27	31	34	38	41	45	48	52	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	lo=12 hi=12 mid=12 41 = 41 → found at 12
3	6	8	12	15	19	21	24	27	31	34	38	41	45	48	52																			
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15																			

Compare the target with the **middle** element. Equal: found. Smaller: it can only be on the left. Larger: only on the right.

This is how you look up a word in a paper dictionary: open it in the middle, see whether your word comes before or after, and throw away the half it cannot be in. Each comparison halves the part still in play. In the figure, 16 values become 8, then 4, then 1: **four** comparisons find 41, where a linear search would need 13.

It works **only** because the values are sorted. On unsorted data the comparison with the middle tells you nothing about either half, and binary search returns wrong answers without any error. Nothing checks that the input is sorted — checking would itself cost $O(n)$ and defeat the purpose. The sorted order is a **precondition**: the caller's promise.

The algorithm, and the promise it keeps

Keep two indices, `lo` and `hi`: the part still in play is `values[lo..hi]`, both ends included.

1. `lo = 0, hi = n - 1`
2. `while lo <= hi: mid = (lo + hi) // 2`

- `values[mid] == target` → return `mid`
 - `values[mid] < target` → `lo = mid + 1`
 - otherwise → `hi = mid - 1`
3. return `-1`

Invariant: if the target is anywhere, it is in `values[lo..hi]`.

An **invariant** is a statement that is true before the loop, stays true after every step, and so is still true when the loop ends. Here: “if the target is in the array at all, it is between `lo` and `hi`”. At the start the range is the whole array, so it is true. Each step discards only elements that cannot be the target — everything up to `mid` when `values[mid]` is too small, everything from `mid` when it is too large — so it stays true. When `lo > hi`, the range is empty, and the invariant tells you the target is nowhere: return `-1`.

This is the whole of `binary_search`. The steps are written out here in words on purpose: turning them into Python, and getting every `+ 1` and `<=` right, is the exercise.

A miss, traced. Searching the same 16 values for 20:

lo	hi	mid	values[mid]	Decision
0	15	7	24	$24 > 20 \rightarrow hi = 6$
0	6	3	12	$12 < 20 \rightarrow lo = 4$
4	6	5	19	$19 < 20 \rightarrow lo = 6$
6	6	6	21	$21 > 20 \rightarrow hi = 5$
6	5			$lo > hi$: return <code>-1</code>

Four comparisons, and the loop ends with `lo = 6`: the position where 20 **would be** inserted. Remember that; it is the idea behind `lower_bound`.

How many comparisons?

Each comparison halves the range: after k comparisons at most $n/2^k$ remain.

n	at most $\lfloor \log_2 n \rfloor + 1$ comparisons
1,000	10
1,000,000	20
1,000,000,000	30

Binary search is $O(\log n)$ — a thousand times more data costs 10 more comparisons.

Why $\lfloor \log_2 n \rfloor + 1$: the range starts with n elements, and a comparison that does not find the target leaves at most half of them (rounded down). After $\lfloor \log_2 n \rfloor$ such halvings one element is left, and one more comparison settles it. For the 16 values above, $\log_2 16 = 4$, so at most 5: searching for 52, the last value, takes exactly 5.

The recurrence, in the language of Lecture 03: $T(n) = T(n/2) + O(1)$, which solves to $O(\log n)$. Space: $O(1)$ for the loop — just `lo`, `hi` and `mid`.

`test_binary_search_does_not_scan_linearly` searches a list of a million integers. A correct binary search reads 20 of them. An accidental linear scan somewhere in your function reads a million and the test becomes slow — which is the point of the test.

Easy to state, hard to get right

- `while lo < hi` with an inclusive `hi` → **misses** the last candidate.
- `hi = mid` with an inclusive `hi` → can **loop for ever**.
- `lo = mid` → can **loop for ever** when `hi = lo + 1`.
- `mid = (lo + hi) / 2` in Java or C → **overflows** when `lo + hi` exceeds the largest int.

Binary search has a famous history of being wrong. Knuth records that the first binary search was published in 1946, and the first that works correctly for **every** n not until 1962 (*TAOCP* vol. 3, §6.2.1). Jon Bentley gave the problem to professional programmers in courses at Bell Labs and IBM, with two hours to write it: about ninety percent of the programs had bugs (*Programming Pearls*, column 4). And in 2006 Joshua Bloch reported that the binary search in Java's own library — `java.util.Arrays.binarySearch`, which he had written himself — failed on arrays of more than about a billion elements, because `(low + high) / 2` overflowed a 32-bit int. The fix is `low + (high - low) / 2`. The bug had been in the library for nine years.

Python integers do not overflow, so the last bug cannot happen here — but you will write binary search in C or Java one day, so name it now.

The other three bugs have one cure: **decide what `lo` and `hi` mean, and keep every line faithful to it**. With `hi` inclusive (the range is `values[lo..hi]`), the loop runs while the range is non-empty, `lo <= hi`, and both updates must move past `mid`, because `mid` has just been ruled out. The infinite loops happen when an update does not move: with `lo = 5`, `hi = 6`, `mid` is 5, and `lo = mid` leaves everything as it was. Test the sizes 0, 1 and 2 by hand: nearly every binary-search bug shows up on one of them. The tests check the empty list and one element (`test_handles_empty_and_single_element`), and the first and last positions (`test_finds_first_and_last_element`).

Recursive binary search

Same idea: search the half that can hold the target, **by a call**.

- $O(\log n)$ time, but also $O(\log n)$ **stack** — the loop is $O(1)$ space.
- Pass `lo` and `hi` down. **Never slice**: `values[mid + 1:]` copies — $O(n)$ per call.

`binary_search_recursive(values, target, lo=0, hi=None)` keeps the same contract. The base case is the empty range, `lo > hi`: return `-1`. `hi=None` is the usual trick for a default that depends on another argument: the first call replaces it with `len(values) - 1`.


The slicing trap. A recursive search that calls itself on `values[:mid]` or `values[mid + 1:]` looks elegant and is wrong twice. It is slow: each slice copies up to $n/2$ elements, so the total is $n/2 + n/4 + \dots = O(n)$ — as slow as a linear search, while looking like a binary one. And it loses the position: the index returned by a call on a slice is an index into the slice, not into the original list, so every call on a right half must add an offset. Passing `lo` and `hi` avoids both problems. (Lecture 03 made the same point about `list_max` on slices.)

The depth is at most about $\log_2 n$ frames — 20 for a million elements — so Python’s 1,000-frame limit is no concern here, unlike the linear recursions of Lecture 03. The loop is still the better tool: same speed, no frames at all.

Bounds: Searching for Positions

lower_bound and upper_bound

Bounds are positions between elements: where x would be inserted



- `lower_bound(v, x)`: the first index with $v[i] \geq x$
- `upper_bound(v, x)`: the first index with $v[i] > x$
- Both may return `len(v)`. Neither returns `-1`.

`binary_search` answers “is it there, and where?”, and with duplicates the “where” is **any** of the matching positions. The bounds answer more useful questions, and always with a position:

- **Where would x go?** `lower_bound(v, x)` is the index at which inserting x keeps the array sorted — the insertion point.
- **How many copies of x ?** `upper_bound(v, x) - lower_bound(v, x)`. In the figure, $4 - 1 = 3$ copies of 2, and $4 - 4 = 0$ copies of 3 (`test_bounds_count_duplicates`).
- **The first occurrence?** `i = lower_bound(v, x)`; it is x if $i < \text{len}(v)$ and $v[i] == x$.
- **How many values lie in $[a, b)$?** `lower_bound(v, b) - lower_bound(v, a)`. A range query in $O(\log n)$ — which no hash table can answer (Week 13).


They are both binary searches that **never stop early**: they keep halving until the range is empty and return where it closed. The usual way to write them uses a **half-open** range `[lo, hi)` — `lo` included, `hi` excluded — starting from `lo = 0`, `hi = len(v)`, so that `len(v)` is a possible answer. The loop runs while `lo < hi`, and the only difference between the two functions is which comparison sends the search right.

Python has them already, in the standard library: `bisect.bisect_left` is `lower_bound`, `bisect.bisect_right` is `upper_bound`, and `bisect.insort` inserts into a sorted list. C++ calls them `std::lower_bound` and `std::upper_bound`, which is where the names in `dsa/searching.py` come from. Write your own first; then use the library.

Three More Searches

Jump search: $O(\sqrt{n})$

Jump search: jump $\sqrt{n}$ at a time, then walk forward inside one block



Jump ahead $\sqrt{n}$ at a time until you pass the target; then walk forward inside that one block.


With blocks of size m , there are at most n/m jumps and then at most m steps inside a block: about $n/m + m$ reads. That sum is smallest when $m = \sqrt{n}$, giving $2\sqrt{n}$ — for a million elements, about 2,000 reads against 20 for binary search. Slower. So why bother?

Because jump search only ever moves **forward**, in big steps and then small ones. On storage where going back is expensive — a tape, or data arriving over a network in order — that matters more than the number of comparisons. It is also simple enough to be hard to get wrong.

In the figure, $n = 25$ and the step is $\sqrt{25} = 5$. The jumps check slots 4, 9, 14 and 19 — the last element of each block — and stop at 39, the first block end that is not smaller than 37. The walk goes forward from slot 15 and finds 37 at slot 18. (`math.isqrt(n)` gives the integer square root.)

Exponential search: $O(\log i)$

Exponential search: double the bound, then binary search inside it



Check slots 1, 2, 4, 8, ... until one passes the target; then binary search between the last two bounds.


If the answer is at index i , the doubling stops after about $\log_2 i$ steps, and the binary search that follows runs on a range of about i elements: $O(\log i)$ in total. When the target is near the **front**, that beats a binary search over the whole array; when it is near the end, it costs about twice as much. In the figure, 23 is at index 11: the bounds 1, 2, 4 and 8 are all smaller, 16 is not, so binary search runs on slots 8 to 16.

Its real use is on sequences whose **length you do not know** — a stream, an unbounded function, a huge sorted file: you cannot pick a middle without an end, but you can keep doubling until you have one. Bentley and Yao published it in 1976 as “an almost optimal algorithm for unbounded

searching”. The doubling is the same trick as the `DynamicArray`’s growth in Lecture 04: double until big enough, and the total work stays proportional to the final size.

Watch the edge cases: check slot 0 first (the doubling starts at 1), and do not let the bound run past `len(values) - 1`.

Interpolation search: guess, don’t halve



$$\text{pos} = \text{lo} + \frac{(\text{target} - v[\text{lo}]) (\text{hi} - \text{lo})}{v[\text{hi}] - v[\text{lo}]}$$

Evenly spread values: $O(\log \log n)$ on average. Skewed: $O(n)$.

Looking up “Zaki” in a phone book, you do not open it in the middle: you open it near the end, because Z is near the end of the alphabet. Interpolation search does exactly that. It assumes the values rise in a straight line from $v[\text{lo}]$ to $v[\text{hi}]$ and computes where the target would sit on that line. On evenly spread values the guess is very close, and the range shrinks so fast that the expected number of steps is $O(\log \log n)$ — about 4 or 5 for a million elements, against 20 for binary search. (The idea is W. W. Peterson’s, 1957; the $\log \log n$ analysis is by Perl, Itai and Avni, 1978.)


The price is the **worst case**. In the right-hand panel, one huge value at the end makes the straight line useless: every guess lands at the far left, the range shrinks by one element per step, and the search becomes a slow linear scan — $O(n)$. That is the docstring’s warning: an average case that hides a bad worst case.

Three traps in the code, all tested by the random-data test:

- **Division by zero** when $v[\text{hi}] == v[\text{lo}]$: handle that range separately.
- **A guess outside the range** when the target is smaller than $v[\text{lo}]$ or larger than $v[\text{hi}]$: stop at once — it is not there.
- **Integer division**: the position must be an `int` index. Use `//`.

Measured

Eight functions, one picture



Left: time. On a million sorted integers, a linear search took about 17 milliseconds and a binary search about 4 **microseconds** — some four thousand times faster, and the gap doubles every time n doubles. The binary line is not perfectly smooth: at these tiny times, memory caches and timer noise show. Its rise over the whole range, from 1 to about 4 μ s while n grew a thousandfold, is the logarithm.

Right: work — how many elements each search reads, on average, counted exactly by passing the searches a sequence that counts every values[i]. The counts come from the reference implementation, and yours may differ by a small constant factor.

- **Linear:** about $n/2$ reads — slope 1 on the log-log plot.
- **Jump:** about $\sqrt{n}$ reads — slope $1/2$. 1,035 reads at $n = 2^{20}$, where $\sqrt{n} = 1,024$.
- **Binary:** 18.8 reads at $n = 2^{20}$, where $\log_2 n = 20$.
- **Exponential:** about twice binary, since it doubles up to the target and then binary searches a range as big again — 37.5 at $n = 2^{20}$.
- **Interpolation, even data:** flat at about 12 reads — about four steps of three reads each ($v[\text{lo}]$, $v[\text{hi}]$ and the guess), at every size. $\log \log n$ is as close to constant as a growing function gets.
- **Interpolation, skewed data:** more than 100,000 reads at $n = 65,536$ — **three times worse than linear search**, because each of its roughly $n/2$ one-slot steps reads three values.

The four searches on sorted data all pass the same tests. Measurement is what tells them apart.

When is sorting first worth it?

Binary search needs sorted data, and sorting costs $O(n \log n)$ (Weeks 9–10).

k searches on n values	Cost
linear search, k times	$O(kn)$
sort once, then binary search k times	$O(n \log n + k \log n)$

Sorting wins once **k is larger than about $\log n$** — for a million values, after a few dozen searches.

For $n = 1,000,000$: one linear search reads up to a million values; a good sort does about $n \log_2 n \approx 20$ million comparisons. So after roughly twenty searches the sort has paid for itself, and every search after that costs 20 reads instead of a million. That is the reasoning behind every database index: pay once to keep data ordered, then answer questions fast for ever.

It is also why **keeping** an array sorted as it changes matters: inserting at `lower_bound` keeps it sorted, but the insertion itself shifts elements — $O(n)$ (Lecture 02). A sorted array is fast to search and slow to change. Weeks 11 and 13 give two answers: the binary search **tree** (fast to search and to change, while it stays balanced) and the **hash table** ($O(1)$ average lookup, but no order — no bounds, no ranges).

The eight functions

Function	Needs	Time	Note
<code>linear_search</code>	nothing	$O(n)$	first match
<code>binary_search</code>	sorted	$O(\log n)$	loop, $O(1)$ space
<code>binary_search_recursive</code>	sorted	$O(\log n)$	$O(\log n)$ stack; never slice
<code>lower_bound</code>	sorted	$O(\log n)$	first $v[i] \geq x$; may be <code>len</code>
<code>upper_bound</code>	sorted	$O(\log n)$	first $v[i] > x$; may be <code>len</code>
<code>jump_search</code>	sorted	$O(\sqrt{n})$	forward only
<code>exponential_search</code>	sorted	$O(\log i)$	unknown length; near the front
<code>interpolation_search</code>	sorted, even	$O(\log \log n)$ avg, $O(n)$ worst	guesses

This Week

Exercises: `dsa/searching.py`

Function	The trap
<code>linear_search</code>	unsorted input; return the first match
<code>binary_search</code>	<code>lo <= hi</code> ; <code>mid ± 1</code> ; <code>-1</code> when missing
<code>binary_search_recursive</code>	pass <code>lo, hi</code> ; no slices
<code>lower_bound, upper_bound</code>	half-open <code>[lo, hi)</code> ; <code>>=</code> against <code>></code>
<code>jump_search</code>	<code>math.isqrt</code> ; the last, partial block
<code>exponential_search</code>	slot 0 first; bound capped at <code>len - 1</code>
<code>interpolation_search</code>	equal ends; target outside <code>[v[lo], v[hi]]</code>

```
pytest tests/test_searching.py -v
```

Forty tests. The four searches on sorted data are run against the same cases — found, missing, empty, one element, first, last — and against linear search on 30 random sorted lists (`test_agrees_with_linear_search_on_random_data`, with a fixed seed so a failure is repeatable). The values in those tests are distinct, so “any matching index” and “the first matching index” agree.

The functions take Python lists as **input**, which the storage rule allows: they only read them and store nothing.

Homework 8 — before Lecture 09

1. **Implement** `dsa/searching.py` until all forty tests pass.
2. **Trace** `binary_search` on `[2, 5, 8, 12, 16, 23, 38, 56, 72, 91]` for 23 and for 60, as tables of `lo`, `hi`, `mid`.
3. **Measure** linear against binary search in `notebooks/08-searching.ipynb`, and count reads with a counting sequence, as in the figure.
4. **Break it.** Change `lo <= hi` to `lo < hi` in your `binary_search`. Which tests fail, and why exactly those?

For item 3, the counting sequence is a class with `__len__` and a `__getitem__` that adds 1 to a counter before returning `self.values[i]`; your searches cannot tell it from a list. `tools/figures_l08.py` has one, `CountingReads`.

For item 4, predict before you run: which **positions** can a `lo < hi` search never examine?

The midterm window

Weeks 1–7 are examinable. For each week, can you:

- **1–2:** define an ADT; give Θ of a loop; prove a Big-O bound?
- **3:** write a recurrence and trace a call stack?
- **4:** explain amortised $O(1)$ append?
- **5–7:** draw a linked list, stack, queue or ring **after every operation**?

The **mock exam for weeks 1–7** is in the question bank.

The exact date and format of the midterm will be announced in the lecture and on the WhatsApp channel. Until then, the safest preparation is the one the final rewards: the weekly question banks, done **before** reading their answers, and the mock exam under exam conditions — in one sitting, closed book, no computer. Searching is not on the weeks 1–7 mock; it is on the final.

Summary

Seven things to keep

1. Linear search needs no order and costs $O(n)$.
2. Binary search needs **sorted** data and costs $O(\log n)$: 20 comparisons for a million.

3. Its invariant: if the target is anywhere, it is in `values[lo..hi]`.
4. Get the boundaries right: `lo <= hi`, `mid ± 1` — test sizes 0, 1, 2.
5. Recursive: pass `lo` and `hi`; a slice turns $O(\log n)$ into $O(n)$.
6. `lower_bound/upper_bound` return **positions**: insertion points and counts of duplicates.
7. Jump $O(\sqrt{n})$, exponential $O(\log i)$, interpolation $O(\log \log n)$ average but $O(n)$ **worst**.

Next

Week 9 — Basic sorting. Bubble, selection and insertion sort — and counting sort, which beats $O(n \log n)$ by never comparing at all.

Sources and further reading

- **D. E. Knuth.** *The Art of Computer Programming*, vol. 3, *Sorting and Searching*, 2nd ed., Addison-Wesley, 1998, §6.1 “Sequential searching” and §6.2.1 “Searching an ordered table” — binary search, its history, and interpolation search.
- **J. Bentley.** *Programming Pearls*, 2nd ed., Addison-Wesley, 1999, column 4, “Writing correct programs” — binary search, invariants, and the programmers who got it wrong.
- **J. Bloch.** “Extra, Extra — Read All About It: Nearly All Binary Searches and Mergesorts are Broken”, Google Research Blog, 2 June 2006.
- **J. L. Bentley and A. C. Yao.** “An almost optimal algorithm for unbounded searching”, *Information Processing Letters* 5(3), 1976 — exponential search.
- **Y. Perl, A. Itai and H. Avni.** “Interpolation search — a $\log \log N$ search”, *Communications of the ACM* 21(7), 1978.
- **Python documentation.** The `bisect` module — `bisect_left`, `bisect_right` and `insort`.

Every figure in this lecture is generated by `tools/figures_l08.py`. The diagrams are fixed traces; the measured figure times a working `dsa/searching.py` and will differ slightly on your machine.