

Advanced Sorting

DSA27 — Lecture 10 · Data Structures and Algorithms

Dr. Haitham A. El-Ghareeb

Week 10 · Fall 2026

Contents

Where We Are	2
From n^2 to $n \log n$	2
Today	2
Divide and Conquer	3
Three steps	3
Merge Sort	3
Split until trivial, merge on the way up	3
Merging two sorted runs	4
_merge: the code	4
merge_sort_steps: divide, conquer, combine	5
How much work? $T(n) = 2T(n/2) + n$	6
The price: $O(n)$ extra space	7
Quicksort	7
Partition, then recurse	7
Lomuto's partition	7
_partition: the code	8
Choosing the pivot	9
The worst case, measured	10
Why the average is $O(n \log n)$	10
Stack depth: recurse on the smaller side	11
Quicksort is not stable	11
Heap Sort	12
The array is the tree	12
Sift down: repair one parent	12
Build-heap in $O(n)$, then sort	13
Heap sort, traced	13
Heap sort: the trade-offs	14
Measured	14
Six sorts, one picture	14
Can anything beat $n \log n$?	15
What Python Actually Uses	16
sorted() and list.sort(): Timsort	16

Stability	16
Equal keys, input order	17
The table to remember	17
This Week	18
Exercises: dsa/sorting.py, second half	18
Homework 10 — before Lecture 11	18
Summary	19
Seven things to keep	19
Next	19
Sources and further reading	19

How to read this document. This is the handout for Lecture 10. It holds everything on the slides, plus what I said out loud. Last week's three sorts all cost $O(n^2)$. This week, three sorts that cost $O(n \log n)$ — merge sort, quicksort and heap sort — and one idea behind the first two: **divide and conquer**. Every trace in this handout was produced by running the reference code, and every number in the measured section was counted, not estimated.

Slides: DSA27-L10-slides.pdf · Code: dsa/sorting.py (the second half) · Tests: tests/test_sorting.py -k "merge or quick or heap"

Where We Are

From n^2 to $n \log n$

Week 9: bubble, selection, insertion — $O(n^2)$ comparisons.

n	$n^2/2$	$n \log_2 n$
1,000	500,000	10,000
1,000,000	500,000,000,000	20,000,000

Today: three sorts on the right-hand column — and the price each one pays.

At a million values, an $O(n^2)$ sort does about 25,000 times more work than an $O(n \log n)$ one. Nothing you can do to the constant closes a gap like that; only a different algorithm does. Week 9's sorts all move an element by **one** place per comparison, or compare every pair; this week's sorts each find a way to let one comparison settle the order of many elements at once.

Sorting in Arabic: الترتيب. Merge sort: الترتيب بالدمج. Quicksort: الترتيب السريع. Heap sort: الترتيب بالكومة.

Today

1. Divide and conquer
2. Merge sort: $T(n) = 2T(n/2) + n$, stable, $O(n)$ extra space
3. Quicksort: partition, pivots, and a worst case you will **measure**

4. Heap sort: the array is the tree — Week 12 arriving early
5. Measured: all six sorts, one picture
6. What Python's `sorted()` actually does
7. Stability, and the table to remember

Divide and Conquer

Three steps

1. **Divide** the problem into smaller problems of the same kind.
2. **Conquer** each one — recursively, until it is trivial.
3. **Combine** their answers into the answer.


Merge sort: divide is free, **combine** does the work. Quicksort: **divide** does the work, combine is free.

Divide and conquer in Arabic: فرق تسد. You have met the pattern twice. Lecture 03's recursions split a problem and trusted the recursive call to solve the smaller part; Lecture 08's binary search is divide and conquer with only **one** of the two halves kept, which is why it costs $\log n$ and not n . This week both halves are kept, and each has to be sorted.

The two sorts in this lecture are mirror images, and that is worth saying out loud before any code. Merge sort splits the array in the middle without looking at it, and does all its thinking when it **merges** two sorted halves. Quicksort thinks first — it **partitions** the array into “small” and “large” around a pivot — and then there is nothing left to combine: small, pivot, large is already in order.

Merge Sort

Split until trivial, merge on the way up



The figure is `merge_sort([38, 27, 43, 3, 9, 82, 10, 19])`. Going down, each range is cut at $\text{mid} = (\text{lo} + \text{hi}) // 2$ — no comparisons, no moves. At the bottom,


```

        j += 1
        k += 1
    while i < mid:
        scratch[k] = a[i]
        i += 1
        k += 1
    while j < hi:
        scratch[k] = a[j]
        j += 1
        k += 1
    for k in range(lo, hi):
        a[k] = scratch[k]

```

This is `_merge` from the reference solution, exactly (without its docstring). The runs are `a[lo:mid]` and `a[mid:hi]` — half-open ranges, as in Lecture 08's bounds. `i` walks the left run, `j` the right run, and `k` the next free slot of `scratch`. After the main loop, **one** of the two runs is exhausted and the other still has a sorted tail; only one of the two “leftover” loops does anything. The last loop copies the merged range back into `a`, because the next merge up the tree reads from `a`.

The `<=` is the whole of stability. When `a[i]` and `a[j]` are equal, `<=` takes from the **left** run — the element that came first in the input. Write `<` and the sort is still correct, and every test that compares numbers still passes; but equal keys come out in the wrong order. See “Stability” below.

merge_sort_steps: divide, conquer, combine

```

a = _copy(values)
scratch = Array(len(a))
yield snapshot(a), ()

def sort(lo, hi):
    if hi - lo <= 1:
        return
    mid = (lo + hi) // 2
    yield from sort(lo, mid)
    yield from sort(mid, hi)
    _merge(a, lo, mid, hi, scratch)
    yield snapshot(a), tuple(range(lo, hi))

yield from sort(0, len(a))
yield snapshot(a), ()

```


The body of `merge_sort_steps` in the reference. `sort(lo, hi)` sorts `a[lo:hi]`; ranges of 0 or 1 elements are the base case. The three steps of divide and conquer are the three lines in the middle: compute `mid`, sort both halves, merge.

Why yield from? `sort` is itself a generator — it yields a snapshot after each merge — so calling `sort(lo, mid)` only **creates** a generator and runs nothing. `yield from` runs it and passes each of its snapshots up to whoever is iterating `merge_sort_steps`. Forget it, write `sort(lo, mid)` on its own line, and nothing below the top level is ever sorted: the `sort` returns the input barely changed, and the tests fail.

One scratch Array, shared by every merge. Allocating a new scratch array in every call would also be correct, but n calls each allocating up to n slots is a lot of memory churn. One `Array(len(a))`, made once, is enough, because a merge of `a[lo:hi]` only ever uses `scratch[lo:hi]`. The extra space is therefore exactly n slots, plus the call stack, which is only $\log_2 n$ frames deep because the ranges halve.

The plain form, `merge_sort(values)`, runs this generator to the end and returns the last state — one implementation for the tests and the animation.

How much work? $T(n) = 2T(n/2) + n$



$\log_2 n$ levels, n work per level: $O(n \log n)$ — best, average and worst case alike.

A **recurrence** (Lecture 03) states the cost of a recursive function in terms of its own cost on smaller inputs. For merge sort on n elements: two sorts of $n/2$ elements, plus a merge that touches all n . So $T(n) = 2T(n/2) + cn$ for some constant c , with $T(1) = c$.

The recursion tree solves it by picture. The top call merges n elements. Its two children merge $n/2$ each — n in total. The four grandchildren merge $n/4$ each — n again. Every level adds up to n , because the ranges on one level partition the array. The ranges halve at each level, so after $\log_2 n$ levels they have one element and the recursion stops. Total: $n \cdot \log_2 n$. *Recursion tree* in Arabic: شجرة الاستدعاءات.

By unrolling, if you prefer algebra: $T(n) = 2T(n/2) + n = 4T(n/4) + 2n = 8T(n/8) + 3n = \dots = 2^k T(n/2^k) + kn$. Stop when $n/2^k = 1$, that is $k = \log_2 n$: $T(n) = nT(1) + n \log_2 n = O(n \log n)$.

Always $n \log n$: merge sort splits in the middle whatever the values are, so no input is bad for it. Sorted input only saves comparisons inside each merge — the left run empties first, after $n/2$ comparisons instead of about n — measured: 5,120 comparisons for 1,024 sorted values against 8,935 for random ones. The **moves** are the same either way: every level copies all n elements to scratch and back.

The price: $O(n)$ extra space

- Merge sort is **not in place**: scratch holds n slots.
- The call stack adds $O(\log n)$.
- Fine for $n = 10^6$ numbers; a real cost on a small device, or for data that barely fits in memory.

Merging two runs in place, without a second array, is possible but complicated and slower in practice; nobody teaches it in a first course, and the standard libraries do not do it. The honest summary is: merge sort buys a guaranteed $O(n \log n)$ and stability with n extra slots. Quicksort and heap sort, next, work inside the array itself.

The same merge is also how you sort data **too big for memory**: sort chunks that fit, write each sorted run to disk, then merge the runs, reading each one front to back. That is *external sorting*, and it is still merge sort — only the runs live in files.

Quicksort

Partition, then recurse

Pick a **pivot**. Rearrange so that


$$\underbrace{< \text{pivot}}_{\text{left part}} \quad \text{pivot} \quad \underbrace{\geq \text{pivot}}_{\text{right part}}$$

The pivot is now in its **final place**. Sort the two parts; nothing to combine.

C. A. R. Hoare invented quicksort in 1959 and published it in 1961 ("Algorithm 64: Quicksort", *Communications of the ACM* 4(7)). *Pivot* in Arabic: المحور. *Partition*: التقسيم.

After a partition, the pivot never moves again: everything smaller is to its left and everything larger or equal is to its right, and sorting those two parts cannot change that. So each partition fixes **one** element for good and splits the rest into two smaller, independent problems. How much smaller depends entirely on the pivot — and that is the story of quicksort.

Lomuto's partition



There are two classic partition schemes. Hoare's original moves two indices towards each other; the one in `dsa/sorting.py` is **Lomuto's**, popularised by Jon Bentley in *Programming Pearls* because it is the easier of the two to get right. It uses the **last** element as the pivot, one index `j` that scans left to right, and one boundary:

- `a[lo:boundary]` holds values **< pivot** (green in the figure);
- `a[boundary:j]` holds values **≥ pivot** (amber);
- `a[j:hi]` has not been looked at yet (white);
- `a[hi]` is the pivot.

Each step looks at `a[j]`. If it is at least the pivot, the amber region simply grows by one: nothing moves. If it is smaller, it is swapped with the first amber value, `a[boundary]`, and `boundary` moves right: the green region grows by one and the amber region shifts along. At the end, one last swap puts the pivot between the two regions, at `boundary`.

The figure traces `Q = [6, 3, 9, 1, 8, 2, 7, 4]`. The pivot is 4; the values 3, 1 and 2 are smaller and end up in `a[0:3]`; the pivot lands at index 3. The partition made 7 comparisons, one per non-pivot element: $n - 1$, always.

`_partition: the code`

```
def _partition(a, lo, hi):
    pivot = a[hi]
    boundary = lo
    for j in range(lo, hi):
        if a[j] < pivot:
            a[boundary], a[j] = a[j], a[boundary]
            boundary += 1
    a[boundary], a[hi] = a[hi], a[boundary]
    return boundary
```

Invariant: `a[lo:boundary] < pivot ≤ a[boundary:j]`.

Note the ranges: `lo` and `hi` here are **inclusive**, `a[lo..hi]`, with the pivot at `hi` — unlike `_merge`, whose `hi` is excluded. Both are fine; what is not fine is mixing them up within one function. The invariant is true before the loop (both regions are empty), each step keeps it true, and at the end `j == hi`, so the whole range except the pivot has been sorted into the two regions. The final swap moves the first amber value to the end — which keeps it amber — and puts the pivot at `boundary`.

`a[j] < pivot` is strict: values **equal** to the pivot go right. With many equal values that is a weakness — a list of a thousand equal values partitions into 0 and 999, the worst case again. Industrial quicksorts use a three-way partition (`<`, `=`, `>`) for this reason; it is not required here.

Choosing the pivot

```
def _choose_pivot(a, lo, hi, strategy, rng):
    if strategy == "first":
        return lo
    if strategy == "last":
        return hi
    if strategy == "random":
        return rng.randint(lo, hi)
    if strategy == "median3":
        mid = (lo + hi) // 2
        first, middle, last = a[lo], a[mid], a[hi]
        if first <= middle <= last or last <= middle <= first:
            return mid
        if middle <= first <= last or last <= first <= middle:
            return lo
        return hi
    raise ValueError(...)
```

The chosen pivot is swapped to `a[hi]`; then `_partition` runs as before.


`quick_sort(values, pivot="median3")` takes the strategy by name, and the reference passes it to this function (its `raise` line is shortened here). The index it returns is swapped with `a[hi]`, so that `_partition` never needs to know how the pivot was chosen:

```
p = _choose_pivot(a, lo, hi, pivot, rng)
a[p], a[hi] = a[hi], a[p]
q = _partition(a, lo, hi)
```

- **first / last:** free, and fine on random data. On sorted data, a disaster — next slide.
- **random:** a random index in `lo..hi`. No fixed input is bad for it; a bad run needs bad luck at every level. The reference seeds its generator (`random.Random(27)`) so that the animation is the same every time.
- **median3:** the median of the first, middle and last values — R. C. Singleton's suggestion (1969). Up to three extra comparisons per partition buy a pivot that is never the smallest or largest of those three; on sorted input it is the exact middle.

The median-of-three is written with comparisons, not with `sorted()`: `dsa/sorting.py` does not call the library's `sort` to write its own.

The worst case, measured



Sorted input, `pivot="first"`: every partition peels off **one** element. $(n-1) + (n-2) + \dots + 1 = n(n-1)/2$ comparisons: $O(n^2)$.

On $[0, 1, \dots, 14]$ with `pivot="first"`, the pivot is the smallest value in its range every time. The “small” part is empty, the “large” part has all the rest, and the next partition is only one element shorter. Fifteen levels for fifteen elements, against four for `median3`, which picks the exact middle of each sorted range and halves it.

Counted, not guessed. The reference `quick_sort(list(range(n)), pivot="first")` makes exactly $n(n-1)/2$ comparisons: 523,776 for $n = 1,024$, and 2,096,128 for $n = 2,048$ — four times as many for twice the input, the signature of n^2 . With `pivot="median3"` the same inputs cost 9,228 and 20,493 comparisons. `pivot="last"` is just as bad as “first” on sorted input (and on reversed input, “first” and “last” swap roles).

This is not a curiosity. Sorted and nearly sorted data are the **most common** inputs a real sort receives — a list sorted yesterday with a few new entries today. A quicksort that picks the first element is at its worst on exactly the data it sees most. That is the single best argument for randomisation in this course: a random pivot makes the running time depend on the coin, not on the input.

Why the average is $O(n \log n)$

- A pivot anywhere in the **middle half** splits no worse than $1/4 : 3/4$.
- Then the largest part shrinks by $3/4$ per level: about $\log_{4/3} n$ levels of $O(n)$ work.
- Half of all pivots are in the middle half. **Random pivots: expected $O(n \log n)$** , about $1.39 n \log_2 n$ comparisons.

The intuition, not the proof. Call a pivot *good* if it lands in the middle half of its range, so that neither part holds more than three quarters of it. After a good pivot, the larger part has at most $\frac{3}{4}$ of the elements, so a chain of good pivots reaches size 1 after $\log_{4/3} n \approx 2.4 \log_2 n$ steps. A random pivot is good half the time, so on average it takes about two partitions to get one good one: about $2 \times 2.4 \log_2 n$ levels, each doing at most n comparisons. That is still $O(n \log n)$ — a bigger constant, not a bigger class. The exact analysis (Hoare’s; see

Sedgewick and Wayne, *Algorithms*, §2.3) gives $2n \ln n \approx 1.39 n \log_2 n$ expected comparisons.

In practice quicksort is often the **fastest** of the three: its inner loop is one comparison and at most one swap, it works in place, and it scans memory in order, which caches love. That is why it is called quick, despite a worst case that is worse than merge sort's.

Stack depth: recurse on the smaller side

```
def sort(lo, hi):
    while lo < hi:
        p = _choose_pivot(a, lo, hi, pivot, rng)
        a[p], a[hi] = a[hi], a[p]
        q = _partition(a, lo, hi)
        yield snapshot(a), (q,)
        if q - lo < hi - q:
            yield from sort(lo, q - 1)
            lo = q + 1
        else:
            yield from sort(q + 1, hi)
            hi = q - 1
```

The **smaller** part is a call; the **larger** part is the next turn of the loop. Stack depth: $O(\log n)$, even in the worst case.

The inner function of `quick_sort_steps`, exactly. The obvious version — partition, then two recursive calls — has a hidden cost. On sorted input with `pivot="first"` its recursion is n deep, one frame per peeled-off element, and Python stops at 1,000 frames: the obvious quicksort crashes with `RecursionError` on a sorted list of a few thousand elements, which is exactly when it is also at its slowest. `test_quick_sort_survives_its_worst_case_input` sorts `list(range(200))`, but with the default pivot and well under the limit, so it cannot catch this; Lab 10 has you try `pivot="first"` on 1,000 values.

The fix is a classic (it is in Sedgewick's *Algorithms*): make a recursive call only for the **smaller** part, and handle the larger part by updating `lo` or `hi` and going round the while loop again — the tail call turned into a loop, as in Lecture 03. Each call is on a range at most half the size of its caller's, so the stack is at most $\log_2 n$ frames deep, whatever the pivots. On sorted input with `pivot="first"` the smaller part is always **empty**, so the stack never grows at all. The time is still $O(n^2)$ on that input; only the space is fixed.

Extra space, then: $O(\log n)$ for the stack, nothing else. Quicksort is in place.

Quicksort is not stable


The partition swaps across long distances: an element can jump over its equal.

`quick_sort` on `[3a, 1b, 3c, 2d, 1e, 2f]` (sorted by number) gives `[1b, 1e, 2d, 2f, 3c, 3a]` — **3c before 3a**.

The letters are there only to tell equal numbers apart: the sort compares the numbers alone. The first partition picks the median of 3a, 3c and 2f, which is 3c (the middle one), and swaps it to the end: [3a, 1b, 2f, 2d, 1e, 3c]. The four smaller values are swapped to the front one by one, pushing 3a along, and the final swap puts the pivot 3c at index 4 — and 3a at index 5, behind its equal. Nothing in the rest of the algorithm puts them back. The stability figure below shows the whole comparison.

Heap Sort

The array is the tree



A **max-heap**: every parent $\geq$ its children. So the maximum is at $a[0]$.

This is Week 12 arriving early; there it becomes a priority queue. Any array can be **read** as a complete binary tree: $a[0]$ is the root, and the children of $a[i]$ are $a[2i + 1]$ and $a[2i + 2]$. The tree is not stored anywhere — there are no nodes and no pointers, only index arithmetic. `viz.draw.draw_array_as_tree` draws any list this way. *Heap* in Arabic: الكومة.

A **max-heap** is an array whose tree has the heap property: every parent is at least as large as each of its children. That does **not** make it sorted — in the figure, 5 is below 7 but 8 is not below 7 — but it does put the maximum at the root, where it can be read in $O(1)$.

The tree is **complete**: every level full except possibly the last, which fills from the left. So its height is $\lfloor \log_2 n \rfloor$, and anything that walks from the root to a leaf costs $O(\log n)$.

Sift down: repair one parent

```
def _sift_down(a, index, size):
    while True:
        largest = index
        left, right = 2 * index + 1, 2 * index + 2
        if left < size and a[left] > a[largest]:
            largest = left
```

```

    if right < size and a[right] > a[largest]:
        largest = right
    if largest == index:
        return
    a[index], a[largest] = a[largest], a[index]
    index = largest
    
```

`_sift_down(a, index, size)` assumes that both subtrees of `index` are already heaps, and only `a[index]` may be too small. It swaps `a[index]` with its larger child until it is at least as large as both children, or has none. **The larger child**, not just any larger one: otherwise the child that moves up could be smaller than its new sibling. `size` says where the heap ends, so that the sorted part at the back of the array is never touched. Each swap goes one level down: at most $\log_2 n$ swaps, $O(\log n)$.


Build-heap in $O(n)$, then sort

```

for index in range(n // 2 - 1, -1, -1):           # build the heap: O(n)
    _sift_down(a, index, n)
for end in range(n - 1, 0, -1):
    a[0], a[end] = a[end], a[0]                  # the max goes to the back
    _sift_down(a, 0, end)                       # restore the heap in a[:end]
    
```

1. **Build** a max-heap: sift down every parent, last to first.
2. **Sort**: swap the max to the back; the heap shrinks by one; sift down.

Heap sort, traced



The two loops of `heap_sort_steps` (the reference also yields a snapshot after the first loop and after each turn of the second).

Build-heap. The leaves — the last $\lceil n/2 \rceil$ slots — are one-element heaps already. Sift down every parent, from the last one, $n // 2 - 1$, back to the root: when a node is sifted, its subtrees are already heaps, which is exactly what `_sift_down` needs. On $H = [4, 10,$

3, 5, 1, 8, 2, 7], the sifts at indices 3, 2, 1 and 0 give [10, 7, 8, 5, 1, 3, 2, 4] — the heap in the previous figure.

Why $O(n)$ and not $O(n \log n)$? $n/2$ sifts, each up to $\log n$ deep, suggests $O(n \log n)$ — true, but not tight. Most nodes are near the bottom, where a sift is short: about $n/4$ nodes can sink 1 level, $n/8$ can sink 2, $n/16$ can sink 3, and only the root can sink $\log_2 n$. The total is at most $n(1/4 + 2/8 + 3/16 + \dots) = n$ swaps. This is R. W. Floyd's improvement (1964) on J. W. J. Williams's original heapsort of the same year, and you will prove it properly in Week 12.

Sort. The maximum is at $a[0]$. Swap it with the last slot of the heap, $a[\text{end}]$, where it belongs for good; the heap is now one shorter, and only its new root may break the heap property — one `_sift_down`, $O(\log n)$. Repeat $n - 1$ times: $O(n \log n)$. The sorted part grows from the back, like selection sort's — heap sort **is** selection sort, with a heap to find the maximum in $O(\log n)$ instead of $O(n)$.


Heap sort: the trade-offs

- $O(n \log n)$ **always** — no bad input, like merge sort.
- $O(1)$ **extra space** — in place, like quicksort; no recursion at all.
- **Not stable**: the swap to the back jumps over equals.
- Slower in practice: it jumps around memory ($i \rightarrow 2i + 1$).

On paper heap sort has the best of both: merge sort's guarantee and quicksort's memory. In practice it is usually the slowest of the three, because its sift reads $a[i]$ and then $a[2i + 1]$, far away for large i , so the processor's caches help it little; and it makes more comparisons (measured below: about $1.7 n \log_2 n$ on random data). Its real job is as a **safety net**: C++'s `std::sort` is *introsort* (Musser, 1997), a quicksort that switches to heap sort when its recursion gets suspiciously deep, so the worst case stays $O(n \log n)$.

Measured

Six sorts, one picture



Each point is a **count** of comparisons, made by the reference `dsa/sorting.py` on values wrapped in a small class that adds 1 to a counter inside `__lt__`, `__le__`, `__gt__` and `__ge__` — the same trick as `CountingReads` in Lab 08. Counts do not depend on your machine.

Left, random input. The three $O(n^2)$ sorts climb with slope 2 on the log-log plot; the three $O(n \log n)$ sorts climb with slope just over 1. At $n = 512$, bubble sort made 130,680 comparisons and merge sort 3,945 — 33 times fewer — and the gap doubles each time n doubles. Among the fast three, at $n = 4,096$: merge sort 43,943, quicksort (median of three) 61,330, heap sort 85,710. The dotted line is $\log_2(n!)$, the fewest comparisons **any** comparison sort can guarantee (see “Can anything beat $n \log n$?” below), and merge sort sits almost on it. So does `sorted()`: 43,939.

Right, sorted input. `pivot="first"` is the red line with slope 2 — $n(n-1)/2$ exactly. `random` and `median3` stay with the $n \log n$ group. Merge sort needs only $\frac{1}{2}n \log_2 n$ here (5,120 at $n = 1,024$), and heap sort does not care: it makes a little more than on random data. And Python’s `sorted()` makes **$n - 1$** comparisons — 1,023 at $n = 1,024$ — because it notices that the input is one sorted run and stops. That is the next section.

Why count instead of time? Because timing the course code can measure something else. `merge_sort` is `merge_sort_steps` run to the end, and a frame made with `list(a)` copies the whole array: after each of the $n - 1$ merges, n values — $O(n^2)$ copying in total, more than the sort itself. With that version, `merge_sort` took 0.36, 1.29 and 4.7 seconds for 1,000, 2,000 and 4,000 values on this machine: four times as long for twice the input, which is what n^2 looks like. So every `_steps` form in the reference takes a snapshot parameter — `list` by default, so the animation gets fresh frames — and the plain form passes a function that returns a itself, uncopied. That is the `snapshot(a)` in the listings above. With it, the same runs took 0.06, 0.12 and 0.27 seconds: a little over twice the time for twice the input, the $n \log n$ shape. Counting comparisons is still the cleaner measure, because a time includes the interpreter and whatever else the machine is doing; Lab 10, Part 8, has you measure both versions. (Python’s `sorted()`, written in C, sorts 4,000 values in under a millisecond.)

Can anything beat $n \log n$?

A comparison sort must tell apart all $n!$ orders of its input. Each comparison has two outcomes, so it needs at least $\log_2(n!) \approx n \log_2 n - 1.44n$ comparisons in the worst case.

Merge sort and heap sort are optimal — up to a constant.

Think of any comparison sort as a game of twenty questions: each comparison is a yes/no question about the input, and the sort must ask enough to know which of the $n!$ possible orders it has been given — otherwise two different inputs get the same treatment and one of them comes out unsorted. k yes/no answers can tell apart at most 2^k cases, so $2^k \geq n!$, i.e. $k \geq \log_2(n!)$. By Stirling’s formula that is about $n \log_2 n - 1.44n$. For $n = 1,024$: $\log_2(1024!) \approx 8,769$, and merge sort used 8,935 on random input.

The bound is about **comparison** sorts only. Counting sort (Week 9) never compares two elements — it uses the values as indices — and runs in $O(n+k)$. That is not a contradiction; it is a different game.

What Python Actually Uses

`sorted()` and `list.sort()`: Timsort

Enrichment — not examinable.

1. Find the **runs** already in the data (reversing descending ones).
2. Extend short runs to 32–64 elements with **insertion sort**.
3. **Merge** the runs, like merge sort — stable, $O(n)$ extra space.

Sorted input: one run, **$n - 1$ comparisons**. Worst case: $O(n \log n)$.

Python's sort is not quicksort. It is **Timsort**, written by Tim Peters for Python 2.3 in 2002 and described in his notes `listsort.txt` in the CPython source. It is a merge sort that takes the data's existing order seriously:

- It scans for **natural runs** — stretches that are already ascending, or strictly descending (which it reverses in place). Real data is full of them.
- A run shorter than a minimum length (between 32 and 64, chosen from n) is extended with **binary insertion sort** — Week 9's insertion sort, which is the fastest thing there is on short or nearly sorted ranges.
- The runs are **merged** in an order that keeps the merges balanced, using a temporary array — merge sort's price. During a merge, when one run keeps winning, it switches to **galloping**: an exponential search (Lecture 08) for how far that run can be copied at once.

It is **stable** — Python guarantees it — and it is what you measured on the right-hand side of the figure: $n - 1$ comparisons on sorted input, because the whole list is one run and there is nothing to merge. Since Python 3.11 the order in which runs are merged follows the *powersort* rule of Munro and Wild (2018). Java uses Timsort to sort objects, and a dual-pivot quicksort for primitive types, where stability cannot be observed. Even Timsort had a bug in its run-merging invariant, found in 2015 by a team **formally verifying** it (de Gouw et al.) — the binary-search story of Lecture 08 again: correct-looking code, broken on an input nobody had tried.

What this means for you: in your own programs, use `sorted()` or `list.sort()`, and `key=` for anything but plain values. In `dsa/sorting.py`, the point is that you can write the algorithms it is made of.

Stability

Equal keys, input order

Stability: equal keys (same number) keep their input order — or not

input	3a	1b	3c	2d	1e	2f	
merge_sort	1b	1e	2d	2f	3a	3c	stable: a before c, b before e, d before f
quick_sort (median3)	1b	1e	2d	2f	3c	3a	NOT stable: equal keys swapped (red)
heap_sort	1e	1b	2f	2d	3c	3a	NOT stable: equal keys swapped (red)

A sort is **stable** if equal keys keep their input order.

Stable in Arabic: مستقر. Stability matters as soon as you sort **records by one field**. Sort students by name, then **stably** by grade: within each grade they are still in name order — two stable sorts give a sort by grade, then name. With an unstable second sort the name order within a grade is lost.

The figure sorts six cards compared by their **number only**; the letters are there so you can see the order. Merge sort keeps a before c, b before e and d before f, because its merge takes from the left run on a tie. Quicksort's partition and heap sort's swap-to-the-back both move elements across long distances, and swap equal pairs (red). Heap sort, here, reversed all three.

Testing stability needs the right values. Sorting (key, tag) **tuples** proves nothing, because Python compares tuples by the tag too when the keys are equal — any correct sort puts (1, "a") before (1, "c"). To see stability, compare by the key alone: a small class whose `__lt__`, `__le__`, `__gt__` and `__ge__` look only at the key, as in `tools/figures_l10.py` (Card), in `test_insertion_sort_is_stable`, and in Lab 10.

The table to remember

Sort	Best	Average	Worst	Extra space	Stable
bubble	n	n^2	n^2	$O(1)$	yes
selection	n^2	n^2	n^2	$O(1)$	no
insertion	n	n^2	n^2	$O(1)$	yes
merge	$n \log n$	$n \log n$	$n \log n$	$O(n)$	yes
quick	$n \log n$	$n \log n$	n^2	$O(\log n)$	no
heap	$n \log n$	$n \log n$	$n \log n$	$O(1)$	no
counting	$n + k$	$n + k$	$n + k$	$O(n + k)$	—
Timsort	n	$n \log n$	$n \log n$	$O(n)$	yes

Times are in comparisons (counting sort: steps), all $O(\cdot)$. Notes on the rows:

- **bubble** best case n : with the early exit of the reference — one pass with no swap on sorted input.

- **quick** extra space $O(\log n)$: with the smaller-side recursion. The naive version needs $O(n)$ stack in the worst case. Its worst case is n^2 for every fixed pivot rule; with random, it is n^2 only with vanishing probability.
- **heap** best case: with distinct keys; if every key is equal each sift stops at once and it is $O(n)$.
- **counting**: k is the largest value plus one. The course's `counting_sort` rebuilds the integers from their counts, so stability does not apply to it; the textbook version that moves whole records is stable.

Which to use? **Guarantee and stability**: merge sort (or Timsort). **Speed in place, on average**: quicksort with a random or median-of-three pivot. **Worst case and no extra memory**: heap sort. **Tiny or nearly sorted arrays**: insertion sort — which is why Timsort and introsort both use it at the bottom.

This Week

Exercises: `dsa/sorting.py`, second half

Function	The trap
<code>merge_sort_steps</code>	yield from the recursive calls; <code><=</code> in the merge; one scratch Array
<code>quick_sort_steps</code>	Lomuto's invariant; swap the pivot to <code>hi</code> ; smaller side first
<code>heap_sort_steps</code>	build from <code>n // 2 - 1</code> down to 0; sift with <code>size = end</code>

```
pytest tests/test_sorting.py -v -k "merge or quick or heap"
```

37 tests: each sort on the eight fixed cases (empty, one, two, five mixed values, duplicates, sorted, reversed, negatives), 20 random lists, “does not mutate the input”, the two generator tests, and quicksort on sorted input. The plain forms are one line each once the `_steps` forms work. Write private helpers — `_merge`, `_partition`, `_choose_pivot`, `_sift_down` — as the lecture does; the tests only call the public names.

Two things the tests do **not** check, and the lab does: stability (sort cards compared by key only) and the pivot strategies (pivot="first" against "median3" on sorted input, counted).

The storage rule holds: copy the input into an Array, sort the Array, and return a list. No `sorted()`, `list.sort()` or `heapq` inside `dsa/sorting.py`.

Homework 10 — before Lecture 11

1. **Implement** merge, quick and heap sort until the 37 tests pass.
2. **Trace** `_partition` on `[5, 8, 1, 9, 3, 7, 2, 6]` as a table of `j`, `a[j]`, boundary and the array.

3. **Draw** the merge tree of `merge_sort([6, 5, 3, 1, 8, 7, 2, 4])`.
4. **Count** comparisons for `quick_sort(list(range(n)), pivot)` with "first" and "median3", $n = 100, 200, 400$. Predict first.
5. **Break it:** change `<=` to `<` in `_merge`. Which tests fail? What else changes?

For item 4, reuse the idea of Lab 08's `CountingReads`: a class whose comparison methods add to a counter (Lab 10, Part 6). Before you run it: $n(n-1)/2$ for "first"; for "median3", a little over $n \log_2 n$.

For item 5, predict before you run. (A hint: which of the tests compares values that are equal **but distinguishable**?)

Summary

Seven things to keep

1. **Divide and conquer:** split, solve the parts recursively, combine.
2. **Merge sort:** $T(n) = 2T(n/2) + n = O(n \log n)$ always; $O(n)$ extra; stable because of one `<=`.
3. **Quicksort:** partition around a pivot; the pivot lands in its final place.
4. `pivot="first"` on sorted input: $n(n-1)/2$ comparisons — **measured**. Random or median-of-three: $O(n \log n)$ on average.
5. Recurse on the **smaller** part: $O(\log n)$ stack, always.
6. **Heap sort:** the array is the tree; build-heap $O(n)$; $O(n \log n)$ always, $O(1)$ extra, not stable.
7. No comparison sort beats $\log_2(n!) \approx n \log_2 n$. Python's `sorted()` is **Timsort**: merge sort that exploits existing runs.

Next

Week 11 — Trees. Binary search trees: the sorted array's fast search, without its $O(n)$ insert. And an in-order traversal gives the values **sorted** — one more sort, for free.

Keep the pivot story in mind. A binary search tree built by inserting sorted values is exactly quicksort with `pivot="first"`: every new value goes to the right of all the others, and the "tree" becomes a linked list, n deep. The same bad input, the same cure (balance), in a different structure.

Sources and further reading

- **C. A. R. Hoare.** "Algorithm 64: Quicksort", *Communications of the ACM* 4(7), 1961; and "Quicksort", *The Computer Journal* 5(1), 1962.
- **J. W. J. Williams.** "Algorithm 232: Heapsort", *Communications of the ACM* 7(6), 1964;
- **R. W. Floyd.** "Algorithm 245: Treesort 3", *Communications of the ACM* 7(12), 1964 — build-heap in linear time.

- **R. C. Singleton.** "Algorithm 347: An efficient algorithm for sorting with minimal storage", *Communications of the ACM* 12(3), 1969 — median of three.
- **D. E. Knuth.** *The Art of Computer Programming*, vol. 3, *Sorting and Searching*, 2nd ed., Addison-Wesley, 1998, §5.2 and §5.3.1 — merge sort (which Knuth traces to John von Neumann, 1945), quicksort, heapsort, and the $\log_2(n!)$ lower bound.
- **J. Bentley.** *Programming Pearls*, 2nd ed., Addison-Wesley, 1999, column 11, "Sorting" — Lomuto's partition.
- **R. Sedgewick and K. Wayne.** *Algorithms*, 4th ed., Addison-Wesley, 2011, §2.2–2.4 — mergesort, quicksort (with the average-case analysis), heapsort.
- **D. R. Musser.** "Introspective sorting and selection algorithms", *Software: Practice and Experience* 27(8), 1997 — introsort.
- **T. Peters.** `Objects/listsort.txt` in the CPython source — Timsort, explained by its author.
- **J. I. Munro and S. Wild.** "Nearly-optimal mergesorts: fast, practical sorting methods that optimally adapt to existing runs", ESA 2018 — powersort.
- **S. de Gouw, J. Rot, F. S. de Boer, R. Bubel and R. Hähnle.** "OpenJDK's `java.util.Collection.sort()` is broken: the good, the bad and the worst case", CAV 2015.

Every figure in this lecture is generated by `tools/figures_l10.py` from the states and counts of a working `dsa/sorting.py` (run it with `python tools/with_solutions.py tools/figures_l10.py` to use the reference).