

Hash Tables

DSA27 — Lecture 13 · Data Structures and Algorithms

Dr. Haitham A. El-Ghareeb

Week 13 · Fall 2026

Contents

Where We Are	2
Finding a value by its key	2
Today	2
The Dream	3
Direct addressing: the key is the index	3
Hashing: squeeze the key into the table	3
Hash Functions	4
Python's hash ()	4
Same run, same hash. Next run — maybe not.	5
What makes a good hash function	5
Collisions	6
Collisions are certain	6
Two ways to live with collisions	7
Chaining	7
A linked list in every bucket	8
The three operations, in words	8
Trace: chaining, capacity 8, max_load 0.75	9
Cost: $1 + \alpha$	9
Open Addressing	10
Linear probing: one key per slot	10
Trace: open addressing, capacity 8, max_load 0.66	11
Primary clustering	11
Delete: why not just empty the slot?	12
Load Factor and Resizing	13
Resize: every key moves	13
Amortised $O(1)$ — Week 4 again	14
Measured	15
Probes against load factor	15
What breaks $O(1)$	16
The Real Thing	17

Python's dict and set — which you now understand	17
Hash table, sorted array or tree?	17
This Week	18
Exercises: dsa/hashmap.py	18
Homework 13 — before Lecture 14	19
Summary	19
Seven things to keep	19
Next	19
Sources and further reading	19

How to read this document. This is the handout for Lecture 13. It holds everything on the slides, plus what I said out loud. Every structure so far finds a value by walking or by halving. This week you build one that finds it by **computing where it is** — the structure behind Python's dict and set, and behind nearly every “look it up by name” in software. You will build it twice, measure it, and then break it on purpose, because its famous $O(1)$ is a promise about the **average**, and you should know exactly who can break it.

Slides: DSA27-L13-slides.pdf · Code: dsa/hashmap.py · Tests: tests/test_hashmap.py

Where We Are

Finding a value by its key

Structure	Find a key	Needs
unsorted array, linked list	$O(n)$	nothing
sorted array (Week 8)	$O(\log n)$	sorted order
binary search tree (Week 11)	$O(\log n)$ if balanced	order, balance
hash table (today)	$O(1)$ on average	a good hash function

Today's question: **can we find a key without comparing it to other keys at all?**

Every search in this course so far worked by **comparison**: is this the key? Is the key smaller or larger? Comparison-based search cannot beat $O(\log n)$ in the worst case — each comparison gives at most one bit of information, and telling n positions apart takes about $\log_2 n$ bits. A hash table escapes that bound the same way counting sort (Week 9) escaped $O(n \log n)$: it does not compare keys, it **computes** with them.

Hash table in Arabic: المفتاح. Key: دالة التجزئة. Hash function: جدول التجزئة.

Today

1. The dream: the key **is** the index
2. Hash functions: what hash() does, and what is and is not reproducible
3. Collisions, and the birthday paradox
4. Chaining: a linked list in every bucket

5. Open addressing: linear probing, clustering, **tombstones**
6. Load factor and resizing — amortised again
7. Measured: probes against load factor; and the keys that break it
8. Python's dict and set: the real thing, which you now understand

The Dream

Direct addressing: the key is the index

Keys are student numbers 0 to 999? Use an Array (1000):


- `put(k, v): a[k] = v —  $O(1)$`
- `get(k): a[k] —  $O(1)$`

No search at all. **But:** the space is the size of the **key range**, not the number of keys.

This is the best possible dictionary, and you have had it since Lecture 02: indexing an Array is $O(1)$, so if the key **is** an index, there is nothing to search. Counting sort used exactly this in Week 9 — one counter per possible value.

It fails as soon as the keys are not small integers. Mansoura student numbers have seven digits: an Array indexed by them needs ten million slots to hold a class of two hundred. A name, a URL or a tuple is not an index at all. We need a function that squeezes **any** key into a **small** range of indices — and we must then live with what that squeezing costs.

Hashing: squeeze the key into the table



$$\text{index} = \text{hash}(\text{key}) \bmod \text{capacity}$$

Two steps, and it pays to keep them apart:

1. **hash(key)** turns the key into an integer. That is Python's job: every hashable object has a `__hash__`, and `hash()` calls it. You do not write a hash function this week; the docstring of `_index` says so.

2. **% capacity** folds that integer into the table: a number from 0 to capacity - 1, an index into the Array. That is the table's job, and it is one line of `dsa/hashmap.py`, given to you:

```
def _index(self, key):
    return hash(key) % len(self._buckets)
```

Python's % always returns a non-negative result for a positive capacity, even for a negative hash: `-5 % 8` is 3. (In C and Java, `-5 % 8` is -5, and a negative index is a crash — the same trap as the ring buffer's step back in Lecture 07.)

The figure uses three seven-digit student numbers. For an `int`, `hash(n)` is just `n`, so the index is simply the last three bits of the number: 2027305, 2027114 and 2026988 land in slots 1, 2 and 4 of an 8-slot table. Ten million possible keys, eight slots. Something has to give — and it is the next section.

Hash Functions

Python's `hash()`

Expression	Value	Every run?
<code>hash(42)</code>	42	yes
<code>hash(-1)</code>	-2	yes
<code>hash(1.0)</code> , <code>hash(True)</code>	1	yes
<code>hash((1, 2))</code>	-3550055125485641917	yes
<code>hash("abc")</code>	a different number each run	no
<code>hash([1, 2])</code>	<code>TypeError: unhashable type: 'list'</code>	—

Every value in that table was printed by Python 3.14 on a 64-bit machine; the course's 3.13 gives the same. Three of them deserve a sentence.

- **`hash(n)` is `n`** for every integer you will meet in this course. (The exact rule is `n` reduced modulo the prime $2^{61} - 1$, so `hash(2**61 - 1)` is 0.) The one oddity is **`hash(-1)` is -2**: inside CPython, -1 is the return value that means “an error happened”, so no object may hash to it.
- **`hash(1.0) == hash(1) == hash(True)`**. Because `1 == 1.0 == True` in Python, and the one rule every hash function must keep is: **equal keys must have equal hashes**. Otherwise `d[1.0]` would look in a different bucket from `d[1]` and miss a key that is, by `==`, the same one.
- **A list is unhashable**. A list can change after it is put in a table; its hash would change with it, and the table would look for it in the wrong bucket. So Python refuses. Tuples of hashable things, strings, numbers and frozensets are hashable; lists, dicts and sets are not. Your `put` inherits this for free: `hash(key)` raises the same `TypeError`.

Same run, same hash. Next run — maybe not.

```
python -c "print(hash('abc'))"      # -60919065288970757
python -c "print(hash('abc'))"      # 2138548942786055565
python -c "print(hash(12345))"      # 12345, every time
```

- **Deterministic within one run:** always — a table depends on it.
- **Across runs:** int, float, tuples of them — yes. **str and bytes — no:** randomised on purpose.
- PYTHONHASHSEED=0 fixes the seed, for reproducing a bug.

Since Python 3.3, the hash of a `str` (and of `bytes`) is **randomised**: it is computed with a secret key chosen when the interpreter starts. Since Python 3.4 the function is `SipHash`, a keyed hash designed for exactly this job (PEP 456), and since 3.11 its faster `SipHash-1-3` variant (`sys.hash_info.algorithm` prints `siphash13`). So:

- **Inside one run**, `hash("abc")` is the same every time you ask. A table could not work otherwise.
- **In the next run** it is almost certainly different. So the order in which your `Chaining-HashMap` iterates string keys can change from one run to the next, and so can which bucket a string key lands in. A bug that depends on where keys land — the lab has one — can pass on one run and fail on the next.
- **Integers are not randomised.** Nor are floats, or tuples made of ints and floats. That is why every trace in this lecture, the lab and the question bank uses **integer keys**: `hash(n) = n`, and anyone can check them by hand.

Why randomise at all? Section 7 answers that: an attacker who can predict your hashes can choose keys that all collide. To make string hashes reproducible — to chase a bug that only shows on some runs — fix the seed before Python starts: `$env:PYTHONHASHSEED = "0"` in PowerShell, `PYTHONHASHSEED=0` in a Linux shell. Unset it afterwards.

(Objects of your own classes hash by **identity** unless you define `__hash__` — based on where the object lives in memory, so that is not reproducible across runs either.)

What makes a good hash function

1. **Consistent:** $a == b \Rightarrow \text{hash}(a) == \text{hash}(b)$. Non-negotiable.
2. **Spreads keys evenly** over the indices — similar keys, different slots.
3. **Uses the whole key**, not just part of it.
4. **Fast** — it runs on every single operation.

A bad one: **sum of the character codes**. `"listen"`, `"silent"` and `"enlist"` all hash to 655.

Rule 1 is about correctness; the other three are about speed. A hash function that returns 0 for everything is **correct** — every table built on it still finds every key — it is just $O(n)$, because every key lands in one bucket. Hold on to that: a bad hash never makes a table *wrong*, only *slow*. (W13-E2 asks you to argue it.)

The sum of character codes breaks rule 3 in a subtle way: it uses every character but ignores their **order**, so every anagram collides. Here is a toy hash you can compute by hand, with the fix:

```
def sum_hash(s):                                # bad: order-blind
    return sum(ord(c) for c in s)

def poly_hash(s):                               # better: position matters
    h = 0
    for c in s:
        h = 31 * h + ord(c)
    return h
```


`sum_hash` gives 655 for all three anagrams; `poly_hash` gives "listen" 3192458695, "silent" 3392640085 and "enlist" 2996453575 (these numbers are deterministic — no randomisation here, you can check them). `poly_hash` is the scheme Java's `String.hashCode` uses (Java then keeps only 32 bits). It is a polynomial in 31, evaluated by Horner's rule: each character is multiplied by a different power of 31 according to its position.

Why % with a power of two is risky. `dsa/hashmap.py` starts at capacity 8 and doubles: the capacity is always a power of two, and `h % 2**k` keeps only the **last k bits** of the hash. If the keys differ only in their high bits — multiples of 1,024, say — they all land in slot 0. Choosing a **prime** capacity spreads such keys; CPython and Java keep powers of two (so % is a fast bit mask) and instead **mix the high bits in**: Java's `HashMap` uses `h ^ (h >>> 16)`, and CPython's probe sequence folds in the high bits as it goes. Section 7 measures what happens when nobody does.

Collisions

Collisions are certain

- **Pigeonhole:** more possible keys than slots $\Rightarrow$ some keys **must** share.
- **Birthday paradox:** it happens **early**.



A **collision** is two different keys with the same index. They are not a sign of a bad hash

function — with ten million student numbers and eight slots, they are forced. The question is only how early they come.

Much earlier than intuition says. With 23 people in a room, the chance that two share a birthday is 50.7%, although there are 365 days. The exact probability that n keys placed at random in m slots all miss each other is

$$P(\text{no collision}) = \frac{m}{m} \cdot \frac{m-1}{m} \cdots \frac{m-n+1}{m}$$

and it drops below one half at about $n \approx 1.18\sqrt{m}$. For $m = 10,000$ slots, that is $n = 119$: a table **one percent full** has probably already had a collision. The crosses in the figure are 400 random trials at each n , and they sit on the formula.

Collision in Arabic: تصادم.

The lesson for the design: **every hash table needs a collision strategy**, not as a rare fallback but as a normal path that runs all the time. There are two families.

Two ways to live with collisions


	Chaining	Open addressing
a slot holds	a list of entries	one entry
on collision	add to the list	probe another slot
load factor	may exceed 1	must stay < 1
delete	unlink a node	leave a tombstone
used by	Java HashMap	Python dict, set

Chaining in Arabic: التسلسل. *Open addressing*: العنونة المفتوحة.

You build both. ChainingHashMap reuses what you learnt in Week 5 — every bucket is a small linked list of Entry nodes. OpenAddressingHashMap keeps everything in the Array itself and needs one genuinely new idea, the tombstone. Their tests are the same eight, run on each class, plus one test that only open addressing needs.

Chaining

A linked list in every bucket



`self._buckets` is an Array; each slot is `None` or the **first** Entry of a chain.

Entry is given to you, and it is a Week 5 node with one extra field:

```
class Entry:
    __slots__ = ("key", "value", "next")

    def __init__(self, key, value, next=None):
        self.key = key
        self.value = value
        self.next = next
```

The figure is the real state of a `ChainingHashMap(capacity=8)` after `put(10, '10')`, `put(22, '22')`, `put(31, '31')`, `put(4, '4')`, `put(15, '15')` and `put(28, '28')`, read back out of its buckets by `tools/figures_l13.py`. The indices are just the keys mod 8: $10 \rightarrow 2$, $22 \rightarrow 6$, $31 \rightarrow 7$, $4 \rightarrow 4$, $15 \rightarrow 7$, $28 \rightarrow 4$. Two collisions, two chains of two.

Why the front? A new key is linked in at the **head** of its chain — `push_front`, $O(1)$, with no walk to the tail. But `put` has to walk the chain anyway, to find out whether the key is already there (then it overwrites rather than adding a duplicate). Adding at the front is simply the shortest code once the walk has found nothing.

The storage rule holds: the buckets are an Array, the chains are node objects. No Python `list` or `dict` anywhere — which would, after all, be cheating on a hash table.

The three operations, in words

get(key): walk the chain at `_index(key)`; return the value of the entry whose `key == key`; at the end of the chain: `default`, or **KeyError**.

put(key, value): walk the same chain; if the key is there, **overwrite** and return. Otherwise push a new Entry on the front; `size += 1`; if `size / capacity > max_load`, **resize** to twice the capacity.

delete(key): walk with **two** references, `prev` and `entry`; unlink the match — the head of the chain is the special case; `size -= 1`. Not found: **KeyError**.

Those are the exercises; turning them into Python is your job. Three details.

- **Compare keys with ==, not is.** Two equal strings may be two different objects. hash found the bucket; == finds the entry in it.
- **get has a sentinel default.** get(key) with no default raises KeyError; get(key, "fallback") returns the fallback — and so must get(key, None). The skeleton uses a private object, _MISSING, as the “no default given” marker, because None is a value a caller may legitimately pass (and store). The given __contains__ calls get(key) and catches KeyError for the same reason: a key whose value is None is still **in** the table.
- **delete is Week 5’s remove**, on a list whose head lives in an Array slot instead of a head field. When the match is the first entry, the bucket slot itself must change; otherwise prev.next skips over the match.

Trace: chaining, capacity 8, max_load 0.75

Operation	Index	Buckets (non-empty)	size	load
put 10, 22, 31, 4	2, 6, 7, 4	2:10 · 4:4 · 6:22 · 7:31	4	0.5
put 15	7	... 7: 15 → 31	5	0.625
put 28	4	4: 28 → 4 ...	6	0.75
put 7	7	7 > 0.75·8: resize to 16	7	0.4375
		4:4 · 6:22 · 7:7 · 10:10 · 12:28 · 15: 31 → 15		
delete 31	15	15: 15	6	0.375

Produced by running the reference solution and printing its buckets. Three things to see.

- **0.75 is allowed; above it is not.** After put 28 the load is exactly 0.75, and the test is >: no resize. put 7 makes it 7/8 = 0.875, and the table doubles.
- **After the resize, every key has a new index**, because the index is key % 16 now: 10 moves from bucket 2 to bucket 10, 28 from 4 to 12, 31 from 7 to 15. 15 and 31 still collide — they differ by 16. The chain order came out reversed (31 → 15), because _resize pushes each moved entry on the front of its new chain; order within a chain means nothing.
- **Delete unlinks the head.** 31 was first in bucket 15, so the bucket slot itself now points at 15.

Cost: $1 + \alpha$

Load factor: $\alpha = n/\text{capacity}$ — the average chain length.

- successful get: about $1 + \alpha/2$ entries examined
- unsuccessful get: about α entries — the whole chain
- keep $\alpha \leq 0.75 \Rightarrow O(1)$ **on average**

Worst case: every key in **one** chain — $O(n)$.

Load factor in Arabic: معامل التحميل.

Assume the hash spreads keys as if at random (the “simple uniform hashing” assumption of the textbooks). Then a chain has, on average, α entries. An unsuccessful search walks the whole chain of its bucket: α entries, plus the $O(1)$ cost of hashing. A successful one stops, on average, halfway along the chain of a key that is there: $1 + \alpha/2$. (CLRS, theorems 11.1 and 11.2.) Section 7 shows both formulas holding in measurement, within a percent.


So a hash table is $O(1)$ **only while α is bounded by a constant**. That is the whole reason for the resize: without it, $\alpha = n/8$ and every operation is $O(n)$ — a linked list with an array of eight heads in front. `test_load_factor_stays_bounded` puts 500 keys and checks that the load factor is not above `max_load + 0.05`; a table that never resizes shows 62.5.

The worst case does not go away: $O(n)$ happens when many keys share a bucket, and **no** resize fixes it, because keys with equal hashes collide at every capacity. That is Section 7.

Open Addressing

Linear probing: one key per slot

Open addressing, linear probing: one key per slot, step on when taken



Two Arrays, `_keys` and `_values`. Start at `hash(key) % capacity`; if that slot is taken by another key, try the **next**: $(i + 1) \% \text{capacity}$.

No nodes, no chains: the table **is** the array. A collision is resolved by walking forward to the next free slot, wrapping round the end with `%` exactly like the ring buffer of Lecture 07. The sequence of slots a key tries is its **probe sequence**; with linear probing it is

$$h, h + 1, h + 2, \dots \pmod{\text{capacity}}, \quad h = \text{hash}(\text{key}) \pmod{\text{capacity}}$$

In `dsa/hashmap.py`, `_probe(key)` is a **generator** (Lab 03) that yields those indices, capacity of them at most — every slot once. `put`, `get` and `delete` then all loop for `i` in `self._probe(key)`, and only the loop body differs. If you later want quadratic probing (the docstring’s challenge), you change `_probe` and nothing else.

The figure is the real state after put 10, 22, 31, 4 and 15 into 8 slots. 15 hashes to 7, which 31 holds; the probe moves to $(7 + 1) \% 8 = 0$, which is free. Note what that means for **get(15)**: it must follow the **same** probe sequence, and it cannot stop at slot 7 just because slot 7 holds something else.

When does a search stop? At the key (found), or at a **never-used** slot, None — because put would have put the key there. get (23) on this table: $23 \% 8 = 7$ holds 31, slot 0 holds 15, slot 1 is None — 23 is not in the table, after 3 probes.

Trace: open addressing, capacity 8, max_load 0.66

Operation	Home	Probes	Slots 0–7	size
put 10, 22, 31, 4	2, 6, 7, 4	1 each	_ _ 10 _ 4 _ 22 31	4
put 15	7	7, 0	15 _ 10 _ 4 _ 22 31	5
delete 31	7	7	15 _ 10 _ 4 _ 22 †	4
get 15	7	7 (†, go on), 0	unchanged	4
put 23	7	load check: rebuild first	23 _ 10 _ 4 _ 22 15	5

† is the tombstone.

Produced by the reference solution. The first two rows are the figure. The interesting rows are the last three.

- **delete 31** does not empty slot 7. It leaves a **tombstone** there — the module's TOMB-STONE object — and the reason is the next row.
- **get 15** starts at 7, as it must. Slot 7 holds a tombstone: "something was here once; keys may have probed past it". So the search goes on, to slot 0, and finds 15. Had delete emptied slot 7, get (15) would have stopped there and said "not in the table" — wrong.
- **put 23** runs the load check first, and **counts the tombstone as used**: $(4 \text{ live} + 1 \text{ tombstone} + 1 \text{ new}) / 8 = 0.75 > 0.66$. But the live keys alone, $(4 + 1) / 8 = 0.625$, are within the limit, so the table does not grow: it **rebuilds at the same capacity**, dropping the tombstone. 15 moves back to its home slot 7, and 23 probes 7, then 0. Why count tombstones at all? See the next slide but one.

Primary clustering

Primary clustering: occupied slots clump into runs

a run of 9



40 random keys in 64 slots (load 0.625): 11 runs; the longest are 9, 8, 7, 4. A key that hashes anywhere into a run lands at its end — and makes it longer.

A key that hashes **anywhere** into a run lands at its **end** — and the run grows. Long runs get longer faster.

This is the price of linear probing's simplicity, and it has a name: **primary clustering**. The figure is a real OpenAddressingHashMap with 64 slots and 40 random integer keys: at a

load of only 0.625, the occupied slots are not spread out but clumped into runs, the longest 9 slots long (and the table is a ring, so the run that ends at slot 63 carries on at slot 0).

Why runs feed themselves: a run of length L is hit by any key whose home is in it, or just before it — $L + 1$ homes out of the capacity. So a long run is a bigger target than a short one, and every hit makes it longer still. A search that starts in a run must walk to its end. Knuth's analysis of linear probing (1963, and *TAOCP* vol. 3, §6.4) gives, for a random hash:

$$\text{hit} \approx \frac{1}{2} \left(1 + \frac{1}{1-\alpha} \right), \quad \text{miss} \approx \frac{1}{2} \left(1 + \frac{1}{(1-\alpha)^2} \right)$$

At $\alpha = 0.5$: 1.5 and 2.5 probes. At $\alpha = 0.9$: 5.5 and 50.5. The square in the miss formula is clustering. **Quadratic probing** ($h, h+1, h+4, h+9, \dots$) and **double hashing** (a step size taken from a second hash) break up the runs; they are the docstring's challenge. And yet linear probing is widely used, because consecutive slots are next to each other in memory, and on real hardware a few extra probes in the same cache line cost less than one jump elsewhere.


Delete: why not just empty the slot?

Why delete cannot just empty a slot


3, 11 and 19 all hash to 3: they sit in slots 3, 4, 5



delete 11 by EMPTYING slot 4 — then `get(19)` probes 3, 4: None, 'not here'. Wrong.



delete 11 with a TOMBSTONE — `get(19)` probes 3, 4 (skip), 5: found.



This is the one subtle idea of the week, and the one the tests check on its own (`test_open_addressing_delete_preserves_probe_chain`). Keys 3, 11 and 19 all hash to 3, so they fill slots 3, 4, 5. A probe for 19 **passes through** slots 3 and 4. Empty slot 4 when 11 is deleted, and 19 becomes unreachable: `get(19)` stops at the None in slot 4, because None means “no key ever probed past here”. The table has not lost 19 — it is still in slot 5 — it has lost the **path** to it.

A tombstone — “deleted, but something lay here” — keeps the path open:

Operation	Meets None	Meets TOMBSTONE
get, delete	stop: not found	skip it , keep probing
put	stop: insert here — or at the first tombstone passed	remember the first one; keep probing in case the key is further on
<code>_resize</code>	—	drop it : only live keys are re-inserted

The put row hides a bug worth naming. A put may **reuse** a tombstone's slot — that is the point of reusing space — but it must not stop at the first tombstone, because the key may already be further along the probe sequence. Insert at the first tombstone without looking further, and the table holds the same key **twice** (W13-B3).

Why count tombstones in the load check. Tombstones are not keys, so `len()` and `load_factor` ignore them. But they are not `None` either, and only `None` stops a probe. A table that sees many puts and deletes fills up with tombstones; if the load check counts only live keys, the table never rebuilds, and eventually there is **no** `None` left: every unsuccessful search walks the whole table, $O(n)$, and a probe loop that only stops at `None` never stops at all. So put checks $(\text{size} + \text{tombstones} + 1) / \text{capacity}$. When the live keys alone are within `max_load`, it rebuilds at the **same** capacity — a sweep that throws the tombstones away; otherwise it doubles.

Tombstone in Arabic: شاهد القبر — literally, a gravestone.

The alternative, without tombstones. For linear probing only, a delete can instead **shift back** the keys after the hole that would otherwise be cut off (“backward-shift deletion”; Knuth’s Algorithm R, §6.4). It leaves no tombstones, at the price of a trickier loop. It is an honest option; the course uses tombstones because they work for every probe sequence.

Load Factor and Resizing

Resize: every key moves

$$\alpha = \frac{\text{size}}{\text{capacity}} \quad \text{chaining: } \alpha \leq 0.75 \quad \text{open addressing: } \alpha \leq 0.66$$

When a put would pass the limit: allocate **twice** the buckets and **re-insert every key**.

You **cannot copy the buckets across**: the index is `hash % capacity`, and the capacity just changed.


This is `_resize(capacity)`, the last method of each class. Its docstring says it in one line: “Every key must be rehashed”. A key in bucket 2 of 8 — say 10 — belongs in bucket 10 of 16. Copy `old[i]` to `new[i]` and the table *looks* fine: all 7 keys are there, `len` is right. But `get(10)` now looks in bucket 10, finds nothing, and raises `KeyError`. The lab makes you watch it happen.

Resizing a chaining table **re-links the existing Entry nodes** into the new buckets instead of creating new ones: walk each old chain, and for each entry compute its new index and push it on the front of that chain. Save `entry.next` **before** you relink the entry — relinking overwrites it, and the rest of the old chain is lost (W13-B2 in the question bank).

For open addressing, `_resize` makes two new Arrays, resets `size` and the tombstone count to 0, and re-inserts every **live** key — skipping None and TOMBSTONE. In the new table there are no tombstones and no duplicates, so each key can simply go into the first None of its probe sequence.

Why different limits? A chain can be longer than 1, so a chaining table still works at $\alpha = 1$ or 2 — only slower. An open-addressing table cannot exceed $\alpha = 1$ at all, and by the formulas above its misses cost 5 probes at 0.7 and 50 at 0.9. Java's HashMap resizes at 0.75; CPython's dict at about two-thirds.

Amortised $O(1)$ — Week 4 again



Left: the load factor of each table after every one of 3,000 puts, measured. It climbs to the limit, halves at the resize, and climbs again — a saw-tooth that never leaves the band between about 0.33 and 0.75.

Right: the total number of keys rehashed so far, divided by the number of keys inserted. One resize is expensive — the chaining table's last one moved 1,537 keys at once — but it never rises above 2 per key. The argument is Lecture 04's exactly: resizes happen at sizes $s, s/2, s/4, \dots$, so the total moved is at most $s + s/2 + s/4 + \dots < 2s$. Over n puts, the rehashing costs $O(n)$ in total: **amortised $O(1)$ per put**. Over the 3,000 puts, the chaining table rehashed 3,075 keys in all and the open-addressing table 5,396.


The DynamicArray copied **values**; a hash table must **rehash** them — call `hash()` again and recompute every index. CPython's dict avoids the repeated `hash()` by storing each key's hash beside it; Entry here does not, which keeps it simple and is fine for `int` keys, whose hash is free. Python's `str` caches its own hash inside the string object for the same reason.

Should a table also **shrink** when keys are deleted? Ours never does. Shrinking at a quarter full keeps the memory proportional to the size with the same amortised bound — the same

answer as for a `DynamicArray` that shrinks (Lecture 04). CPython's `dict` does not shrink on delete either.

Measured

Probes against load factor




Every marker is measured on the reference `dsa/hashmap.py`: tables of 16,384 slots, filled with random integer keys to each load factor without resizing, then searched for keys that are there and for keys that are not. For open addressing the count is the number of slots `_probe` yielded; for chaining, the number of entries compared. The dashed lines are the formulas of the previous sections. They agree to within a few percent everywhere.

α	linear probing, hit	linear probing, miss	chaining, hit	chaining, miss
0.5	1.51	2.56	1.25	0.52
0.7	2.21	6.01	1.35	0.70
0.9	5.58	39.74	1.47	0.90
0.95	8.67	108.83	1.49	0.96

Read two things off it. First, **below about 0.7 everything is a handful of probes** — that is what $O(1)$ looks like, and why the limits sit where they do. Second, **open addressing falls off a cliff** as $\alpha \rightarrow 1$, while chaining degrades gently: the miss curve of linear probing grows like $1/(1 - \alpha)^2$. A chaining miss can cost less than one comparison, because many buckets are empty.

What breaks $O(1)$



Keys $i \times 2^{20}$ all end in twenty zero bits: in **every** table of $2^k \leq 2^{20}$ slots they share **one** bucket. Each get becomes $O(n)$.

“ $O(1)$ on average” hides an assumption: the keys spread over the buckets. The figure times one get on tables of n keys, for two kinds of keys.

- **Random keys** (the flat lines): about 1.3 μ s per get for chaining and 3 μ s for open addressing, whatever n is — $O(1)$. (Open addressing is slower here only by a constant: a generator call per probe and two Arrays to read.)
- **Keys $i \times 2^{20}$** (the rising lines): every one is a multiple of 2^{20} , so key % capacity is 0 for any power-of-two capacity up to 2^{20} , and resizing never helps. Every key lands in one chain, or one cluster, and each get costs $O(n)$: about 190 μ s per get at $n = 4,000$ for chaining, 150 times slower than random keys, and the gap doubles with n . Building the table is $O(n^2)$.
- **Python’s dict with the same bad keys** stays at about 0.15 μ s. Its probe sequence mixes the high bits of the hash in (next slide), so keys that agree in their low bits do not stay together.

Adversarial keys. A bad hash is an accident; an attacker is not. In 2003 Crosby and Wallach showed that anyone who can predict a program’s hash function can send it keys that all collide — a denial of service from a few kilobytes of input. In December 2011 Klink and Wälde showed it working against the web frameworks of the day: one HTTP POST with thousands of colliding form-field names kept a server’s CPU busy for minutes, because the framework stored the fields in a hash table. The fix in Python was the **hash randomisation** of Section 2 — optional from 3.2.3, on by default since 3.3, with SipHash since 3.4 (PEP 456): if the attacker cannot predict `hash("...")`, they cannot choose colliding strings. Integers were left unrandomised, as you have just seen, and Java took another route: since Java 8, a HashMap chain longer than 8 turns into a balanced tree, so even a total collision costs $O(\log n)$, not $O(n)$.

So when someone says a hash table is $O(1)$, the honest version is: **$O(1)$ expected, if the hash function spreads the keys; $O(n)$ worst case, and the worst case can be forced**

by whoever chooses the keys.

The Real Thing

Python's dict and set — which you now understand

- **Open addressing**, capacity a power of two, **8** slots to start.
- Grows at about $\frac{2}{3}$ full (your `max_load = 0.66`).
- A deleted key leaves a **dummy** — a tombstone.
- Probes $i = (5 \cdot i + 1 + \text{perturb}) \bmod 2^k$, with `perturb` feeding in the high bits of the hash.
- Since Python 3.7, a dict keeps **insertion order**. A set does not.

Everything on that slide is in `Objects/dictobject.c` of CPython, and every idea in it is one you have now built.

- **Compact dict (3.6)**. A dict keeps two arrays: a small sparse **index table** — the hash table proper, probed exactly like yours — and a dense **entries array** of (hash, key, value) in insertion order. The index table holds only positions in the entries array. Iterating a dict walks the entries array, so it comes out in insertion order, and the language has guaranteed that since 3.7. Your `ChainingHashMap` iterates in bucket order, which depends on the hashes — for string keys, a different order every run.
- **Perturbation**. Instead of $i + 1$, CPython steps with $i = (5 \cdot i + 1 + \text{perturb}) \% 2^k$ and shifts `perturb` right by 5 bits each time; `perturb` starts as the full hash. The first probes use the high bits that `%` threw away, which is why the bad keys above did not hurt it — and once `perturb` reaches 0, $5 \cdot i + 1$ still visits every slot.
- **Dummy entries** are CPython's tombstones, for the same reason as yours; and a resize drops them, as yours does.
- **Each entry stores its hash**, so a resize never calls `hash()` again, and two keys are compared with `==` only when their full hashes are equal.
- A **set** is a separate hash table of keys only, with the same ideas; it probes a few neighbouring slots linearly before perturbing, and it has no insertion order.

in on a list is $O(n)$; in on a set or a dict is $O(1)$ on average. Lab 03 told you that set membership is “ $O(1)$ average”; now you know where the “average” comes from.

Hash table, sorted array or tree?

	Hash table	Sorted array	Balanced BST
find by key	$O(1)$ avg	$O(\log n)$	$O(\log n)$
insert, delete	$O(1)$ avg, amortised	$O(n)$	$O(\log n)$
min, max, in order	$O(n)$ — no order	$O(1), O(n)$	$O(\log n), O(n)$
range $[a, b)$, next larger	no	$O(\log n)$	$O(\log n)$
worst case	$O(n)$	guaranteed	guaranteed

Hash table	Sorted array	Balanced BST
------------	--------------	--------------

The hash table wins at exactly one thing — “is this key here, and what is its value?” — and that one thing is the most common question in software. It loses everything that depends on **order**. `lower_bound` (Week 8) can answer “how many values lie in $[a, b)$?” in $O(\log n)$; a hash table must look at every key, because hashing deliberately destroys order: 2027114 and 2027115 land far apart on purpose. So a database keeps both: a hash index for equality lookups, a tree (a B-tree) for ranges and sorting.

Choose by the questions you will ask. “Have I seen this before?”, “count each word”, “the value for this ID”: hash table. “Everything between these two dates”, “the next appointment after 10:00”, “the smallest”: sorted array or tree.

This Week

Exercises: `dsa/hashtable.py`

Class	Method	The trap
Chaining	<code>get</code>	<code>==</code> not <code>is</code> ; <code>KeyError</code> only without a default
Chaining	<code>put</code>	overwrite, never duplicate; resize above 0.75
Chaining	<code>delete</code>	the head of a chain; <code>size -= 1</code>
Chaining	<code>_resize</code>	rehash every key; save next first
Open addr.	<code>_probe</code>	% capacity; every slot once
Open addr.	<code>put, get</code>	stop at <code>None</code> , not at a tombstone
Open addr.	<code>delete</code>	leave <code>TOMBSTONE</code>
Open addr.	<code>_resize</code>	drop tombstones; reset both counts

```
pytest tests/test_hashtable.py -v
```

Every method is $O(1)$ on average except the two `_resizes`, which are $O(n)$ and make `put` amortised $O(1)$. Seventeen tests: eight run on each class (`-k Chaining` or `-k OpenAddressing` selects them), and one, `test_open_addressing_delete_preserves_probe_chain`, only for open addressing. `-k` matches case-insensitively, but the underscore matters: `-k OpenAddressing` does **not** select the tombstone test; use `-k "OpenAddressing or open_addressing"` for all nine.

Given to you, and worth reading first: `Entry`, both `__init__`s, `_index`, `load_factor`, `__len__`, `__contains__` and `__iter__`. `__iter__` of the chaining table walks every chain, the loop you will write in `_resize`.

Next week needs this. Week 14’s Graph stores its adjacency list in **your** Chaining-HashMap (node $\rightarrow$ neighbours), and BFS and DFS keep their visited set in one. A hash map that loses a key after a resize breaks graph search too.

Homework 13 — before Lecture 14

1. **Implement** `dsa/hashmap.py` until all 17 tests pass.
2. **Trace** a `ChainingHashMap(capacity=4)` through `put 5, 9, 13, 2`, then `delete 9`: the chains after each step, and when it resizes.
3. **Trace** an `OpenAddressingHashMap(capacity=8)` through `put 3, 11, 19`, `delete 11`, `get 19`, `put 27`. Show every probe.
4. **Break it.** Replace `TOMBSTONE` with `None` in your `delete`. Which test fails, and exactly why?
5. **Measure** probes against load factor in `notebooks/13-hash-tables.ipynb`.

For item 2: `max_load` is 0.75, so the fourth `put` makes $4/4 = 1.0$ and the table doubles to 8 — then every index changes.

For item 5, count probes by subclassing: a class that inherits from `OpenAddressingHashMap`, overrides `_probe`, and adds 1 to a counter for every index it passes on from `super()._probe(key)`. `tools/figures_l13.py` does exactly that, in `figure_probes`.

Summary

Seven things to keep

1. **index = hash(key) % capacity** — computed, not searched.
2. Equal keys **must** hash equal. `str` hashes change **between runs**; `int` hashes do not.
3. Collisions are certain and early: $\sim 1.18\sqrt{m}$ keys (birthday paradox).
4. **Chaining**: a linked list per bucket; cost about $1 + \alpha$.
5. **Linear probing**: next free slot; clustering; delete leaves a **tombstone**.
6. Keep α bounded: **double and rehash every key** — amortised $O(1)$.
7. $O(1)$ is the **average** for well-spread keys; the worst case is $O(n)$, and an adversary can force it.

Next

Week 14 — Graphs. The most general structure in the course — stored in the hash map you just built — and the two searches that walk it: breadth-first and depth-first.

Sources and further reading

- **D. E. Knuth.** *The Art of Computer Programming*, vol. 3, *Sorting and Searching*, 2nd ed., Addison-Wesley, 1998, §6.4 “Hashing” — chaining, open addressing, the analysis of linear probing, and deletion without tombstones (Algorithm R).
- **T. H. Cormen, C. E. Leiserson, R. L. Rivest and C. Stein.** *Introduction to Algorithms*, 3rd ed., MIT Press, 2009, chapter 11, “Hash Tables” — theorems 11.1–11.2 (chaining) and 11.6–11.8 (open addressing).
- **S. A. Crosby and D. S. Wallach.** “Denial of Service via Algorithmic Complexity Attacks”, *12th USENIX Security Symposium*, 2003.

- **A. Klink and J. Wälde.** “Efficient Denial of Service Attacks on Web Application Platforms”, 28th Chaos Communication Congress (28C3), 2011.
- **C. Heimes.** PEP 456, “Secure and interchangeable hash algorithm”, 2013 — SipHash for `str` and `bytes` in Python 3.4.
- **CPython source.** `Objects/dictobject.c` — the long comment at the top explains the compact layout, the perturbation probe and the dummy entries; `Objects/setobject.c` for `set`.

Every figure in this lecture is generated by `tools/figures_113.py`. The tables drawn are real states of a working `dsa/hashmap.py`, read back from its slots; the three measured figures will differ slightly on your machine.