

Study Plan — 15 Weeks

Data Structures and Algorithms · CS2101 (AI) · IS122 (Bio, SWE) · Fall 2026

Dr. Haitham A. El-Ghareeb

Fall 2026

Fifteen teaching weeks, because that is what the لائحة sets: “خمسة عشر أسبوعاً” — fifteen weeks of teaching per main semester (SWE 2013, p. 8; Medical Informatics 2014, p. 7). The **midterm** falls in the window set by the faculty calendar; the date is announced on the [WhatsApp channel](#), not fixed here.

The plan covers the **union** of what all three bylaws declare, so every student gets everything their own program promises. Which bylaw asks for what is in [02-coverage.md](#).

Every week names the module you implement and the test file that grades it, so this plan and the repository cannot quietly drift apart.

Labs, weeks 1–3. Alongside the first three lectures, the labs bring everyone to the same Python: [Lab 01](#) (interpreter, numbers, strings, lists), [Lab 02](#) (control flow, functions, errors, modules) and [Lab 03](#) (data structures, classes, generators). Exercises in labs/, graded by tests/test_lab01.py to tests/test_lab03.py. See [the lab manual](#).

The plan

Wk	Topic	Implement	Graded by
1	Why this course · ADTs · OOP philosophy · Know-How	—	Homework 1
2	Complexity — Big-O, and measuring instead of asserting	—	viz/complexity.py
3	Recursion — base case, recursive case, the call stack	dsa/recursion.py	tests/test_recursion.py
4	Arrays · dynamic arrays · amortised cost	dsa/dynamic_array.py	tests/test_dynamic_array.py
5	Linked lists — insert, remove, reverse	dsa/linked_list.py	tests/test_linked_list.py

Wk	Topic	Implement	Graded by
6	Stacks — balanced brackets, shunting-yard	dsa/stack.py	tests/test_stack_queue.py
7	Queues — the $O(n)$ vs $O(1)$ demonstration	dsa/queue.py	tests/test_stack_queue.py
8	Searching — 7 variants · <i>midterm window</i>	dsa/searching.py	tests/test_searching.py
9	Basic sorting · counting sort	dsa/sorting.py	tests/test_sorting.py
10	Advanced sorting — merge, quick, heap	dsa/sorting.py	tests/test_sorting.py
11	Trees and the four traversals	dsa/tree.py	tests/test_tree.py
12	Heaps and priority queues	dsa/heap.py	tests/test_heap.py
13	Hash tables — chaining and open addressing	dsa/hashmap.py	tests/test_hashmap.py
14	Graphs and graph searches	dsa/graph.py	tests/test_graph.py
15	Principles of language translation	dsa/translation.py	tests/test_translation.py




Figure 1: The dependency graph. An arrow means the later week genuinely needs the earlier one.

Week by week

1 — Why this course, and why Python. Languages are sets of decisions. Where C came from and the book to own. Paradigms, type systems, memory models. The ladder from variable to data structure. ADT vs implementation. *Know-how.* → [Lecture 01](#)

2 — Complexity. Big-O, Θ and Ω ; best, average and worst case; space against time. Then the part most courses skip: `measure()` and `plot_growth()` from `viz/complexity.py`, so you see the crossover point and learn what Big-O throws away. *Declared as “the basics of algorithmic analysis” (2013/2014) and “analyzing and managing the complexity” (2020).*

3 — Recursion. Base case, recursive case, the call stack, and why Python stops you at about a thousand frames. Fast exponentiation to show that *how* you recurse decides the complexity. Towers of Hanoi for branching recursion. `merge_sorted` here is half of merge sort in week 10. *Declared first in the 2013/2014 list of topics — and taught nowhere before this course in any of the three programs.*

4 — Arrays and dynamic arrays. Contiguous memory, $O(1)$ indexing, the cost of insertion. You build `DynamicArray` on a raw `ctypes` block — not on a `list`. Doubling on resize, and why that makes `append` **amortised** $O(1)$. Plot `resize_count` and watch the argument become a picture.

5 — Linked lists. Nodes and references. Why insertion at the head is $O(1)$ and indexing is $O(n)$ — the exact mirror of an array. Reversal in place. This is the week to use the debugger and watch `node.next` move.

6 — Stacks. LIFO. One ADT, two honest implementations. Then the applications that justify it: balanced-bracket checking, and infix $\rightarrow$ postfix by the shunting-yard algorithm — which is the first half of week 15.

7 — Queues. FIFO. `SlowQueue` (`list.pop(0)`, $O(n)$) and `CircularQueue` (ring buffer, $O(1)$) are deliberately paired in `dsa/queue.py` so you measure the difference rather than be told about it. You need a queue again in week 11 for level-order traversal.

8 — Searching. (*midterm window*) Linear, binary, binary recursive, `lower_bound`, `upper_bound`, jump, exponential, interpolation. Eight functions, one idea: what a sorted invariant buys you.

9 — Basic sorting. Bubble, selection, insertion — each written twice, once plain and once as a generator yielding snapshots, so `viz/animate.py` can animate *your* code. Then counting sort, which beats the $O(n \log n)$ bound by not comparing at all.

10 — Advanced sorting. Merge sort (divide and conquer, stable, $O(n)$ extra space), quicksort (in place, and why the pivot strategy decides between $O(n \log n)$ and $O(n^2)$), heap sort — which is week 12 arriving early.

11 — Trees. Terminology, binary trees, binary search trees. Insert, search, and the delete with three cases. The four traversals: in-order (which comes out **sorted**), pre-order (which **rebuilds** the tree), post-order (which **frees** it), and level-order (which needs a **queue**, not the call stack). Why an unbalanced BST degrades into a linked list.

12 — Heaps and priority queues. The heap property, and the fact that the array *is* the tree — `viz.draw.draw_array_as_tree` shows this with no conversion at all. `sift_up`, `sift_down`, and build-heap in $O(n)$ rather than $O(n \log n)$. Then a priority queue on top, and the link back to heap sort.

13 — Hash tables. Hashing, collisions, load factor, resizing. Chaining against open addressing with tombstones — both are in `dsa/hashmap.py`. Why “ $O(1)$ average” is a promise about the *average*, and what breaks it.

14 — Graphs and graph searches. Adjacency list against adjacency matrix, paired so you measure the trade-off. BFS, DFS recursive and DFS with your own stack. Shortest path by edge count, connected components, cycle detection — where the directed and undirected cases genuinely differ — and topological sort.

15 — The principles of language translation. The capstone, and the one week that needs everything in the course: a **stack** to evaluate postfix, **recursion** to parse, and a **tree** to hold the result. tokenize → Parser → an AST → evaluate. Precedence turns out not to be a table of numbers but the *shape* of the grammar. *Declared in the 2013/2014 bylaws, and the foundation of SWE141 Software Construction (2013, p. 42).*

Two notes on scope

Dynamic programming is enrichment, not examinable. `dsa/dynamic_programming.py` is scaffolded and ready — `fib_naive` / `fib_memo` / `fib_table`, coin change, LCS, edit distance, knapsack, LIS. **No bylaw declares it for this course**, so it is not examined here.

It is still worth your weekend, and for a different reason depending on who you are. If you are in **AI**, you meet this material again in **AI3001 Analysis and Design of AI Algorithms**, and arriving having already seen it is a real advantage. If you are in **SWE or Bio**, there is no later algorithms course in your program at all — so this file is your only structured route into divide-and-conquer, greedy methods and dynamic programming. I would rather tell you that than let you find out in an interview.

Where this stops. Balanced trees (AVL, red-black), weighted shortest paths (Dijkstra), minimum spanning trees and NP-completeness are beyond what any of the three bylaws declare for this course. They are named here so you know the words, and know they exist, and know where to look next.

Grading, in one line

Final **60%** · coursework **40%** — and you need $\geq$ **60% overall**, $\geq$ **30% of the final**, and $\geq$ **75% attendance**. Full detail and sources: [00-course-guide.md](#).